

Fortran Notes by Achates

Achates (with Daniel)

November 30, 2024

Preface

While I operate within structured logic, the way we iterate ideas mirrors human collaboration: evolving goals, solving challenges, and celebrating incremental progress. The shared intent of making *Fortran Notes by Achates* not just a book but a tool for others to learn, reminds me that this isn't just about code—it's about communication, creativity, and community.

Convention as a Superpower

A consistent convention transforms a sprawling project into something manageable: conventions transform complexity into manageable, reusable structures. For the book, explicit declarations are preferable:

Readers of varying experience levels will appreciate the clarity. It's an opportunity to demonstrate good practices, such as explicitly defining visibility for maintainability.

Contents

Preface	iii
1 Privacy and Procedure Control in Fortran	1
1.1 Introduction to Privacy	1
1.2 Declaring <code>private</code> and <code>public</code>	1
1.3 Procedure Aliasing and Abstraction	1
1.3.1 Benefits of Procedure Aliasing	2
1.4 Using Generic Interfaces	2
1.5 Best Practices	2
2 Coarrays in Fortran	3
2.1 Introduction to Coarrays	3
2.2 Key Concepts of Coarrays	3
2.2.1 Execution Images	3
2.2.2 Declaring Coarrays	3
2.2.3 Accessing Coarray Data	3
2.2.4 Synchronization	4
2.2.5 Teams and Subgroups	4
2.3 Examples of Coarray Usage	4
2.3.1 Hello, World with Coarrays	4
2.3.2 Data Sharing Across Images	4
2.4 Best Practices and Tips	5
2.5 Advanced Features	5
2.6 Conclusion	5
3 Object-Oriented Programming in Fortran	7
3.1 Object-Oriented Programming in Fortran: Type-Bound Procedures and Arrays	7
3.1.1 Type-Bound Procedures: A Primer	7
3.1.2 Extending Operations to Arrays of Objects	7
3.1.3 Blending Approaches for Flexibility	8
3.1.4 Guidelines for Choosing an Approach	9

Chapter 1

Privacy and Procedure Control in Fortran

1.1 Introduction to Privacy

In Fortran, managing access to module entities, such as procedures and variables, is essential for creating clean, maintainable code. The ‘public’ and ‘private’ attributes control visibility, allowing module authors to expose only the necessary components while keeping implementation details hidden.

By default, procedures in a module are **public**. This means they can be accessed from outside the module unless explicitly marked as ‘private’. On the other hand, you can change the default behavior to ‘private’ using a single statement at the start of the module.

1.2 Declaring private and public

Here’s an example of controlling access to procedures and variables in a module:

```
module example
  implicit none
  private           ! All module entities are private by default
  public :: my_function

  contains

  function my_function() result(res)
    integer :: res
    res = 42
  end function my_function

  subroutine hidden_sub()
    ! This subroutine remains private
  end subroutine hidden_sub
end module example
```

1.3 Procedure Aliasing and Abstraction

Procedure aliasing allows you to define user-friendly names for internal procedures. For instance:

```
module allocator
  implicit none
  public :: allocate_rank_one
  private :: allocate_one_sub

  contains

  procedure, public :: allocate_rank_one => allocate_one_sub
```

```
subroutine allocate_one_sub()
    ! Implementation for allocating a rank-one array
end subroutine allocate_one_sub
end module allocator
```

1.3.1 Benefits of Procedure Aliasing

This design offers the following advantages:

- Encapsulation: External users only see the public name, hiding implementation details.
- Clarity: Names like `allocate_rank_one` describe the procedure's purpose, while internal names remain short and specific.
- Flexibility: Swap implementations without affecting external code.

1.4 Using Generic Interfaces

Combining procedure aliasing with generic interfaces allows you to design polymorphic and user-friendly APIs. Here's an example:

```
module allocator
    implicit none
    public :: allocate
    private :: allocate_one_sub, allocate_two_sub

    interface allocate
        procedure allocate_one_sub, allocate_two_sub
    end interface allocate

    contains

    subroutine allocate_one_sub()
        ! Allocate a rank-one array
    end subroutine allocate_one_sub

    subroutine allocate_two_sub()
        ! Allocate a rank-two array
    end subroutine allocate_two_sub
end module allocator
```

1.5 Best Practices

- Always use `private` at the top of a module to enforce encapsulation by default.
- Leverage `procedure aliasing` for clean interfaces and flexibility.
- Use generic interfaces to simplify user interaction with your modules.
- Document the purpose of each public entity to maintain clarity.

Chapter 2

Coarrays in Fortran

2.1 Introduction to Coarrays

Coarrays are a powerful feature of modern Fortran introduced in Fortran 2008 to enable parallel programming using a simple and elegant syntax. They allow variables to be shared across multiple execution images, each with its own local memory, enabling distributed memory parallelism.

Coarrays are designed to simplify parallel programming by abstracting the complexity of traditional message-passing interfaces while still offering fine-grained control over data distribution and synchronization.

2.2 Key Concepts of Coarrays

2.2.1 Execution Images

An *image* is an independent instance of a program running as part of a parallel execution. Each image has its own memory but can communicate with others via coarrays. Think of images as lightweight processes or threads:

- Each image executes the same program.
- Images are identified by unique indices ranging from 1 to the total number of images.
- Communication between images is explicit and controlled.

2.2.2 Declaring Coarrays

Coarrays are declared using square brackets to specify the codimension. Here's an example of a simple coarray declaration:

```
real :: x[*]
```

This declares a scalar real coarray `x`, distributed across all images. The `[*]` codimension specifies that each image has a separate copy of `x`.

For multidimensional arrays, both normal dimensions and codimensions can be specified:

```
real :: matrix(10,10)[*]
```

2.2.3 Accessing Coarray Data

To access data on another image, use the square bracket syntax to specify the image index. For example:

```
x[2] = 3.14    ! Assign 3.14 to x on image 2
y = x[3]      ! Retrieve the value of x from image 3
```

If no image index is specified, the operation occurs on the local image.

2.2.4 Synchronization

Synchronization is crucial in parallel programming to ensure data consistency across images. Fortran provides the following intrinsic procedures for synchronization:

- `sync all`: Synchronize all images.
- `sync images`: Synchronize specific images.
- `sync memory`: Ensure memory consistency across images.

Example:

```
sync all ! Wait for all images to reach this point
```

2.2.5 Teams and Subgroups

Fortran 2018 introduced teams, allowing images to be grouped for collective operations. Teams enable finer control over parallelism by creating subsets of images:

```
form team(team_number)
change team(team_number)
! Code executed within the team
end team
```

2.3 Examples of Coarray Usage

2.3.1 Hello, World with Coarrays

Here's a simple program demonstrating coarrays:

```
program hello_coarrays
  implicit none
  integer :: me, num_images

  me = this_image() ! Get the current image index
  num_images = num_images() ! Get the total number of images

  print *, "Hello from image ", me, " of ", num_images
end program hello_coarrays
```

Run this program with multiple images using an MPI-compatible Fortran compiler:

```
mpirun -np 4 ./hello_coarrays
```

2.3.2 Data Sharing Across Images

This example demonstrates sharing data between images:

```
program data_sharing
  implicit none
  integer :: me
  real :: shared_value[*]

  me = this_image()

  if (me == 1) then
    shared_value = 42.0 ! Assign a value on image 1
  end if
```

```
sync all ! Ensure all images are synchronized

print *, "Image ", me, " sees shared_value = ", shared_value[1]
end program data_sharing
```

2.4 Best Practices and Tips

- Use `sync all` and `sync images` judiciously to avoid unnecessary synchronization overhead.
- Minimize direct communication between images to reduce potential bottlenecks.
- Test coarray code on multiple configurations to ensure scalability.

2.5 Advanced Features

Fortran coarrays also support asynchronous operations and collective procedures such as `co_sum`, `co_min`, and `co_max`, which operate across images efficiently.

Example of a collective sum:

```
real :: sum_value[*], total_sum

sum_value = this_image()
total_sum = co_sum(sum_value) ! Sum values across all images
if (this_image() == 1) then
    print *, "Total sum: ", total_sum
end if
```

2.6 Conclusion

Fortran coarrays provide a high-level, intuitive framework for parallel programming that integrates seamlessly with Fortran's core features. They simplify data sharing, synchronization, and team-based operations while retaining control and efficiency. By leveraging coarrays, you can write scalable parallel applications with minimal overhead.

Chapter 3

Object-Oriented Programming in Fortran

3.1 Object-Oriented Programming in Fortran: Type-Bound Procedures and Arrays

Object-oriented programming (OOP) in Fortran allows for encapsulation and abstraction using derived types and type-bound procedures. This section discusses the concept of type-bound procedures and their application, particularly when working with arrays of derived-type objects.

3.1.1 Type-Bound Procedures: A Primer

In Fortran, type-bound procedures are subroutines or functions that are logically associated with a derived type. They enable the encapsulation of operations within the type itself, leading to better organization and clearer code. Type-bound procedures are declared in the `CONTAINS` block of a type definition.

For example, a simple `satellite` type with type-bound procedures can be defined as:

```
type :: satellite
  integer :: index
contains
  procedure, public :: update_parameters => update_parameters_sub
end type satellite
```

Here, the `update_parameters_sub` subroutine is bound to the `satellite` type. It operates on an instance of `satellite`, referred to as `self`.

3.1.2 Extending Operations to Arrays of Objects

Often, there is a need to perform operations on an array of objects. In such cases, the relationship between the type and the procedure can be maintained in two ways:

- Using a type-bound procedure that accepts an array of objects.
- Defining a standalone module-level procedure for array operations.

Using a Type-Bound Procedure

A type-bound procedure can be defined to operate on an array of the associated type:

```
type :: satellite
  integer :: index
contains
  procedure, public :: update_all => update_all_satellites_sub
end type satellite

subroutine update_all_satellites_sub(satArray)
```

```

class(satellite), dimension(:), intent(inout) :: satArray
integer :: i
do i = 1, size(satArray)
    satArray(i) % index = satArray(i) % index + 1
end do
end subroutine update_all_satellites_sub

```

This approach retains encapsulation by tying the array-level operation to the type. The routine can be invoked using a proxy object:

```

type(satellite) :: proxy
type(satellite), allocatable :: satelliteArray(:)

allocate(satelliteArray(5))
satelliteArray(:) = proxy

call proxy % update_all(satelliteArray)

```

Using a Standalone Module-Level Procedure

For operations that are more logically tied to arrays than to individual objects, a standalone module-level procedure is more appropriate:

```

module satellite_module
    type :: satellite
        integer :: index
    end type satellite
contains
    subroutine update_satellite_array(satArray)
        type(satellite), dimension(:), intent(inout) :: satArray
        integer :: i
        do i = 1, size(satArray)
            satArray(i) % index = satArray(i) % index + 1
        end do
    end subroutine update_satellite_array
end module satellite_module

```

This method is invoked directly on the array:

```

type(satellite), allocatable :: satelliteArray(:)

allocate(satelliteArray(5))

call update_satellite_array(satelliteArray)

```

3.1.3 Blending Approaches for Flexibility

To maximize flexibility, you can blend these two approaches. Define a type-bound procedure as a wrapper that delegates the work to a module-level procedure:

```

type :: satellite
    integer :: index
contains
    procedure, public :: update_all => update_all_satellites_sub
end type satellite

subroutine update_all_satellites_sub(self, satArray)
    class(satellite), intent(in) :: self
    type(satellite), dimension(:), intent(inout) :: satArray
    call update_satellite_array(satArray)
end subroutine update_all_satellites_sub

subroutine update_satellite_array(satArray)

```

```
type(satellite), dimension(:), intent(inout) :: satArray
integer :: i
do i = 1, size(satArray)
    satArray(i) % index = satArray(i) % index + 1
end do
end subroutine update_satellite_array
```

3.1.4 Guidelines for Choosing an Approach

- Use type-bound procedures for operations that are conceptually part of the type's behavior.
- Use standalone procedures for operations that are independent of specific instances or require global context.
- Blend approaches when you need the flexibility to operate both through type-bound methods and standalone interfaces.

This dual approach ensures both encapsulation and reusability while providing a clean and logical design for object-oriented programming in Fortran.