

Unix Tools for Probing Executable Files

Daniel Topa
daniel.topa@hii-tds.com

Mission Technologies
Huntington Ingalls Industries
Kirtland AFB, NM

October 7, 2024

Abstract

Linux offers many tools for the exploration of executable files, some of which depend upon the application being compiled with debug information. A brief survey follows. The `man` pages are included, making this document useful in siloed computing networks.

Contents

1	Overview	2
2	Command Examples	2
2.1	<code>env</code>	2
2.2	<code>ldd</code>	5
2.3	<code>lddconfig</code>	5
2.3.1	Display First Five Current Libraries In The Cache	6
2.3.2	Recursively Display Libraries For Each Directory	6
2.4	<code>locate</code>	6
2.5	<code>lsof</code>	6
2.5.1	<code>lsof</code> on Process ID	6
2.5.2	<code>lsof</code> on User	7
2.6	<code>objdump</code>	8
2.7	<code>readelf</code>	8
2.8	<code>nm</code>	8
2.8.1	Find Headers Only	9
2.8.2	Find Everything	9
2.9	<code>strace</code>	9
2.9.1	Trace System From <code>HelloWorld</code>	9
2.9.2	Trace System Calls To A Given Path	10
2.9.3	Inventory time, calls, and errors for every system call	10
2.9.4	Identify Information Associated With File Descriptors	10
2.10	<code>strings</code>	11
2.10.1	Applied To Compiled Object File	11
2.10.2	Applied To Object Class	12

2.11	<code>gdb</code>	14
3	Manual Pages	15
3.1	<code>env</code> : Show <code>ENVIRONMENT</code> variables.	15
3.2	<code>ldd</code> : Print Shared Object Dependencies	17
3.3	<code>lddconfig</code> : Configure Dynamic Linker Run-time Bindings	18
3.4	<code>locate</code> : List File in Databases	19
3.5	<code>lsof</code> : Show Open Files	22
3.6	<code>objdump</code> : Display Information From Object Files	54
3.7	<code>readelf</code> : Display Information On ELF Files	65
3.8	<code>nm</code> : List Symbols From Object Files	71
3.9	<code>strace</code> : Trace System Calls and Signals	76
3.10	<code>strings</code> : Print Sequences Of Printable Characters	91

1 Overview

Unix provides powerful tools for probing executable files. The following section shows sample usage for each command and the final section contains the information from the `manual` page. The final element is the GNU debugger and not a formal element of Unix.

1. `env`
2. `ldd`
3. `lddconfig`
4. `locate`
5. `objdump`
6. `lsof`
7. `readelf`
8. `nm`
9. `strace`
10. `strings`
11. `gdb`

The goal is to be able to resolve the workings of an executable file exploiting the ELF structure show in figures 1. The next figure, 2, shows the relationship between source files, header files, shared objects, and the executable program.

2 Command Examples

2.1 `env`

A good starting point in debugging is to list environment variables with the command `env` to show where your application is looking for files and libraries. Most prominent are `RPATH` and `LD_LIBRARY_PATH`. These variables can be manipulated using `chrpath` or `patchelf`.

```
dantopa@Quaxolotl.local:~ $ env
build=23A344
bold=
d_python=/Volumes/T7-Touch/repos/github/python
MAC=ac:de:48:00:11:22
dirLsIo=/Users/dantopa/Mathematica_files/io/projects/least squares
gf=/Volumes/T7-Touch/repos/github/f
dbase=dantopa/base-centos:8.5
TERM_PROGRAM=Apple_Terminal
gflags=-g -ffpe-trap=denormal,invalid,zero -fbacktrace -Wall -Wextra -Waliasing -Wsurprising -Wimplicit-procedure -Wi
```




Figure 1: The structure of a Unix ELF file.

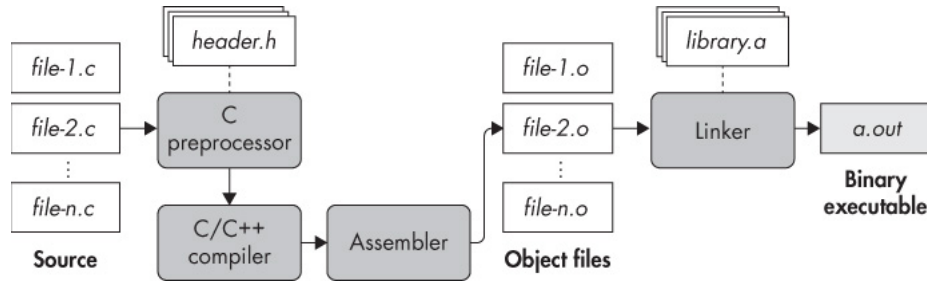


Figure 2: Connecting source files, object files, libraries, and binary executables.

```

dirDockerLocker=
centos_version=8.5
my_log=/Users/dantopa/.info/modules-list.txt
SHELL=/bin/bash
TERM=xterm-256color
machine=MacBookPro16,1
cppflags=-g -Wall -Wextra -Og -pedantic -fmax-errors=5# -fdiagnostics-color=auto
local_spack=/Volumes/T7-Touch/spackitivity/spack-quaxolotl-darwin
TMPDIR=/var/folders/ld/z0sr6fhn0rz4tnd14nfm941m0000gn/T/
reconstructor=/Volumes/T7-Touch/repos/github/reconstructor
repo_results_docker=/Volumes/T7-Touch/repos/github/builds/results-docker
drive_int=Macintosh HD
TERM_PROGRAM_VERSION=453
astra=/Volumes/T7-Touch/repos/github/astra-spack-mirror
gitlab=/Volumes/T7-Touch/repos/gitlab
TERM_SESSION_ID=178C575A-C566-4D30-90CD-2FEB7C9D4E30
os=darwin
gflags48=-g -ffpe-trap=denormal,invalid,zero -fbacktrace -Wall -Wextra -Waliasing -Wsurprising -Wimplicit-procedure -
normal=
USER=dantopa
debian_version=11.3
dir_config=MacBookPro16,1-(quaxolotl)/darwin-23.0.0/Sonoma-14.0.0
jop=/Volumes/T7-Touch/repos/github/jop
lrepos=/Users/dantopa/repos
moniker=quaxolotl
SSH_AUTH_SOCK=/private/tmp/com.apple.launchd.XDYmj6zGeg/Listeners
dcker=/Volumes/T7-Touch/repos/github/docker/unified
repos_bitbucket= icons placemat
bitbucket=/Volumes/T7-Touch/repos/bitbucket
repo_scripts_spack=/Volumes/T7-Touch/repos/github/builds/scripts-spack
drive_ext=T7-Touch
gflags45=-g -ffpe-trap=denormal,invalid,zero -fbacktrace -Wall -Wextra -Waliasing -Wsurprising -Wimplicit-procedure -
repo_build=/Volumes/T7-Touch/repos/github/builds
PATH=/opt/local/bin:/opt/local/sbin:/opt/local/bin:/opt/local/sbin:/usr/local/bin:/System/Cryptexes/App/usr/bin:/usr/
scratch=/Users/dantopa/scratch
volume_ext=/Volumes/T7-Touch
PWD=/Users/dantopa
repo_results_spack=/Volumes/T7-Touch/repos/github/builds/results-spack
mySpack=/Volumes/T7-Touch/spackitivity
capulin_gitlab=dantopa@wtrw:cp-fe:/users/dantopa/repos/gitlab
LANG=en_US.UTF-8
github=/Volumes/T7-Touch/repos/github
serial_num=C02CR18HMD6
  
```

```

capulin=dantopa@wtrw:cp-fe:
configuration=/Volumes/T7-Touch/repos/bitbucket/mac-configurations/MacBookPro16,1-(quaxolotl)/darwin-23.0.0/Sonoma-14
repos_github= builds conferences gbs gf jop d_nursery python reconstructor
bash_scripts=/Volumes/T7-Touch/repos/github/gitlab-bash-scripts
XPC_FLAGS=0x0
stools=/Volumes/T7-Touch/repos/bitbucket/spack-tools
dirDropbox=/spackitivity/mirror
platform=mac
amzn_version=20230
XPC_SERVICE_NAME=0
llnl_ssh=ssh -l topal -X
SHLVL=1
HOME=/Users/dantopa
allflags=-Wall -Wextra -Waliasing -Wsurprising -Wimplicit-procedure -Wintrinsics-std -Wfunction-elimination -Wc-binding
icons=/Volumes/T7-Touch/repos/bitbucket/icons
dscience=dantopa/science-centos:8.5
builds=/Volumes/T7-Touch/repos/github/builds
LOGNAME=dantopa
Wflags=-Wall -Wextra -Waliasing -Wsurprising -Wimplicit-procedure -Wintrinsics-std -Wfunction-elimination -Wc-binding
gflags8=-g -ffpe-trap=denormal,invalid,zero -fbacktrace -Wall -Wextra -Waliasing -Wsurprising -Wimplicit-procedure -W
locker=/Users/dantopa/.info
firmware_v=1715.60.5.0.0 (iBridge: 19.16.10647.0.0,0)
placemat=/Volumes/T7-Touch/repos/bitbucket/placemat
core=/Volumes/T7-Touch/repos/github/gitlab-bash-scripts/core-scripts
vrepos=/Volumes/T7-Touch/repos
gbs=/Volumes/T7-Touch/repos/github/gitlab-bash-scripts
ubuntu_version=22.04
llnl_moniker=topal
file_docker_log=/Volumes/T7-Touch/repos/github/builds/results-docker/vm-log-book.txt
DISPLAY=/private/tmp/com.apple.launchd.0JarZcoW6h/org.xquartz:0
repo_scripts_docker=/Volumes/T7-Touch/repos/github/builds/scripts-docker
d_nursery=/Volumes/T7-Touch/repos/github/nursery-slide-decks
owner=native
SECURITYSESSIONID=24745
dirPyVirtualEnv=/Users/dantopa/scratch/python-virtual-environments
dirPyTruth=/Users/dantopa/scratch/python-virtual-environments/truth
bash_file=.quaxolotl.sh
_=/usr/bin/env
PS1=[\ 033[00;35m]\ u\[\ 033[00m\] @\[\ 033[00;35m\] H\[\ 033[00m\] :\[\ 033[00;33m\] W\[\ 033[00m\] $\ 033[00;3

```

2.2 ldd

The command `ldd` prints shared object dependencies, in this example, for the executable `bash`:

```

root@69cb14a32689:/# ldd /bin/bash
linux-vdso.so.1 (0x00007ffe64317000)
libtinfo.so.6 => /lib/x86_64-linux-gnu/libtinfo.so.6 (0x00007f842112d000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f8420f04000)
/lib64/ld-linux-x86-64.so.2 (0x00007f84212e3000)

```

Symbolic links (symlinks) are highlighted with blue color.

2.3 lddconfig

`lddconfig`: A tool to configure run-time bindings for the dynamic linker. May reside in `/sbin/lddconfig`.

2.3.1 Display First Five Current Libraries In The Cache

```
# ldconfig -p | head -5
916 libs found in cache '/etc/ld.so.cache'
    libzephyr.so.4 (libc6) => /usr/lib/libzephyr.so.4
    libzbar.so.0 (libc6) => /usr/lib/libzbar.so.0
    libz.so.1 (libc6) => /lib/libz.so.1
    libz.so (libc6) => /usr/lib/libz.so
```

2.3.2 Recursively Display Libraries For Each Directory

```
# ldconfig -v | head
/usr/lib/mesa:
    libGL.so.1 -> libGL.so.1.2
/usr/lib/i686-linux-gnu:
    liblouis.so.2 -> liblouis.so.2.2.0
/usr/lib/alsa-lib:
    libasound_module_ctl_oss.so -> libasound_module_ctl_oss.so
    libasound_module_ctl_bluetooth.so -> libasound_module_ctl_bluetooth.so
    libasound_module_pcm_bluetooth.so -> libasound_module_pcm_bluetooth.so
    libasound_module_pcm_vdownmix.so -> libasound_module_pcm_vdownmix.so
    libasound_module_rate_speexrate.so -> libasound_module_rate_speexrate_medium.so
```

2.4 locate

The **locate** command lists files in a prebuilt database of files generated by the **updatedb** command or by a daemon and compressed using incremental encoding.

```
dantopa@92bc4c447e32:/$ locate libc.so.6
/usr/lib/x86_64-linux-gnu/libc.so.6
/usr/lib32/libc.so.6
```

2.5 lsof

This command does an **ls** on open files. The example show how to query both a user and a process id (pid).

2.5.1 lsof on Process ID

The **lsof** command shows open files, here for the bash process with PID = 10932:

```
dantopa@92bc4c447e32:~$ ps
  PID TTY          TIME CMD
 10932 pts/1    00:00:00 bash
 11152 pts/1    00:00:00 ps
dantopa@92bc4c447e32:~$ lsof -p 10932
COMMAND  PID   USER   FD   TYPE DEVICE SIZE/OFF  NODE NAME
bash     10932 dantopa cwd    DIR   0,71    4096  6820049 /home/dantopa
bash     10932 dantopa rtd    DIR   0,71    4096  61653409 /
bash     10932 dantopa txt    REG   0,71  1396520  62702252 /usr/bin/bash
bash     10932 dantopa mem    REG  254,1        62702252 /usr/bin/bash (path dev=0,71)
bash     10932 dantopa mem    REG  254,1    63095938 /usr/lib/x86_64-linux-gnu/libc.so.6 (path dev=0,71)
bash     10932 dantopa mem    REG  254,1    1190606 /usr/lib/x86_64-linux-gnu/libtinfo.so.6.3 (path dev=0,71)
bash     10932 dantopa mem    REG  254,1    63095935 /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2 (path dev=0,71)
bash     10932 dantopa 0u     CHR  136,1        0t0      4 /dev/pts/1
bash     10932 dantopa 1u     CHR  136,1        0t0      4 /dev/pts/1
bash     10932 dantopa 2u     CHR  136,1        0t0      4 /dev/pts/1
bash     10932 dantopa 255u   CHR  136,1        0t0      4 /dev/pts/1
```

2.5.2 lsof on User

These are open files for user dantopa:

dantopa@92bc4c447e32:~\$ lsof -u dantopa

COMMAND	PID	USER	FD	TYPE	DEVICE	SIZE/OFF	NODE	NAME
bash	10921	dantopa	cwd	DIR	0,71	4096	61653409	/
bash	10921	dantopa	rtd	DIR	0,71	4096	61653409	/
bash	10921	dantopa	txt	REG	0,71	1396520	62702252	/usr/bin/bash
bash	10921	dantopa	mem	REG	254,1		62702252	/usr/bin/bash (path dev=0,71)
bash	10921	dantopa	mem	REG	254,1		63095938	/usr/lib/x86_64-linux-gnu/libc.so.6 (path dev=0,71)
bash	10921	dantopa	mem	REG	254,1		1190606	/usr/lib/x86_64-linux-gnu/libtinfo.so.6.3 (path dev=0,71)
bash	10921	dantopa	mem	REG	254,1		63095935	/usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2 (path dev=0,71)
bash	10921	dantopa	0u	CHR	136,0	0t0	3	/dev/pts/0
bash	10921	dantopa	1u	CHR	136,0	0t0	3	/dev/pts/0
bash	10921	dantopa	2u	CHR	136,0	0t0	3	/dev/pts/0
bash	10921	dantopa	255u	CHR	136,0	0t0	3	/dev/pts/0
bash	10932	dantopa	cwd	DIR	0,33	704	1572	/repos/github/vault-fortran/Xmodern-fortran/tau/apex
bash	10932	dantopa	rtd	DIR	0,71	4096	61653409	/
bash	10932	dantopa	txt	REG	0,71	1396520	62702252	/usr/bin/bash
bash	10932	dantopa	mem	REG	254,1		62702252	/usr/bin/bash (path dev=0,71)
bash	10932	dantopa	mem	REG	254,1		63095938	/usr/lib/x86_64-linux-gnu/libc.so.6 (path dev=0,71)
bash	10932	dantopa	mem	REG	254,1		1190606	/usr/lib/x86_64-linux-gnu/libtinfo.so.6.3 (path dev=0,71)
bash	10932	dantopa	mem	REG	254,1		63095935	/usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2 (path dev=0,71)
bash	10932	dantopa	0u	CHR	136,1	0t0	4	/dev/pts/1
bash	10932	dantopa	1u	CHR	136,1	0t0	4	/dev/pts/1
bash	10932	dantopa	2u	CHR	136,1	0t0	4	/dev/pts/1
bash	10932	dantopa	255u	CHR	136,1	0t0	4	/dev/pts/1
lsof	11139	dantopa	cwd	DIR	0,33	704	1572	/repos/github/vault-fortran/Xmodern-fortran/tau/apex
lsof	11139	dantopa	rtd	DIR	0,71	4096	61653409	/
lsof	11139	dantopa	txt	REG	0,71	167544	709329	/usr/bin/lsof
lsof	11139	dantopa	mem	REG	254,1		709329	/usr/bin/lsof (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63095951	/usr/lib/x86_64-linux-gnu/libresolv.so.2 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		1190531	/usr/lib/x86_64-linux-gnu/libkeyutils.so.1.9 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63096020	/usr/lib/x86_64-linux-gnu/libkrb5support.so.0.1 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63096026	/usr/lib/x86_64-linux-gnu/libcom_err.so.2.1 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63096018	/usr/lib/x86_64-linux-gnu/libk5crypto.so.3.1 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63096022	/usr/lib/x86_64-linux-gnu/libkrb5.so.3.3 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		1190578	/usr/lib/x86_64-linux-gnu/libpcr2-8.so.0.10.4 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63096024	/usr/lib/x86_64-linux-gnu/libgssapi_krb5.so.2.2 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63095938	/usr/lib/x86_64-linux-gnu/libc.so.6 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		1190588	/usr/lib/x86_64-linux-gnu/libselinux.so.1 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		1190608	/usr/lib/x86_64-linux-gnu/libtirpc.so.3.0.0 (path dev=0,71)
lsof	11139	dantopa	mem	REG	254,1		63095935	/usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2 (path dev=0,71)
lsof	11139	dantopa	0u	CHR	136,1	0t0	4	/dev/pts/1
lsof	11139	dantopa	1u	CHR	136,1	0t0	4	/dev/pts/1
lsof	11139	dantopa	2u	CHR	136,1	0t0	4	/dev/pts/1
lsof	11139	dantopa	3r	DIR	0,74	0	1	/proc
lsof	11139	dantopa	4r	DIR	0,74	7	123326	/proc/11139/fd
lsof	11139	dantopa	5w	FIFO	0,11	0t0	123331	pipe
lsof	11139	dantopa	6r	FIFO	0,11	0t0	123332	pipe
lsof	11140	dantopa	cwd	DIR	0,33	704	1572	/repos/github/vault-fortran/Xmodern-fortran/tau/apex
lsof	11140	dantopa	rtd	DIR	0,71	4096	61653409	/
lsof	11140	dantopa	txt	REG	0,71	167544	709329	/usr/bin/lsof
lsof	11140	dantopa	mem	REG	254,1		709329	/usr/bin/lsof (path dev=0,71)
lsof	11140	dantopa	mem	REG	254,1		63095951	/usr/lib/x86_64-linux-gnu/libresolv.so.2 (path dev=0,71)
lsof	11140	dantopa	mem	REG	254,1		1190531	/usr/lib/x86_64-linux-gnu/libkeyutils.so.1.9 (path dev=0,71)
lsof	11140	dantopa	mem	REG	254,1		63096020	/usr/lib/x86_64-linux-gnu/libkrb5support.so.0.1 (path dev=0,71)
lsof	11140	dantopa	mem	REG	254,1		63096026	/usr/lib/x86_64-linux-gnu/libcom_err.so.2.1 (path dev=0,71)

```

lsof    11140 dantopa mem    REG  254,1      63096018 /usr/lib/x86_64-linux-gnu/libk5crypto.so.3.1 (path dev=0,7
lsof    11140 dantopa mem    REG  254,1      63096022 /usr/lib/x86_64-linux-gnu/libkrb5.so.3.3 (path dev=0,71)
lsof    11140 dantopa mem    REG  254,1      1190578 /usr/lib/x86_64-linux-gnu/libpcre2-8.so.0.10.4 (path dev=0
lsof    11140 dantopa mem    REG  254,1      63096024 /usr/lib/x86_64-linux-gnu/libgssapi_krb5.so.2.2 (path dev=
lsof    11140 dantopa mem    REG  254,1      63095938 /usr/lib/x86_64-linux-gnu/libc.so.6 (path dev=0,71)
lsof    11140 dantopa mem    REG  254,1      1190588 /usr/lib/x86_64-linux-gnu/libselinux.so.1 (path dev=0,71)
lsof    11140 dantopa mem    REG  254,1      1190608 /usr/lib/x86_64-linux-gnu/libtirpc.so.3.0.0 (path dev=0,71)
lsof    11140 dantopa mem    REG  254,1      63095935 /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2 (path dev=0
lsof    11140 dantopa 4r    FIFO  0,11      0t0     123331 pipe
lsof    11140 dantopa 7w    FIFO  0,11      0t0     123332 pipe

```

2.6 objdump

The **objdump** command shows dependent shared objects, typically libraries. Two versions of the shared library for the GNU standard C library – one 32 bit, the other 64 bit – are located.

```

dantopa@92bc4c447e32:/$ locate libc.so.6
/usr/lib/x86_64-linux-gnu/libc.so.6
/usr/lib32/libc.so.6

```

2.7 readelf

The **readelf** command displays information about ELF files, or Executable and Linkable Format files which are a standard file format for executable files, object code, shared libraries, and core dumps.¹ This example lists the header file for the command **bash**.

```

dantopa@92bc4c447e32:~$ file /bin/bash
/bin/bash: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x
BuildID[sha1]=7a6408ba82a2d86dd98f1f75ac8edcb695f6fd60, for GNU/Linux 3.2.0, stripped
dantopa@92bc4c447e32:~$ readelf -h /bin/bash
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
  Class:                                ELF64
  Data:                                      2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  DYN (Position-Independent Executable file)
  Machine:                               Advanced Micro Devices X86-64
  Version:                               0x1
  Entry point address:                   0x32ef0
  Start of program headers:              64 (bytes into file)
  Start of section headers:              1394600 (bytes into file)
  Flags:                                  0x0
  Size of this header:                   64 (bytes)
  Size of program headers:               56 (bytes)
  Number of program headers:             13
  Size of section headers:               64 (bytes)
  Number of section headers:             30
  Section header string table index:     29

```

2.8 nm

The **nm** command shows dependent shared objects and executables;

¹For an ELF cheatsheet see <https://gist.github.com/x0nu11byt3/bcb35c3de461e5fb66173071a2379779>.

2.8.1 Find Headers Only

```
dantopa@Quaxolotl.local:z2-debug $ nm z2 | grep header
0000000100000000 T __mh_execute_header
```

2.8.2 Find Everything

```
dantopa@Quaxolotl.local:z2-debug $ nm z2
000000010000323c T __Z6moveupPii
000000010000346f T __Z6updatedPP7__sFILE
00000001000032cd T __Z8movedownPii
000000010000335d T __Z9coldstartv
0000000100003e98 s __ZNSt8__detail30__integer_to_chars_is_unsignedIjEE
0000000100003e99 s __ZNSt8__detail30__integer_to_chars_is_unsignedImEE
0000000100003e9a s __ZNSt8__detail30__integer_to_chars_is_unsignedIyEE
0000000100000000 T __mh_execute_header
U _drand48
U _exit
U _exp
U _fclose
U _fopen
U _getcwd
0000000100008000 S _lnk
0000000100003ae8 T _main
U _perror
U _printf
U _srand48
U mcount
```

2.9 strace

The **strace** command reveals how an executable interacts with the operating system.

2.9.1 Trace System From HelloWorld

```
[root@nickdev ~]$ strace ./hello_world
execve("./hello_world", [ "./hello_world" ], [ /* 50 vars */ ]) = 0
brk(0) = 0xa7e000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f7583380000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY) = 3
fstat(3, st_mode=S_IFREG|0644, st_size=77737, ...) = 0
mmap(NULL, 77737, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f758336d000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/lib/libc.so.6", O_RDONLY) = 3
read(3, "\ 177ELF\ 2\ 1\ 1\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 3\ 0>\ 0\ 1\ 0\ 0\ 0\ 0\ 355\ 1\ 0\ 0\ 0\ 0\ 0\ 0"... , 832) = 832
fstat(3, st_mode=S_IFREG|0755, st_size=1432968, ...) = 0
mmap(NULL, 3541032, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7f7582e04000
mprotect(0x7f7582f5c000, 2093056, PROT_NONE) = 0
mmap(0x7f758315b000, 20480, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x157000) = 0x7f758315b000
mmap(0x7f7583160000, 18472, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x7f7583160000
close(3) = 0
mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f758336c000
mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f758336b000
mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f758336a000
arch_prctl1(ARCH_SET_FS, 0x7f758336b700) = 0
```

```

mprotect(0x7f758315b000, 16384, PROT_READ) = 0
mprotect(0x7f7583382000, 4096, PROT_READ) = 0
munmap(0x7f758336d000, 77737) = 0
fstat(1, st_mode=S_IFCHR|0600, st_rdev=makedev(136, 6), ...) = 0
mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f758337f000
write(1, "Hello, world\n", 13Hello, world
) = 13
exit_group(0) = ?

```

2.9.2 Trace System Calls To A Given Path

```

root@169e8b2c1ae3:/# strace -P /etc/ld.so.cache ls /dev/null
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
newfstatat(3, "", st_mode=S_IFREG|0644, st_size=135191, ..., AT_EMPTY_PATH) = 0
mmap(NULL, 135191, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f03bba95000
close(3) = 0
/dev/null
+++ exited with 0 +++

```

2.9.3 Inventory time, calls, and errors for every system call

```

root@169e8b2c1ae3:/# strace -c ls > /dev/null

```

% time	seconds	usecs/call	calls	errors	syscall
71.76	0.013546	6773	2		getdents64
7.85	0.001482	247	6		openat
4.88	0.000922	922	1		execve
4.44	0.000839	49	17		mmap
1.84	0.000347	43	8		close
1.48	0.000279	39	7		mprotect
1.40	0.000265	37	7		newfstatat
1.26	0.000237	47	5		read
0.94	0.000178	44	4		pread64
0.77	0.000145	48	3		brk
0.57	0.000108	36	3	3	ioctl
0.49	0.000092	46	2	2	statfs
0.47	0.000088	44	2	2	access
0.34	0.000065	32	2	1	arch_prctl
0.34	0.000065	65	1		getrandom
0.32	0.000061	61	1		munmap
0.18	0.000034	34	1		rseq
0.17	0.000032	32	1		set_robust_list
0.16	0.000031	31	1		write
0.16	0.000031	31	1		set_tid_address
0.16	0.000031	31	1		prlimit64
100.00	0.018878	248	76	8	total

2.9.4 Identify Information Associated With File Descriptors

```

root@169e8b2c1ae3:/# strace -yy cat /dev/null
execve("/usr/bin/cat", ["cat", "/dev/null"], 0x7fffb8b235d0 /* 10 vars */) = 0
brk(NULL) = 0x5611c6a38000
arch_prctl(0x3001 /* ARCH_??? */, 0x7ffeede990c0) = -1 EINVAL (Invalid argument)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f5c648b8000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD< />, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3< /etc/ld.so.cache>
newfstatat(3< /etc/ld.so.cache>, "", st_mode=S_IFREG|0644, st_size=135191, ..., AT_EMPTY_PATH) = 0

```



```
integer(kind=4)
m-precision-definitions.f08
```

2.10.2 Applied To Object Class

```
dantopa@Quaxolotl.local:bravo $ strings m-cl-soln-basic.o
Error allocating %lu bytes
In file 'm-cl-soln-basic.f08', around line 99
Index '%ld' of dimension 1 of array 'sizes' above upper bound of %ld
At line 98 of file m-cl-soln-basic.f08
Loop iterates infinitely
Index '%ld' of dimension 1 of array 'strides' below lower bound of %ld
Index '%ld' of dimension 1 of array 'strides' above upper bound of %ld
Index '%ld' of dimension 1 of array 'sizes' below lower bound of %ld
Recursive call to nonrecursive procedure 'solnscalingparameter_sub'
At line 69 of file m-cl-soln-basic.f08
m-cl-soln-basic.f08
Array bound mismatch for dimension 1 of array 'u' (%ld/%ld)
At line 85 of file m-cl-soln-basic.f08
Recursive call to nonrecursive procedure 'solncomputesigma_sub'
At line 53 of file m-cl-soln-basic.f08
Array bound mismatch for dimension 1 of array 'me' (%ld/%ld)
At line 60 of file m-cl-soln-basic.f08
Array bound mismatch for dimension 1 of array 'measurements' (%ld/%ld)
At line 61 of file m-cl-soln-basic.f08
Recursive call to nonrecursive procedure 'solnbasic_sub'
At line 33 of file m-cl-soln-basic.f08
Allocatable actual argument 'measurements' is not allocated
At line 38 of file m-cl-soln-basic.f08
Potential error: || u || =
Subroutine: solnScalingParameter_sub
Module:      mClassSolutionBasic
Small size may magnify errors in solution parameter a
GNU Fortran2008 13.2.0 -fPIC -feliminate-unused-debug-symbols -mmacosx-version-min=14.0.0 -mtune=core2 -g -Og -fno-re
m-cl-soln-basic.f08
/Volumes/T7-Touch/repos/github/f/projects/fireball/bravo
sigma
count
phimin
phi2min
sigmasc
residualerror
epsilon
alloc
real(kind=8)
toolkitallocation
requestedgb
alloc_status
sizeelementbits
numelements
numrows
numcols
alloc_message
mykind
integer(kind=4)
__def_init_mclasssolutionbasic_Solnbasic
__mclasssolutionbasic_MOD__def_init_mclasssolutionbasic_Solnbasic
integer(kind=8)
```

```

__vtype_mclasssolutionbasic_Solnbasic
_hash
_size
_extends
_def_init
_copy
_final
_deallocate
solncomputesigma
solnscalingparameter
__class_mclasssolutionbasic_Solnbasic_t
_data
_vptr
data
maxradius
time
radius
alloc
__vtab_mclasssolutionbasic_Solnbasic
__mclasssolutionbasic_MOD__vtab_mclasssolutionbasic_Solnbasic
logical(kind=4)
solnbasic_sub
!#__mclasssolutionbasic_MOD_solnbasic_sub
_descriptor
solncomputesigma_sub
5*__mclasssolutionbasic_MOD_solncomputesigma_sub
solnscalingparameter_sub
E.__mclasssolutionbasic_MOD_solnscalingparameter_sub
normu
__copy_mclasssolutionbasic_Solnbasic
__mclasssolutionbasic_MOD__copy_mclasssolutionbasic_Solnbasic
__final_mclasssolutionbasic_Solnbasic
__mclasssolutionbasic_MOD__final_mclasssolutionbasic_Solnbasic
array
byte_stride
fini_coarray
idx2
ignore
is_contiguous
nelem
offset
ptr2
sizes
strides
__result__final_mclasssoluti
logical(kind=1)
?
4
__def_init_mclasssolutionbasic_Solnbasic
__vtab_mclasssolutionbasic_Solnbasic
__def_init_mclasssolutionbasic_Solnbasic
__vtab_mclasssolutionbasic_Solnbasic
solnbasic_sub
solncomputesigma_sub
solnscalingparameter_sub
__copy_mclasssolutionbasic_Solnbasic
__final_mclasssolutionbasic_Solnbasic
real(kind=8)

```

```

integer(kind=4)
mclasssolutionbasic.toolkitallocation
mclasssolutionbasic.solnbasic
integer(kind=8)
mclasssolutionbasic._.class_mclasssolutionbasic_Solnbasic_t
mclasssolutionbasic.data
mclasssolutionbasic._.vtype_mclasssolutionbasic_Solnbasic
logical(kind=4)
logical(kind=1)
m-cl-soln-basic.f08
measurements
descriptor
solnbasic

```

2.11 gdb

The application **gdb** is an open-source code debugger. When codes are compiled with debug symbols (e.g. `gcc -d ...`) the debugger is a powerful. But even for codes compiled without debug symbols, **gdb** provides helpful information. For example, the following memory error is traced to the Qt library.

```

dantopa@dtopa-latitude-5491:bin $ gdb -ex bt ./MMViz_4.1.12 core
GNU gdb (Ubuntu 9.0.90.20200105-0ubuntu1) 9.0.90.20200105-git
Copyright (C) 2019 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

```

For help, type "help".

Type "apropos word" to search for commands related to "word"...

Reading symbols from ./MMViz_4.1.12...

(No debugging symbols found in ./MMViz_4.1.12)

[New LWP 1885649]

[New LWP 1885710]

[Thread debugging using libthread_db enabled]

Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Core was generated by './MMViz_4.1.12'.

Program terminated with signal SIGSEGV, Segmentation fault.

--Type <RET> for more, q to quit, c to continue without paging--c

#0 0x00000000042730b in

vector_and_utility_module_mp_real_vector_norm.A ()

[Current thread is 1 (Thread 0x7f4673884e00 (LWP 1885649))]

#0 0x00000000042730b in

vector_and_utility_module_mp_real_vector_norm.A ()

#1 0x000000000545b78 in

sie_geometry_module_mp_sie_geometry_tri_compute.A ()

#2 0x000000000643b1d in

mmviz_geometry_module_mp_readgeometry.A ()

#3 0x000000000746e37 in

MMViz::loadFile(QString const&) ()

#4 0x000000000757004 in

MMViz::qt_metacall(QMetaObject::Call, int, void**) ()

#5 0x00007f46754c8f3b in

```

    QMetaObject::activate(QObject*, int, int, void**) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtCore.so.4
#6 0x0000000007567c4 in
    currentUI::loadFile(QString) ()
#7 0x00000000006a9dbd in
    currentUI::createGeometry() ()
#8 0x0000000000756428 in
    currentUI::qt_metacall(QMetaObject::Call, int, void**) ()
#9 0x00007f46754c8f3b in
    QMetaObject::activate(QObject*, int, int, void**) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtCore.so.4
#10 0x00007f46760b3fc9 in
    QAbstractButtonPrivate::click() ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#11 0x00007f46760b418b in
    QAbstractButton::mouseReleaseEvent(QMouseEvent*) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#12 0x00007f4675e9fc2f in
    QWidget::event(QEvent*) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#13 0x00007f4675e6d599 in
    QApplicationPrivate::notify_helper(QObject*, QEvent*) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#14 0x00007f4675e6cef8 in
    QApplication::notify(QObject*, QEvent*) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#15 0x00007f4675eb4095 in
    QETWidget::translateMouseEvent(_XEvent const*) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#16 0x00007f4675ead60f in
    QApplication::x11ProcessEvent(_XEvent*) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#17 0x00007f4675ec5e45 in
    QEventDispatcherX11::processEvents(QFlags<QEventLoop::ProcessEventsFlag>) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtGui.so.4
#18 0x00007f46754b5be7 in
    QEventLoop::processEvents(QFlags<QEventLoop::ProcessEventsFlag>) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtCore.so.4
#19 0x00007f46754b5d17 in
    QEventLoop::exec(QFlags<QEventLoop::ProcessEventsFlag>) ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtCore.so.4
#20 0x00007f46754b92cd in
    QCoreApplication::exec() ()
    from /home/dantopa/Dropbox/2nd-generation/RCS-project/4.1.12/Linux64/bin/libQtCore.so.4
#21 0x0000000000720479
    in main ()

```

3 Manual Pages

3.1 env: Show ENVIRONMENT variables.

NAME

env — run a program in a modified environment

SYNOPSIS

env [OPTION]... [-] [NAME=VALUE]... [COMMAND [ARG]...]

DESCRIPTION

Set each NAME to VALUE in the environment and run COMMAND.

Mandatory arguments to long options are mandatory for short options too.

```

-a, --argv0=ARG
    pass ARG as the zeroth argument of COMMAND

-i, --ignore-environment
    start with an empty environment

-0, --null
    end each output line with NUL, not newline

-u, --unset=NAME
    remove variable from the environment

-C, --chdir=DIR
    change working directory to DIR

-S, --split-string=S
    process and split S into separate arguments; used to pass
    multiple arguments on shebang lines

--block-signal[=SIG]
    block delivery of SIG signal(s) to COMMAND

--default-signal[=SIG]
    reset handling of SIG signal(s) to the default

--ignore-signal[=SIG]
    set handling of SIG signal(s) to do nothing

--list-signal-handling
    list non default signal handling to stderr

-v, --debug
    print verbose information for each processing step

--help display this help and exit

--version
    output version information and exit

A mere - implies -i. If no COMMAND, print the resulting
environment.

SIG may be a signal name like 'PIPE', or a signal number like
'13'. Without SIG, all known signals are included. Multiple
signals can be comma-separated. An empty SIG argument is a
no-op.

Exit status:
125 if the env command itself fails
126 if COMMAND is found but cannot be invoked
127 if COMMAND cannot be found
- the exit status of COMMAND otherwise

OPTIONS
-S/--split-string usage in scripts
The -S option allows specifying multiple parameters in a script.
Running a script named 1.pl containing the following first line:

    #!/usr/bin/env -S perl -w -T
...

Will execute perl -w -T 1.pl .

Without the '-S' parameter the script will likely fail with:

/usr/bin/env: 'perl-w-T': No such file or directory

See the full documentation for more details.

--default-signal[=SIG] usage
This option allows setting a signal handler to its default
action, which is not possible using the traditional shell trap
command. The following example ensures that seq will be
terminated by SIGPIPE no matter how this signal is being handled
in the process invoking the command.

    sh -c 'env --default-signal=PIPE seq inf | _head -n1'

NOTES
POSIX's _exec(3p) _pages_ says:
"many existing applications wrongly assume that they start
with certain signals set to the default action and/or
unblocked. ... Therefore, it is best not to block or ignore
signals across execs without explicit reason to do so, and
especially not to block signals across execs of arbitrary
(not closely cooperating) programs."

AUTHOR
-----Written by Richard Mlynarik, David MacKenzie, and Assaf Gordon.
REPORTING BUGS
-----GNU coreutils online help:

```



```

<https://www.gnu.org/software/coreutils/>
Report any translation bugs to
<https://translationproject.org/team/>
COPYRIGHT
Copyright (C) 2024 Free Software Foundation, Inc. License GPLv3+:
GNU GPL version 3 or later <https://gnu.org/licenses/gpl.html>.
This is free software: you are free to change and redistribute
it. There is NO WARRANTY, to the extent permitted by law.
SEE ALSO
sigaction(2), sigprocmask(2), signal(7)

Full documentation <https://www.gnu.org/software/coreutils/env>
or available locally via: info '(coreutils) env invocation'
COLOPHON
This page is part of the coreutils (basic file, shell and text
manipulation utilities) project. Information about the project
can be found at http://www.gnu.org/software/coreutils/. If you
have a bug report for this manual page, see
http://www.gnu.org/software/coreutils/. This page was obtained
from the tarball coreutils-9.5.tar.xz fetched from
http://ftp.gnu.org/gnu/coreutils/ on 2024-06-14. If you
discover any rendering problems in this HTML version of the page,
or you believe there is a better or more up-to-date source for
the page, or you have corrections or improvements to the
information in this COLOPHON (which is not part of the original
manual page), send a mail to man-pages@man7.org

GNU coreutils 9.5 March 2024 ENV(1)

```

3.2 ldd: Print Shared Object Dependencies

NAME

ldd — print shared object dependencies

SYNOPSIS

ldd [option]... file...

DESCRIPTION

ldd prints the shared objects (shared libraries) required by each program or shared object specified on the **command** line. An example of its use and output is the following:

```

$ ldd /bin/ls
linux-vdso.so.1 (0x00007ffcc3563000)
libselinux.so.1 => /lib64/libselinux.so.1 (0x00007f87e5459000)
libcap.so.2 => /lib64/libcap.so.2 (0x00007f87e5254000)
libc.so.6 => /lib64/libc.so.6 (0x00007f87e4e92000)
libpcre.so.1 => /lib64/libpcre.so.1 (0x00007f87e4c22000)
libdl.so.2 => /lib64/libdl.so.2 (0x00007f87e4a1e000)
/lib64/ld-linux-x86-64.so.2 (0x00005574bf12e000)
libattr.so.1 => /lib64/libattr.so.1 (0x00007f87e4817000)
libpthread.so.0 => /lib64/libpthread.so.0 (0x00007f87e45fa000)

```

In the usual case, ldd invokes the standard dynamic linker (see ld.so(8)) with the LD_TRACE_LOADED_OBJECTS environment variable set to 1. This causes the dynamic linker to inspect the program's dynamic dependencies, and find (according to the rules described in ld.so(8)) and load the objects that satisfy those dependencies. For each dependency, ldd displays the location of the matching object and the (hexadecimal) address at which it is loaded. (The linux-vdso and ld-linux shared dependencies are special; see vdso(7) and ld.so(8).)

Security

Be aware that in some circumstances (e.g., where the program specifies an ELF interpreter other than ld-linux.so), some versions of ldd may attempt to obtain the dependency information by attempting to directly execute the program, which may lead to the execution of whatever code is defined in the program's ELF interpreter, and perhaps to execution of the program itself. (Before glibc 2.27, the upstream ldd implementation did this for example, although most distributions provided a modified version that did not.)

Thus, you should never employ ldd on an untrusted executable, since this may result in the execution of arbitrary code. A safer alternative when dealing with untrusted executables is:

```
$ objdump -p /path/to/program | grep NEEDED
```

Note, however, that this alternative shows only the direct dependencies of the executable, while ldd shows the entire dependency tree of the executable.

OPTIONS

```

--version
    Print the version number of ldd.

--verbose
-v
    Print all information, including, for example, symbol
    versioning information.

```

```

--unused
-u      Print unused direct dependencies.  (Since glibc 2.3.4.)

--data-relocs
-d      Perform relocations and report any missing objects (ELF
        only).

--function-relocs
-r      Perform relocations for both data objects and functions,
        and report any missing objects or functions (ELF only).

--help  Usage information.

```

BUGS

ldd does not work on a.out shared libraries.

ldd does not work with some extremely old a.out programs which were built before ldd support was added to the compiler releases. If you use ldd on one of these programs, the program will attempt to run with `argc = 0` and the results will be unpredictable.

SEE ALSO

pldd(1), sprof(1), ld.so(8), ldconfig(8)

COLOPHON

This page is part of the man-pages (Linux kernel and C library user-space interface documentation) project. Information about the project can be found at <https://www.kernel.org/doc/man-pages/>. If you have a bug report for this manual page, see <https://git.kernel.org/pub/scm/docs/man-pages/man-pages.git/tree/CONTRIBUTING>. This page was obtained from the tarball man-pages-6.9.1.tar.gz fetched from <https://mirrors.edge.kernel.org/pub/linux/docs/man-pages/> on 2024-06-26. If you discover any rendering problems in this HTML version of the page, or you believe there is a better or more up-to-date source for the page, or you have corrections or improvements to the information in this COLOPHON (which is not part of the original manual page), send a mail to man-pages@man7.org

Linux man-pages 6.9.1 2024-05-02 ldd(1)

3.3 ldconfig: Configure Dynamic Linker Run-time Bindings

NAME

ldconfig - configure dynamic linker run-time bindings

SYNOPSIS

```

/sbin/ldconfig [-nNvVX] [-C cache] [-f conf] [-r root]
               directory ...

/sbin/ldconfig -l [-v] library ...

/sbin/ldconfig -p

```

DESCRIPTION

ldconfig creates the necessary links and cache to the most recent shared libraries found in the directories specified on the command line, in the file `/etc/ld.so.conf`, and in the trusted directories, `/lib` and `/usr/lib`. On some 64-bit architectures such as x86-64, `/lib` and `/usr/lib` are the trusted directories for 32-bit libraries, while `/lib64` and `/usr/lib64` are used for 64-bit libraries.

The cache is used by the run-time linker, `ld.so` or `ld-linux.so`. `ldconfig` checks the header and filenames of the libraries it encounters when determining which versions should have their links updated. `ldconfig` should normally be run by the superuser as it may require write permission on some root owned directories and files.

`ldconfig` will look only at files that are named `lib*.so*` (for regular shared objects) or `ld-*.so*` (for the dynamic loader itself). Other files will be ignored. Also, `ldconfig` expects a certain pattern to how the symbolic links are set up, like this example, where the middle file (`libfoo.so.1` here) is the SONAME for the library:

```

libfoo.so -> libfoo.so.1 -> libfoo.so.1.12

```

Failure to follow this pattern may result in compatibility issues after an upgrade.

OPTIONS

```

--format=fmt
-c fmt  (Since glibc 2.2) Use cache format fmt, which is one of
        old, new, or compat. Since glibc 2.32, the default is
        new. Before that, it was compat.

-C cache
        Use cache instead of /etc/ld.so.cache.

```

```

-f conf
    Use conf instead of /etc/ld.so.conf.

--ignore-aux-cache
-i      (Since glibc 2.7) Ignore auxiliary cache file.

-l      (Since glibc 2.2) Interpret each operand as a library name
        and configure its links. Intended for use only by
        experts.

-n      Process only the directories specified on the command
        line; don't process the trusted directories, nor those
        specified in /etc/ld.so.conf. Implies -N.
-----N-----Don't rebuild the cache. Unless -X is also specified,
        links are still updated.

--print-cache
-p      Print the lists of directories and candidate libraries
        stored in the current cache.

-r root
        Change to and use root as the root directory.

--verbose
-v      Verbose mode. Print current version number, the name of
        each directory as it is scanned, and any links that are
        created. Overrides quiet mode.

--version
-V      Print program version.

-X      Don't update links. Unless -N is also specified, the
        cache is still rebuilt.
FILES
-----/lib/ld.so
        is the run-time linker/loader.
-----/etc/ld.so.conf
        contains a list of directories, one per line, in which to
        search for libraries.
-----/etc/ld.so.cache
        contains an ordered list of libraries found in the
        directories specified in /etc/ld.so.conf, as well as those
        found in the trusted directories.
SEE ALSO
-----ldd(1), ld.so(8)
COLOPHON
-----This page is part of the man-pages (Linux kernel and C library
        user-space interface documentation) project. Information about
        the project can be found at
        https://www.kernel.org/doc/man-pages/. If you have a bug report
        for this manual page, see
        https://git.kernel.org/pub/scm/docs/man-pages/man-pages.git/tree/CONTRIBUTING.
        This page was obtained from the tarball man-pages-6.9.1.tar.gz
        fetched from
        https://mirrors.edge.kernel.org/pub/linux/docs/man-pages/on
        2024-06-26. If you discover any rendering problems in this HTML
        version of the page, or you believe there is a better or more up-
        to-date source for the page, or you have corrections or
        improvements to the information in this COLOPHON (which is not
        part of the original manual page), send a mail to
        man-pages@man7.org
Linux man-pages_6.9.1-----2024-05-02-----ldconfig(8)

```

3.4 locate: List File in Databases

```

NAME
    locate - list files in databases that match a pattern

SYNOPSIS
    locate [-d path | --database=path] [-e | -E | --[non-]existing]
    [-i | --ignore-case] [-0 | --null] [-c | --count] [-w |
    --wholename] [-b | --basename] [-l N | --limit=N] [-S |
    --statistics] [-r | --regex] [--regextype R] [--max-database-age
    D] [-P | -H | --nofollow] [-L | --follow] [--version] [-A |
    --all] [-p | --print] [--help] pattern...

DESCRIPTION
    This manual page documents the GNU version of locate. For each
    given pattern, locate searches one or more databases of file
    names and displays the file names that contain the pattern.
    Patterns can contain shell-style metacharacters: '*', '?', and
    '['. The metacharacters do not treat '/' or '.' specially.
    -----Therefore, a pattern 'foo*bar' can match a file name that
    contains 'foo3/bar', and a pattern 'aduck*' can match a file name
    that contains 'lake/ducky'. Patterns that contain
    -----metacharacters should be quoted to protect them from expansion by
    -----the shell.

```

```

-----If a pattern is a plain string, it contains no metacharacters.
-----locate displays all file names in the database that contain that
-----string anywhere. If a pattern does contain metacharacters,
-----locate only displays file names that match the pattern exactly.
-----As a result, patterns that contain metacharacters should usually
-----begin with a '*', and will most often end with one as well. The
-----exceptions are patterns that are intended to explicitly match the
-----beginning or end of a file name.

The file name databases contain lists of files that were on the
system when the databases were last updated. The system
administrator can choose the file name of the default database,
the frequency with which the databases are updated, and the
directories for which they contain entries; see updatedb(1).

If locate's output is going to a terminal, unusual characters in
the output are escaped in the same way as for the print action
of the find command. If the output is not going to a terminal,
file names are printed exactly as-is.
OPTIONS
-----0, null
-----Use ASCII_NUL as a separator, instead of newline.

-----A, all
-----Print only names which match all non-option arguments, not
-----those matching one or more non-option arguments.

-----b, basename
-----Results are considered to match if the pattern specified
-----matches the final component of the name of a file as
-----listed in the database. This final component is usually
-----referred to as the 'base name'.

-c, --count
-----Instead of printing the matched filenames, just print the
-----total number of matches we found, unless --print (-p) is
-----also present.

-d path, --database=path
-----Instead of searching the default file name database,
-----search the file name databases in path, which is a colon-
-----separated list of database file names. You can also use
-----the environment variable LOCATE_PATH to set the list of
-----database files to search. The option overrides the
-----environment variable if both are used. Empty elements in
-----the path are taken to be synonyms for the file name of the
-----default database. A database can be supplied on stdin,
-----using '-' as an element of path. If more than one element
-----of path is '-', later instances are ignored (and a warning
-----message is printed).

The file name database format changed starting with GNU
find and locate version 4.0 to allow machines with
different byte orderings to share the databases. This
version of locate can automatically recognize and read
databases produced for older versions of GNU locate or
Unix versions of locate or find. Support for the old
locate database format will be discontinued in a future
release.

-e, --existing
-----Only print out such names that currently exist (instead of
-----such names that existed when the database was created).
-----Note that this may slow down the program a lot, if there
-----are many matches in the database. If you are using this
-----option within a program, please note that it is possible
-----for the file to be deleted after locate has checked that
-----it exists, but before you use it.

-E, --non-existing
-----Only print out such names that currently do not exist
----- (instead of such names that existed when the database was
----- created). Note that this may slow down the program a lot,
----- if there are many matches in the database.

--help Print a summary of the options to locate and exit.

-i, --ignore-case
-----Ignore case distinctions in both the pattern and the file
----- names.

-l N, --limit=N
-----Limit the number of matches to N. If a limit is set via
----- this option, the number of results printed for the -c
----- option will never be larger than this number.

-L, --follow
-----If testing for the existence of files (with the -e or -E
----- options), consider broken symbolic links to be non-
----- existing. This is the default.

--max-database-age D

```

Normally, locate will issue a warning message when it searches a database which is more than 8 days old. This option changes that value to something other than 8. The effect of specifying a negative value is undefined.

-m, --mmap
Accepted but does nothing, for compatibility with BSD locate.

-P, -H, --nofollow
If testing for the existence of files (with the **-e** or **-E** options), treat broken symbolic links as if they were existing files. The **-H** form of this option is provided purely for similarity with **find**; the use of **-P** is recommended over **-H**.

-p, --print
Print search results when they normally would not, because of the presence of **--statistics** (**-S**) or **--count** (**-c**).

-r, --regex
The pattern specified on the **command** line is understood to be a regular expression, as opposed to a glob pattern. The Regular expressions work in the same way as in **emacs** except for the fact that **"** will match a newline. GNU **find** uses the same regular expressions. Filenames whose full paths match the specified regular expression are printed (or, in the case of the **-c** option, counted). If you wish to anchor your regular expression at the ends of the full path name, then as is usual with regular expressions, you should use the characters **^** and **\$** to signify this.

--regextype R
Use regular expression dialect **R**. Supported dialects include **'findutils-default'**, **'posix-awk'**, **'posix-basic'**, **'posix-egrep'**, **'posix-extended'**, **'posix-minimal-basic'**, **'awk'**, **'ed'**, **'egrep'**, **'emacs'**, **'gnu-awk'**, **'grep'** and **'sed'**. See the **Texinfo** documentation for a detailed explanation of these dialects.

-s, --stdio
Accepted but does nothing, for compatibility with BSD locate.

-S, --statistics
Print various statistics about each locate database and then exit without performing a search, unless **--non-option** arguments are given. For compatibility with BSD, **-S** is accepted as a synonym for **--statistics**. However, the output of **locate -S** is different for the GNU and BSD implementations of **locate**.

--version
Print the version number of **locate** and exit.

-w, --wholename
Match against the whole name of the file as listed in the database. This is the default.

ENVIRONMENT

LOCATE_PATH
Colon-separated list of databases to search. If the value has a leading or trailing colon, or has two colons in a row, you may get results that vary between different versions of **locate**.

HISTORY

The **locate** program started life as the BSD **fast-find** program, contributed to BSD by James A. Woods. This was described by his paper **Finding Files Fast** which was published in **Usenix**; login:, Vol. 8, No. 1, February/March, 1983, pp. 8-10. When the **find** program began to assume a default **--print** action if no action was specified, this changed the interpretation of **find** pattern. The BSD developers therefore moved the fast **find** functionality into **locate**. The GNU implementation of **locate** appears to be derived from the same code.

Significant changes to locate in reverse order:

- 4.3.7 ----- Byte-order independent support for old database format
- 4.3.3 ----- **locate -i** supports multi-byte characters correctly
- Introduced **--max-db-age**
- 4.3.2 ----- Support for the **slocate** database format
- 4.2.22 ----- Introduced the **--all** option
- 4.2.15 ----- Introduced the **--regex** option
- 4.2.14 ----- Introduced options **-L**, **-P**, **-H**
- 4.2.12 ----- Empty items in **LOCATE_PATH** now indicate the default database
- 4.2.11 ----- Introduced the **--statistics** option
- 4.2.4 ----- Introduced **--count** and **--limit**
- 4.2.0 ----- Glob characters cause matching against the whole file name
- 4.0 ----- Introduced the **LOCATE02** database format
- 3.7 ----- **Locate** can search multiple databases

BUGS

The **locate** database correctly handles filenames containing

-----newlines, but only if the system's **sort** command has a working **-z** option. If you suspect that **locate** may need to **return** filenames containing newlines, consider using its **—null** option.

REPORTING BUGS

GNU findutils online **help**:
[<https://www.gnu.org/software/findutils/#get-help>](https://www.gnu.org/software/findutils/#get-help)
 Report any translation bugs to
[<https://translationproject.org/team/>](https://translationproject.org/team/)
 Report any other issue via the form at the GNU Savannah bug tracker:
[<https://savannah.gnu.org/bugs/?group=findutils>](https://savannah.gnu.org/bugs/?group=findutils)
 General topics about the GNU findutils package are discussed at the bug-findutils mailing list:
[<https://lists.gnu.org/mailman/listinfo/bug-findutils>](https://lists.gnu.org/mailman/listinfo/bug-findutils)

COPYRIGHT

Copyright (C) 1994–2024 Free Software Foundation, Inc. License GPLv3+: GNU GPL version 3 or later
[<https://gnu.org/licenses/gpl.html>.](https://gnu.org/licenses/gpl.html)
 This is free software: you are free to change and redistribute it. There is NO WARRANTY, to the extent permitted by law.

SEE ALSO

find(1), **updatedb(1)**, **xargs(1)**, **glob(3)**, **locatedb(5)**

Full documentation
[<https://www.gnu.org/software/findutils/locate>](https://www.gnu.org/software/findutils/locate)
 or available locally via: **info locate**

COLOPHON

This page is part of the findutils (find utilities) project. Information about the project can be found at <http://www.gnu.org/software/findutils/>. If you have a bug report for this manual page, see <https://savannah.gnu.org/bugs/?group=findutils>. This page was obtained from the project's upstream Git repository <https://git.savannah.gnu.org/findutils.git> on 2024-06-14. (At that time, the date of the most recent commit that was found in the repository was 2024-06-03.) If you discover any rendering problems in this HTML version of the page, or you believe there is a better or more up-to-date source for the page, or you have corrections or improvements to the information in this COLOPHON (which is not part of the original manual page), send a mail to man-pages@man7.org

-----LOCATE(1)

3.5 lsuf: Show Open Files

NAME

lsuf — list open files

SYNOPSIS

```
lsuf [ -?abChlnNOPRtUvVX ] [ -A A ] [ -c c ] [ +c c ] [ +|-d d ]
[ +|-D D ] [ +|-e s ] [ +|-E ] [ +|-f [cfgGn] ] [ -F [f] ] [ -g
[s] ] [ -i [i] ] [ -k k ] [ -K k ] [ +|-L [l] ] [ +|-m m ] [ +|-M
[ ] [ -o [o] ] [ -p s ] [ +|-r [t[m<fmt>]] ] [ -s [p:s] ] [ -S [t]
[ -T [t] ] [ -u s ] [ +|-w ] [ -x [fl] ] [ -z [z] ] [ -Z [Z] ]
[ - ] [ names ]
```

DESCRIPTION

Lsuf revision 4.91 lists on its standard output file information about files opened by processes for the following UNIX dialects:

Apple Darwin 9 and Mac OS X 10.[567]
 FreeBSD 8.[234], 9.0 and 1[012].0 for AMD64-based systems
 Linux 2.1.72 and above for x86-based systems
 Solaris 9, 10 and 11

(See the DISTRIBUTION section of this manual page for information on how to obtain the latest lsuf revision.)

An open file may be a regular file, a directory, a block special file, a character special file, an executing text reference, a library, a stream or a network file (Internet socket, NFS file or UNIX domain socket.) A specific file or all the files in a file system may be selected by path.

Instead of a formatted display, lsuf will produce output that can be parsed by other programs. See the **-F**, option description, and the OUTPUT FOR OTHER PROGRAMS section for more information.

In addition to producing a single output list, lsuf will run in repeat mode. In repeat mode it will produce output, delay, then repeat the output operation until stopped with an interrupt or quit signal. See the **+|-r [t[m<fmt>]]** option description for more information.

OPTIONS

In the absence of any options, lsuf lists all open files belonging to all active processes.

If any list request option is specified, other list requests must

be specifically requested — e.g., if `-U` is specified for the listing of UNIX socket files, NFS files won't be listed unless `-N` is also specified; or if `-a` user list is specified with the `-u` option, UNIX domain socket files, belonging to users not in the list, won't be listed unless the `-U` option is also specified.

Normally list options that are specifically stated are ORed — i.e., specifying the `-i` option without an address and the `-ufoo` option produces a listing of all network files OR files belonging to processes owned by user `'foo'`. The exceptions are:

- 1) the `'^'` (negated) login name or user ID (UID), specified with the `-u` option;
- 2) the `'^'` (negated) process ID (PID), specified with the `-p` option;
- 3) the `'^'` (negated) process group ID (PGID), specified with the `-g` option;
- 4) the `'^'` (negated) command, specified with the `-c` option;
- 5) the `'^'` (negated) TCP or UDP protocol state names, specified with the `-s[p:s]` option.

Since they represent exclusions, they are applied without ORing or ANDing and take effect before any other selection criteria are applied.

The `-a` option may be used to AND the selections. For example, specifying `-a, -U, -ufoo` produces a listing of only UNIX socket files that belong to processes owned by user `'foo'`.

Caution: the `-a` option causes all list selection options to be ANDed; it can't be used to cause ANDing of selected pairs of selection options by placing it between them, even though its placement there is acceptable. Wherever `-a` is placed, it causes the ANDing of all selection options.

Items of the same selection set — command names, file descriptors, network addresses, process identifiers, user identifiers, zone names, security contexts — are joined in a single ORed set and applied before the result participates in ANDing. Thus, for example, specifying `-i@aaa.bbb, -i@ccc.ddd, -a, -ufoo, -gfoo` will select the listing of files that belong to either login `'foo'` OR `'ggg'` AND have network connections to either host `aaa.bbb` OR `ccc.ddd`.

Options may be grouped together following a single prefix — e.g., the option set `'-a -b -C'` may be stated as `-abC`. However, since values are optional following `+-f, -F, -g, -i, +-L, -O, +-r, -s, -S, -T, -x` and `-z`, when you have no values for them be careful that the following character isn't ambiguous. For example, `-Fn` might represent the `-F` and `-n` options, or it might represent the `n` field identifier character following the `-F` option. When ambiguity is possible, start a new option with a `'-'` character — e.g., `'-F -n'`. If the next option is a file name, follow the possibly ambiguous option with `'--'` — e.g., `'-F -- name'`.

Either the `+` or the `-` prefix may be applied to a group of options. Options that don't take on separate meanings for each prefix — e.g., `-i` — may be grouped under either prefix. Thus, for example, `+'+M -i'` may be stated as `+'+Mi'` and the group means the same as the separate options. Be careful of prefix grouping when one or more options in the group does take on separate meanings under different prefixes — e.g., `+-M; '-iM'` is not the same request as `+'-i+M'`. When in doubt, use separate options with appropriate prefixes.

?-h These two equivalent options select a usage (help) output list. `-lsof` displays a shortened form of this output when it detects an error in the options supplied to it, after it has displayed messages explaining each error. (Escape the `'?'` character as your shell requires.)

`-a` causes list selection options to be ANDed, as described above.

`-A A` is available on systems configured for AFS whose AFS kernel code is implemented via dynamic modules. It allows the `ls` user to specify `A` as an alternate name list file where the kernel addresses of the dynamic modules might be found. See the `ls` FAQ (The FAQ section gives its location.) for more information about dynamic modules, their symbols, and how they affect `ls`.

`-b` causes `ls` to avoid kernel functions that might block — `lstat(2)`, `readlink(2)`, and `stat(2)`.

See the BLOCKS AND TIMEOUTS and AVOIDING KERNEL BLOCKS sections for information on using this option.

```

-c c    selects the listing of files for processes executing the
        command that begins with the characters of c. Multiple
        commands may be specified, using multiple -c options.
        They are joined in a single ORed set before participating
        in AND option selection.

        If c begins with a '^', then the following characters
        specify a command name whose processes are to be ignored
        (excluded.)

        If c begins and ends with a slash ('/'), the characters
        between the slashes are interpreted as a regular
        expression. Shell meta-characters in the regular
        expression must be quoted to prevent their interpretation
        by the shell. The closing slash may be followed by these
        modifiers:

        b the regular expression is a basic one.
        i ignore the case of letters.
        x the regular expression is an extended one
        (default).

        See the lsof FAQ (The FAQ section gives its location.)
        for more information on basic and extended regular
        expressions.

        The simple command specification is tested first. If that
        test fails, the command regular expression is applied. If
        the simple command test succeeds, the command regular
        expression test isn't made. This may result in 'no
        command found for regex:' messages when lsof's -V option
        is specified.

        +c w defines the maximum number of initial characters of the
        name, supplied by the UNIX dialect, of the UNIX command
        associated with a process to be printed in the COMMAND
        column. (The lsof default is nine.)

        Note that many UNIX dialects do not supply all command
        name characters to lsof in the files and structures from
        which lsof obtains command name. Often dialects limit the
        number of characters supplied in those sources. For
        example, Linux 2.4.27 and Solaris 9 both limit command
        name length to 16 characters.

        If w is zero ('0'), all command characters supplied to
        lsof by the UNIX dialect will be printed.

        If w is less than the length of the column title,
        'COMMAND', it will be raised to that length.

        -C disables the reporting of any path name components from
        the kernel's name cache. See the KERNEL NAME CACHE
        section for more information.

+ d s    causes lsof to search for all open instances of directory
        s and the files and directories it contains at its top
        level. +d does NOT descend the directory tree, rooted at
        s. The +D D option may be used to request a full-descent
        directory tree search, rooted at directory D.

        Processing of the +d option does not follow symbolic links
        within s unless the -x or -x l option is also specified.
        Nor does it search for open files on file system mount
        points on subdirectories of s unless the -x or -x f
        option is also specified.

        Note: the authority of the user of this option limits it
        to searching for files that the user has permission to
        examine with the system stat(2) function.

- d s    specifies a list of file descriptors (FDs) to exclude from
        or include in the output listing. The file descriptors
        are specified in the comma-separated set s - e.g.,
        'cwd,1,3', '^6,2'. (There should be no spaces in the
        set.)

        The list is an exclusion list if all entries of the set
        begin with '^'. It is an inclusion list if no entry
        begins with '^'. Mixed lists are not permitted.

        A file descriptor number range may be in the set as long
        as neither member is empty, both members are numbers, and
        the ending member is larger than the starting one - e.g.,
        '0-7' or '3-10'. Ranges may be specified for
        exclusion if they have the '^' prefix - e.g., '^0-7'
        excludes all file descriptors 0 through 7.

        Multiple file descriptor numbers are joined in a single
        ORed set before participating in AND option selection.

```



```

-----When there are exclusion and inclusion members in the set,
-----ls of reports them as errors and exits with a non-zero
-----return code.

-----See the description of File Descriptor (FD) output values
-----in the OUTPUT section for more information on file
-----descriptor names.

-----+D_ causes ls of to search for all open instances of directory
-----D_ and all the files and directories it contains to its
-----complete depth.

-----Processing of the +D_ option does not follow symbolic links
-----within D_ unless the -x_ or -x_1_ option is also specified.
-----Nor does it search for open files on file system mount
-----points on subdirectories of D_ unless the -x_ or -x_1_ f
-----option is also specified.

-----Note: the authority of the user of this option limits it
-----to searching for files that the user has permission to
-----examine with the system stat(2) function.

-----Further note: ls of may process this option slowly and
-----require a large amount of dynamic memory to do it. This
-----is because it must descend the entire directory tree,
-----rooted at D_, calling stat(2) for each file and directory,
-----building a list of all the files it finds, and searching
-----that list for a match with every open file. When
-----directory D_ is large, these steps can take a long time, so
-----use this option prudently.

-----D_ directs ls of's use of the device cache file. The use of
this option is sometimes restricted. See the DEVICE CACHE
FILE section and the sections that follow it for more
information on this option.

-D must be followed by a function letter; the function
letter may optionally be followed by a path name. Ls of
recognizes these function letters:

    ? - report device cache file paths
    b - build the device cache file
    i - ignore the device cache file
    r - read the device cache file
    u - read and update the device cache file

The b, r, and u functions, accompanied by a path name, are
sometimes restricted. When these functions are
restricted, they will not appear in the description of the
-D option that accompanies -h or -? option output. See
the DEVICE CACHE FILE section and the sections that follow
it for more information on these functions and when
they're restricted.

-----The -?_ function reports the read-only and write paths that
-----ls of can use for the device cache file, the names of any
-----environment variables whose values ls of will examine when
-----forming the device cache file path, and the format for the
-----personal device cache file path. (Escape the '?'
character as your shell requires.)

When available, the b, r, and u functions may be followed
by the device cache file's path. The standard default is
ls of hostname in the home directory of the real user ID
that executes ls of, but this could have been changed when
ls of was configured and compiled. (The output of the -h
and -? options show the current default prefix--e.g.,
'. ls of'.) The suffix, hostname, is the first component
of the host's name returned by gethostname(2).

When available, the b function directs ls of to build a new
device cache file at the default or specified path.

The i function directs ls of to ignore the default device
cache file and obtain its information about devices via
direct calls to the kernel.

The r function directs ls of to read the device cache at
the default or specified path, but prevents it from
creating a new device cache file when none exists or the
existing one is improperly structured. The r function,
when specified without a path name, prevents ls of from
updating an incorrect or outdated device cache file, or
creating a new one in its place. The r function is always
available when it is specified without a path name
argument; it may be restricted by the permissions of the
ls of process.

When available, the u function directs ls of to read the
device cache file at the default or specified path, if
possible, and to rebuild it, if necessary. This is the
default device cache file function when no -D option has

```

been specified.

+|-e s exempts the file system whose path name is s from being subjected to kernel **function** calls that might block. The +e option exempts stat(2), lstat(2) and most readlink(2) kernel **function** calls. The -e option exempts only stat(2) and lstat(2) kernel **function** calls. Multiple file systems may be specified with separate +|-e specifications and each may have readlink(2) calls exempted or not.

This option is currently implemented only for Linux.

CAUTION: this option can easily be mis-applied to other than the file system of interest, because it uses path name rather than the more reliable device and inode numbers. (Device and inode numbers are acquired via the potentially blocking stat(2) kernel call and are thus not available, but see the +|-m m option as a possible alternative way to supply device numbers.) Use this option with great care and fully specify the path name of the file system to be exempted.

When open files on exempted file systems are reported, it may not be possible to obtain all their information. Therefore, some information columns will be blank, the characters 'UNKN' preface the values in the TYPE column, and the applicable exemption option is added in parentheses to the end of the NAME column. (Some device number information might be made available via the +|-m m option.)

+|-E +E specifies that Linux pipe, Linux UNIX socket and Linux pseudoterminal files should be displayed with endpoint information and the files of the endpoints should also be displayed. Note: UNIX socket file endpoint information is only available when the compile flags line of -v output contains HASUXSOCKET, and pseudoterminal endpoint information is only available when the compile flags line contains HASPTYEPT.

Pipe endpoint information is displayed in the NAME column in the form 'PID,cmd,FDmode', where PID is the endpoint process ID; cmd is the endpoint process **command**; FD is the endpoint file's descriptor; and mode is the endpoint file's access mode.

Pseudoterminal endpoint information is displayed in the NAME column as "->/dev/ptsmin PID,cmd,FDmode" or "PID,cmd,FDmode". The first form is for a master device; the second, for a slave device. min is a slave device's minor device number; and PID, cmd, FD and mode are the same as with pipe endpoint information. Note: pseudoterminal endpoint information is only available when the compile flags line of -V output contains HASPTYEPT.

UNIX socket file endpoint information is displayed in the NAME column in the form "type=TYPE->INO=INODE-PID,cmd,FDmode", where TYPE is the socket type; INODE is the inode number of the connected socket; and PID, cmd, FD and mode are the same as with pipe endpoint information. Note: UNIX socket file endpoint information is available only when the compile flags line of -v output contains HASUXSOCKET.

Multiple occurrences of this information can appear in a file's NAME column.

-E specifies that Linux pipe and Linux UNIX socket files should be displayed with endpoint information, but not the files of the endpoints.

+|-f [cfgGn] f by itself clarifies how path name arguments are to be interpreted. When followed by c, f, g, G, or n in any combination it specifies that the listing of kernel file structure information is to be enabled ('+') or inhibited ('-').

Normally a path name argument is taken to be a file system name if it matches a mounted-on directory name reported by mount(8), or if it represents a block device, named in the mount output and associated with a mounted directory name. When +f is specified, all path name arguments will be taken to be file system names, and lsof will complain if any are not. This can be useful, for example, when the file system name (mounted-on device) isn't a block device. This happens for some CD-ROM file systems.

When -f is specified by itself, all path name arguments will be taken to be simple files. Thus, for example, the "-f -/" arguments direct lsof to search for open files with a "/" path name, not all open files in the "/" (root)

```

-----file _system.

-----Be careful to make sure +f and -f are properly terminated
-----and aren't followed by a character (e.g., of the file or
-----file system name) that might be taken as a parameter. For
-----example, use '---' after +f and -f as in these examples.

$ lsuf +f --- /file/system/name
$ lsuf -f --- /file/name

The listing of information from kernel file structures,
requested with the +f [cfgGn] option form, is normally
inhibited, and is not available in whole or part for some
dialects - e.g., /proc-based Linux kernels below 2.6.22.
When the prefix to f is a plus sign ('+'), these
-----characters request file structure information:

-----c----file _structure _use _count_(not _Linux)
-----f----file _structure _address_(not _Linux)
-----g----file _flag _abbreviations_(Linux_2.6.22_and_up)
-----G----file _flags _in _hexadecimal_(Linux_2.6.22_and_up)
-----n----file _structure _node _address_(not _Linux)

-----When the prefix is minus('-') the same characters disable
-----the listing of the indicated values.

File structure addresses, use counts, flags, and node
addresses may be used to detect more readily identical
files inherited by child processes and identical files in
use by different processes. Lsuf column output can be
sorted by output columns holding the values and listed to
identify identical file use, or lsuf field output can be
parsed by an AWK or Perl post-filter script, or by a C
program.

-F f specifies a character list, f, that selects the fields to
be output for processing by another program, and the
character that terminates each output field. Each field
to be output is specified with a single character in f.
The field terminator defaults to NL, but may be changed to
NUL (000). See the OUTPUT FOR OTHER PROGRAMS section for
a description of the field identification characters and
the field output process.

When the field selection character list is empty, all
standard fields are selected (except the raw device field,
security context and zone field for compatibility reasons)
and the NL field terminator is used.

When the field selection character list contains only a
zero ('0'), all fields are selected (except the raw device
-----field for compatibility reasons) and the NUL terminator
-----character is used.

-----Other combinations of fields and their associated field
-----terminator character must be set with explicit entries in
-----f, as described in the OUTPUT FOR OTHER PROGRAMS section.

-----When a field selection character identifies an item lsuf
-----does not normally list - e.g., PPID, selected with -R-
-----specification of the field character - e.g., '-FR' -
-----also selects the listing of the item.

-----When the field selection character list contains the
-----single character '?', lsuf will display a help list of the
-----field identification characters. (Escape the '?'
-----character as your shell requires.)

-----g-[s] excludes or selects the listing of files for the processes
-----whose optional process group identification (PGID) numbers
-----are in the comma-separated set s - e.g., '-123' or
-----'123,456' - (There should be no spaces in the set.)

-----PGID numbers that begin with '^' (negation) represent
-----exclusions.

Multiple PGID numbers are joined in a single ORed set
before participating in AND option selection. However,
PGID exclusions are applied without ORing or ANDing and
take effect before other selection criteria are applied.

The -g option also enables the output display of PGID
numbers. When specified without a PGID set that's all it
-----does.

-----i-[i] selects the listing of files any of whose Internet address
-----matches the address specified in i. If no address is
-----specified, this option selects the listing of all Internet
-----and x.25 (HP-UX) network files.

-----If -i4 or -i6 is specified with no following address, only
-----files of the indicated IP version, IPv4 or IPv6, are

```

```

-----displayed. (An IPv6 specification may be used only if the
-----dialects supports IPv6, as indicated by '[46]' and
-----'IPv[46]' in lsof's -h or -? output.) Sequentially
specifying -i4, followed by -i6 is the same as specifying
-i, and vice-versa. Specifying -i4, or -i6 after -i is
the same as specifying -i4 or -i6 by itself.

Multiple addresses (up to a limit of 100) may be specified
with multiple -i options. (A port number or service name
range is counted as one address.) They are joined in a
single ORed set before participating in AND option
selection.

An Internet address is specified in the form (Items in
square brackets are optional.):

[46][protocol][@hostname|hostaddr][:service|port]

where:
    46 specifies the IP version, IPv4 or IPv6
        that applies to the following address.
        '6' may be specified only if the UNIX
        dialect supports IPv6. If neither '4' nor
        '6' is specified, the following address
        applies to all IP versions.
    protocol is a protocol name - TCP, UDP
    hostname is an Internet host name. Unless a
        specific IP version is specified, open
        network files associated with host names
        of all versions will be selected.
    hostaddr is a numeric Internet IPv4 address in
        dot form; or an IPv6 numeric address in
        colon form, enclosed in brackets, if the
        UNIX dialect supports IPv6. When an IP
        version is selected, only its numeric
        addresses may be specified.
    service is an /etc/services name - e.g., smtp -
        or a list of them.
    port is a port number, or a list of them.

IPv6 options may be used only if the UNIX dialect supports
IPv6. To see if the dialect supports IPv6, run lsof and
specify the -h or -? (help) option. If the displayed
description of the -i option contains '[46]' and
'IPv[46]', IPv6 is supported.

IPv4 host names and addresses may not be specified if
network file selection is limited to IPv6 with -i 6. IPv6
host names and addresses may not be specified if network
file selection is limited to IPv4 with -i 4. When an open
IPv4 network file's address is mapped in an IPv6 address,
the open file's type will be IPv6, not IPv4, and its
display will be selected by '6', not '4'.

At least one address component - 4, 6, protocol, hostname,
hostaddr, or service - must be supplied. The '@'
character, leading the host specification, is always
required, as is the ':', leading the port specification.
Specify either hostname or hostaddr. Specify either
service name list or port number list. If a service name
list is specified, the protocol may also need to be
specified if the TCP, UDP and UDPLITE port numbers for the
service name are different. Use any case - lower or upper
- for protocol.

Service names and port numbers may be combined in a list
whose entries are separated by commas and whose numeric
range entries are separated by minus signs. There may be
no embedded spaces, and all service names must belong to
the specified protocol. Since service names may contain
embedded minus signs, the starting entry of a range can't
be a service name; it can be a port number, however.

-----Here are some sample addresses:

-----i6 - IPv6 only
-----TCP:25 - TCP and port 25
-----@1.2.3.4 - Internet IPv4 host address 1.2.3.4
-----@[3ffe:1ebc::1]:1234 - Internet IPv6 host address
-----[3ffe:1ebc::1], port 1234
-----UDP:who - UDP who service port
-----TCP@lsof.itap:513 - TCP, port 513 and host name lsof.itap
-----tcp@foo:1-10,smtp,99 - TCP, ports 1 through 10,
-----service_name smtp, port 99, host_name foo
-----tcp@bar:1-smtp - TCP, ports 1 through smtp, host bar
-----:time - either TCP, UDP or UDPLITE time service port

-----Kk - selects the listing of tasks (threads) of processes, on
-----dialects where task (thread) reporting is supported. (If
-----help output - i.e., the output of the -h or -? options -
-----shows this option, then task (thread) reporting is
-----supported by the dialect.)

```

```

-----If -K is followed by a value, -k, it must be 'i'. That
-----causes lsof to ignore tasks, particularly in the default,
-----list everything case when no other options are specified.

-----When -K and -a are both specified on Linux, and the tasks
-----of a main process are selected by other options, the main
-----process will also be listed as though it were a task, but
-----without a task ID. (See the description of the TID column
-----in the OUTPUT section.)

-----Where the FreeBSD version supports threads, all threads
-----will be listed with their IDs.

-----In general threads and tasks inherit the files of the
-----caller, but may close some and open others, so lsof always
-----reports all the open files of threads and tasks.

-----k_k specifies a kernel name list file, -k, in place of /vmunix,
-----/mach, etc. -k is not available under AIX on the IBM
-----RISC/System 6000.

-----l inhibits the conversion of user ID numbers to login names.
-----It is also useful when login name lookup is working
-----improperly or slowly.

-----+|-L[1]
-----enables ('+') or disables ('-') the listing of file link
-----counts, where they are available -- e.g., they aren't
-----available for sockets, or most FIFOs and pipes.

When +L is specified without a following number, all link
counts will be listed. When -L is specified (the
default), no link counts will be listed.

When +L is followed by a number, only files having a link
count less than that number will be listed. (No number
may follow -L.) A specification of the form '+L1' will
select open files that have been unlinked. A
specification of the form '+aL1 <file-system>' will
select unlinked open files on the specified file system.

For other link count comparisons, use field output (-F)
and a post-processing script or program.

+|-m m specifies an alternate kernel memory file or activates
mount table supplement processing.

The option form -m m specifies a kernel memory file, m, in
place of /dev/kmem or /dev/mem -- e.g., a crash dump file.

The option form +m requests that a mount supplement file
be written to the standard output file. All other options
are silently ignored.

There will be a line in the mount supplement file for each
mounted file system, containing the mounted file system
directory, followed by a single space, followed by the
device number in hexadecimal "0x" format -- e.g.,

    / 0x801

Lsof can use the mount supplement file to get device
numbers for file systems when it can't get them via
stat(2) or lstat(2).

-----The option form +m m identifies m as a mount supplement
-----file.

-----Note: the +m and -m m options are not available for all
-----supported dialects. Check the output of lsof's -h or -?
-----options to see if the +m and -m m options are available.

+|-M Enables (+) or disables (-) the reporting of portmapper
registrations for local TCP, UDP and UDPLITE ports, where
port mapping is supported. (See the last paragraph of
this option description for information about where
portmapper registration reporting is supported.)

The default reporting mode is set by the lsof builder with
the HASPMAPENABLED #define in the dialect's machine.h
header file; lsof is distributed with the HASPMAPENABLED
#define deactivated, so portmapper reporting is disabled
by default and must be requested with +M. Specifying
lsof's -h or -? option will report the default mode.
Disabling portmapper registration when it is already
disabled or enabling it when already enabled is
acceptable. When portmapper registration reporting is
enabled, lsof displays the portmapper registration (if
any) for local TCP, UDP or UDPLITE ports in square
brackets immediately following the port numbers or service
names -- e.g., ':1234[name]' or ':name[100083]'. The

```

```

registration information may be a name or number,
depending on what the registering program supplied to the
portmapper when it registered the port.

When portmapper registration reporting is enabled, lsof
may run a little more slowly or even become blocked when
access to the portmapper becomes congested or stopped.
Reverse the reporting mode to determine if portmapper
registration reporting is slowing or blocking lsof.

For purposes of portmapper registration reporting lsof
considers a TCP, UDP or UDPLITE port local if: it is found
in the local part of its containing kernel structure; or
if it is located in the foreign part of its containing
kernel structure and the local and foreign Internet
addresses are the same; or if it is located in the foreign
part of its containing kernel structure and the foreign
Internet address is INADDR_LOOPBACK (127.0.0.1). This
rule may make lsof ignore some foreign ports on machines
with multiple interfaces when the foreign Internet address
is on a different interface from the local one.

See the lsof FAQ (The FAQ section gives its location.)
for further discussion of portmapper registration
reporting issues.

Portmapper registration reporting is supported only on
dialects that have RPC header files. (Some Linux
distributions with Glibc 2.14 do not have them.) When
portmapper registration reporting is supported, the -h or
--help output will show the +|-M option.

-n inhibits the conversion of network numbers to host names
for network files. Inhibiting conversion may make lsof
run faster. It is also useful when host name lookup is
not working properly.

-N selects the listing of NFS files.

-o directs lsof to display file offset at all times. It
causes the SIZE/OFF output column title to be changed to
OFFSET. Note: on some UNIX dialects lsof can't obtain
accurate or consistent file offset information from its
kernel data sources, sometimes just for particular kinds
of files (e.g., socket files.) Consult the lsof FAQ (The
FAQ section gives its location.) for more information.

The -o and -s options are mutually exclusive; they can't
both be specified. When neither is specified, lsof
displays whatever value size or offset is appropriate
and available for the type of the file.

-o-o defines the number of decimal digits (o) to be printed
after the '0t' for a file offset before the form is
switched to '0x...' . An -o value of zero (unlimited)
directs lsof to use the '0t' form for all offset output.

This option does NOT direct lsof to display offset at all
times; specify -o- (without a trailing number) to do that.
-o-o only specifies the number of digits after '0t' in
either mixed size and offset or offset only output. Thus,
for example, to direct lsof to display offset at all times
with a decimal digit count of 10, use:

-o-o-10
or
-o10

The default number of digits allowed after '0t' is
normally 8, but may have been changed by the lsof builder.
Consult the description of the -o-o option in the output
of the -h or --? option to determine the default that is
in effect.

-O directs lsof to bypass the strategy it uses to avoid being
blocked by some kernel operations - i.e., doing them in
forked child processes. See the BLOCKS_AND_TIMEOUTS and
AVOIDING_KERNEL_BLOCKS sections for more information on
kernel operations that may block lsof.

While use of this option will reduce lsof startup
overhead, it may also cause lsof to hang when the kernel
doesn't respond to a function. Use this option
cautiously.

-p s
excludes or selects the listing of files for the processes
whose optional process IDentification (PID) numbers are in
the comma-separated set s - e.g., '123' or '123,456'.
(There should be no spaces in the set.)

PID numbers that begin with '-' (negation) represent
exclusions.

```

```

-----Multiple process ID numbers are joined in a single ORed
-----set before participating in AND option selection.
-----However, PID exclusions are applied without ORing or
-----ANDing and take effect before other selection criteria are
-----applied.

-----P-----inhibits the conversion of port numbers to port names for
-----network files. Inhibiting the conversion may make lsof
-----run a little faster. It is also useful when port name
-----lookup is not working properly.

-----+|-r[<fmt>]
-----puts lsof in repeat mode. There lsof lists open files as
-----selected by other options, delays t seconds (default
-----fifteen), then repeats the listing, delaying and listing
-----repetitively until stopped by a condition defined by the
-----prefix to the option.

-----If the prefix is a '-', repeat mode is endless. Lsof must
-----be terminated with an interrupt or quit signal.

-----If the prefix is '+', repeat mode will end the first cycle
-----no open files are listed and of course when lsof is
-----stopped with an interrupt or quit signal. When repeat
-----mode ends because no files are listed, the process exit
-----code will be zero if any open files were ever listed; one,
-----if none were ever listed.

-----Lsof marks the end of each listing: if field output is in
-----progress (the -F option has been specified), the default
-----marker is 'm'; otherwise the default marker is
-----"====='. The marker is followed by a NL character.

-----The optional "m<fmt>" argument specifies a format for the
-----marker line. The <fmt> characters following 'm' are
-----interpreted as a format specification to the strftime(3)
-----function, when both it and the localtime(3) function are
-----available in the dialect's C library. Consult the
-----strftime(3) documentation for what may appear in its
-----format specification. Note that when field output is
-----requested with the -F option, <fmt> cannot contain the NL
-----format, '%n'. Note also that when <fmt> contains spaces
-----or other characters that affect the shell's interpretation
-----of arguments, <fmt> must be quoted appropriately.

-----Repeat mode reduces lsof startup overhead, so it is more
-----efficient to use this mode than to call lsof repetitively
-----from a shell script, for example.

-----To use repeat mode most efficiently, accompany +|-r with
-----specification of other lsof selection options, so the
-----amount of kernel memory access lsof does will be kept to a
-----minimum. Options that filter at the process level, e.g.,
-----c, -g, -p, -u are the most efficient selectors.

-----Repeat mode is useful when coupled with field output (see
-----the -F option description) and a supervising awk or Perl
-----script, or a C program.

-----R-----directs lsof to list the Parent Process Identification
-----number in the PPID column.

-----s-[p:s]
-----s alone directs lsof to display file size at all times.
-----It causes the SIZE/OFF output column title to be changed
-----to SIZE. If the file does not have a size, nothing is
-----displayed.

-----The optional -s-p:s form is available only for selected
-----dialects, and only when the -h or -? help output lists
-----it.

-----When the optional form is available, the s may be followed
-----by a protocol name (p), either TCP or UDP, a colon (':')
-----and a comma-separated protocol state name list; the option
-----causes open TCP and UDP files to be excluded if their
-----state name(s) are in the list (s) preceded by a '^'; or
-----included if their name(s) are not preceded by a '^'.

-----Dialects that support this option may support only one
-----protocol. When an unsupported protocol is specified, a
-----message will be displayed indicating state names for the
-----protocol are unavailable.

-----When an inclusion list is defined, only network files with
-----state names in the list will be present in the lsof
-----output. Thus, specifying one state name means that only
-----network files with that lone state name will be listed.

-----Case is unimportant in the protocol or state names, but
-----there may be no spaces and the colon (':') separating the

```

```

-----protocol_name_(p)_and_the_state_name_list_(s)_is_required.

-----If_only_TCP_and_UDP_files_are_to_be_listed,_as_controlled
-----by_the_specified_exclusions_and_inclusions,_the_-i_option
-----must_be_specified,_too. If only a single protocol's files
are to be listed, add its name as an argument to the -i
option.

For example, to list only network files with TCP state
LISTEN, use:

-iTCP -sTCP:LISTEN

Or, for example, to list network files with all UDP states
except Idle, use:

-iUDP -sUDP:Idle

State names vary with UNIX dialects, so it's not possible
to provide a complete list. Some common TCP state names
are: _CLOSED, _IDLE, _BOUND, _LISTEN, _ESTABLISHED, _SYN_SENT,
_SYN_RCDV, _ESTABLISHED, _CLOSE_WAIT, _FIN_WAIT1, _CLOSING,
_LAST_ACK, _FIN_WAIT_2, and _TIME_WAIT. Two common UDP state
names are _Unbound and _Idle.

-----See_the_1s_of_FAQ_(The_FAQ_section_gives_its_location.)
-----for_more_information_on_how_to_use_protocol_state
-----exclusion_and_inclusion,_including_examples.

-----The_-o_(without_a_following_decimal_digit_count)_and_-s
-----option_(without_a_following_protocol_and_state_name_list)
-----are_mutually_exclusive;_they_can't_both_be_specified.
When neither is specified, 1s of displays whatever value -
size or offset - is appropriate and available for the type
of file.

Since some types of files don't have true sizes -- sockets,
FIFOs, pipes, etc. -- 1s of displays for their sizes the
content amounts in their associated kernel buffers, if
possible.

-----S_[t]_specifies_an_optional_time-out_seconds_value_for_kernel
-----functions -- 1stat(2), readlink(2), and stat(2) -- that
-----might_otherwise_deadlock. The minimum for -t is two; the
-----default is fifteen; when no value is specified, the default
-----is used.

-----See_the_BLOCKS_AND_TIMEOUTS_section_for_more_information.

-----T_[t]_controls_the_reporting_of_some_TCP/TPI_information,_also
-----reported_by_netstat(1),_following_the_network_addresses.
-----In_normal_output_the_information_appears_in_parentheses,
-----each_item_except_TCP_or_TPI_state_name_identified_by_a
-----keyword,_followed_by_=' ', separated from others by a
single space:

    <TCP or TPI state name>
    QR=<read queue length>
    QS=<send queue length>
    SO=<socket options and values>
    SS=<socket states>
    TF=<TCP flags and values>
    WR=<window read length>
    WW=<window write length>

Not all values are reported for all UNIX dialects. Items
values (when available) are reported after the item name
and '='.

When the field output mode is in effect (See OUTPUT FOR
OTHER PROGRAMS.) each item appears as a field with a 'T'
leading character.

-----T_with_no_following_key_characters_disables_TCP/TPI
-----information_reporting.

-----T_with_following_characters_selects_the_reporting_of
-----specific_TCP/TPI_information:

-----f_____selects_reporting_of_socket_options,
-----states_and_values,_and_TCP_flags_and
-----values.
-----q_____selects_queue_length_reporting.
-----s_____selects_connection_state_reporting.
-----w_____selects_window_size_reporting.

-----Not_all_selections_are_enabled_for_some_UNIX_dialects.
-----State_may_be_selected_for_all_dialects_and_is_reported_by
-----default. The -h or -? help output for the -T option will
-----show_what_selections_may_be_used_with_the_UNIX_dialect.

-----When_-T_is_used_to_select_information--i.e.,_it_is

```



```

-----followed by one or more selection characters--the
-----displaying of state is disabled by default, and it must be
-----explicitly selected again in the characters following -T.
----- (In effect, then, the default is equivalent to -Ts.) --For
-----example, if queue lengths and state are desired, use -Tqs.

-----Socket options, socket states, some socket values, TCP
-----flags and one TCP value may be reported (when available in
-----the UNIX dialect) in the form of the names that commonly
-----appear after SO, so, SS, TCP and TF in the dialect's
-----header files -- most often <sys/socket.h>,
-----<sys/socketvar.h> and <netinet/tcp_var.h>. Consult those
-----header files for the meaning of the flags, options, states
-----and values.

'SO=' precedes socket options and values; 'SS=',
socket states; and 'TF=', TCP flags and values.

If a flag or option has a value, the value will follow an
 '=' and the name -- e.g., 'SO=LINGER=5', 'SO=QLIM=5',
'TF=MSS=512'. The following seven values may be
reported:

Name
Reported Description (Common Symbol)
KEEPALIVE keep alive time (SO_KEEPAIVE)
LINGER linger time (SO_LINGER)
MSS maximum segment size (TCP_MAXSEG)
PQLEN partial listen queue connections
QLEN established listen queue connections
QLIM established listen queue limit
RCVBUF receive buffer length (SO_RCVBUF)
SNDBUF send buffer length (SO_SNDBUF)

Details on what socket options and values, socket states,
and TCP flags and values may be displayed for particular
UNIX dialects may be found in the answer to the 'Why
doesn't lsof report socket options, socket states, and TCP
-----flags and values for my dialect?' and 'Why doesn't lsof
report the partial listen queue connection count for my
dialect?' questions in the lsof FAQ (The FAQ section
gives its location.)

-t specifies that lsof should produce terse output with
process identifiers only and no header -- e.g., so that the
output may be piped to kill(1). -t selects the -w option.

-u s selects the listing of files for the user whose login
names or user ID numbers are in the comma-separated set s
-- e.g., 'abe', or '548,root'. (There should be no
spaces in the set.)

Multiple login names or user ID numbers are joined in a
single ORed set before participating in AND option
selection.

If a login name or user ID is preceded by a '^', it
-----becomes a negation--i.e., files of processes owned by the
-----login name or user ID will never be listed. --A negated
-----login name or user ID selection is neither ANDed nor ORed
-----with other selections; it is applied before all other
-----selections and absolutely excludes the listing of the
-----files of the process. --For example, to direct lsof to
-----exclude the listing of files belonging to root processes,
-----specify '-u^root' or '-u^0'.

-----U-----selects the listing of UNIX-domain socket files.

-----V-----selects the listing of lsof version information,
-----including: revision number; when the lsof binary was
-----constructed; who constructed the binary and where; the
-----name of the compiler used to construct the lsof binary;
-----the version number of the compiler when readily available;
-----the compiler and loader flags used to construct the lsof
-----binary; and system information, typically the output of
-----uname's -a option.

-V directs lsof to indicate the items it was asked to list
and failed to find -- command names, file names, Internet
addresses or files, login names, NFS files, PIDs, PGIDs,
and UIDs.

When other options are ANDed to search options, or
compile-time options restrict the listing of some files,
lsof may not report that it failed to find a search item
when an ANDed option or compile-time option prevents the
listing of the open file containing the located search
item.

For example, 'lsof -V -iTCP@foobar -a -d 999' may not
report a failure to locate open files at 'TCP@foobar'

```

and may not list any, if none have a file descriptor number of 999. A similar situation arises when HASSECURITY and HASNOSOCKSECURITY are defined at compile time and they prevent the listing of open files.

+|-w Enables (+) or disables (-) the suppression of warning messages.

The lsof builder may choose to have warning messages disabled or enabled by default. The default warning message state is indicated in the output of the -h or -? option. Disabling warning messages when they are already disabled or enabling them when already enabled is acceptable.

The -t option selects the -w option.

-x [fl] may accompany the +d and +D options to direct their processing to cross over symbolic links and/or file system mount points encountered when scanning the directory (+d) or directory tree (+D).

If -x is specified by itself without a following parameter, cross-over processing of both symbolic links and file system mount points is enabled. Note that when -x is specified without a parameter, the next argument must begin with '-' or '+'. The optional 'f' parameter enables file system mount point cross-over processing; 'l', symbolic link cross-over processing.

The -x option may not be supplied without also supplying a +d or +D option.

-X This is a dialect-specific option.

AIX: This IBM AIX RISC/System 6000 option requests the reporting of executed text file and shared library references.

WARNING: because this option uses the kernel readx() function, its use on a busy AIX system might cause an application process to hang so completely that it can neither be killed nor stopped. I have never seen this happen or had a report of its happening, but I think there is a remote possibility it could happen.

By default use of readx() is disabled. On AIX 5L and above lsof may need setuid-root permission to perform the actions this option requests.

The lsof builder may specify that the -X option be restricted to processes whose real UID is root. If that has been **done**, the -X option will not appear in the -h or -? **help** output unless the real UID of the lsof process is root. The default lsof distribution allows any UID to specify -X, so by default it will appear in the **help** output.

When AIX readx() use is disabled, lsof may not be able to report information for all text and loader file references, but it may also avoid exacerbating an AIX kernel directory search kernel error, known as the Stale Segment ID bug.

The readx() **function**, used by lsof or any other program to access some sections of kernel virtual memory, can trigger the Stale Segment ID bug. It can cause the kernel's _dir_search() function to believe erroneously that part of an in-memory copy of a file system directory has been zeroed. Another application process, distinct from lsof, asking the kernel to search the directory - e.g., by using open(2) - can cause _dir_search() to loop forever, thus hanging the application process.

----- Consult the lsof FAQ (The FAQ section gives its location) and the 00README file of the lsof distribution for a more complete description of the Stale Segment ID bug, its APAR, and methods for defining readx() use when compiling lsof.

-----Linux: This Linux option requests that lsof skip the reporting of information on all open TCP, UDP and UDPLITE IPv4 and IPv6 files.

----- This Linux option is most useful when the system has an extremely large number of open TCP, UDP and UDPLITE

-----files, the processing of whose information in the
 -----/proc/net/tcp* and /proc/net/udp* files would take lsof
 -----a long time, and whose reporting is not of interest.

-----Use this option with care and only when you are sure
 -----that the information you want lsof to display isn't
 -----associated with open TCP, UDP or UDPLITE socket files.

Solaris 10 and above:
 This Solaris 10 and above option requests the reporting
 of cached paths for files that have been deleted — i.e.,
 removed with rm(1) or unlink(2).

The cached path is followed by the string “ (deleted)”
 to indicate that the path by which the file was opened
 has been deleted.

Because intervening changes made to the path — i.e.,
 renames with mv(1) or rename(2) — are not recorded in
 the cached path, what lsof reports is only the path by
 which the file was opened, not its possibly different
 final path.

—z [z] specifies how Solaris 10 and higher zone information is
 to be handled.

Without a following argument — e.g., NO z — the option
 specifies that zone names are to be listed in the ZONE
 output column.

The —z option may be followed by a zone name, z. That
 causes lsof to list only open files for processes in
 that zone. Multiple —z z option and argument pairs may
 be specified to form a list of named zones. Any open
 file of any process in any of the zones will be listed,
 subject to other conditions specified by other options
 and arguments.

—Z [Z] specifies how SELinux security contexts are to be
 handled. It and 'Z' field output character support are
 inhibited when SELinux is disabled in the running Linux
 kernel. See OUTPUT FOR OTHER PROGRAMS for more
 information on the 'Z' field output character.

Without a following argument — e.g., NO Z — the option
 specifies that security contexts are to be listed in the
 SECURITY-CONTEXT output column.

The —Z option may be followed by a wildcard security
 context name, Z. That causes lsof to list only open
 files for processes in that security context. Multiple
 —Z Z option and argument pairs may be specified to form
 a list of security contexts. Any open file of any
 process in any of the security contexts will be listed,
 subject to other conditions specified by other options
 and arguments. Note that Z can be A:B:C or *:B:C or
 A:B:* or *:*:C to match against the A:B:C context.

— The double minus sign option is a marker that signals
 the end of the keyed options. It may be used, for
 example, when the first file name begins with a minus
 sign. It may also be used when the absence of a value
 for the last keyed option must be signified by the
 presence of a minus sign in the following option and
 before the start of the file names.

names These are path names of specific files to list.
 Symbolic links are resolved before use. The first name
 may be separated from the preceding options with the
 “—” option.

If a name is the mounted-on directory of a file system
 or the device of the file system, lsof will list all the
 files open on the file system. To be considered a file
 system, the name must match a mounted-on directory name
 in mount(8) output, or match the name of a block device
 associated with a mounted-on directory name. The +|-f
 option may be used to force lsof to consider a name a
 file system identifier (+f) or a simple file (—f).

If name is a path to a directory that is not the
 mounted-on directory name of a file system, it is
 treated just as a regular file is treated — i.e., its
 listing is restricted to processes that have it open as
 a file or as a process-specific directory, such as the
 root or current working directory. To request that lsof
 look for open files inside a directory name, use the +d
 s and +D D options.

If a name is the base name of a family of multiplexed
 files — e.g., AIX's /dev/pt[cs] — lsof will list all the
 -----associated multiplexed files on the device that are open

```

-----e.g., /dev/pt[cs]/1, /dev/pt[cs]/2, etc.

-----If a_name is a_UNIX_domain_socket_name, _ls of will
-----usually search for it by the characters of the_name
-----alone--exactly as it is specified and is recorded in
-----the_kernel_socket_structure. (See the next paragraph
-----for an exception to that rule for Linux.) --Specifying a
-----relative_path--e.g., ./file--in place of the_file's
-----absolute_path--e.g., /tmp/file--won't work because
-----_ls of must match the characters you specify with _what_it
-----finds in the_kernel_UNIX_domain_socket_structures.

-----If a_name is a_Linux_UNIX_domain_socket_name, in one
-----case _ls of is able to search for it by its_device_and
-----inode_number, allowing _name to be a relative_path. --The
-----case requires that the_absolute_path--i.e., one
-----beginning with a_slash_('/') be used by the_process that
-----created the_socket, and hence be stored in the
-----/proc/net/unix_file; and it requires that _ls of be able
-----to obtain the_device_and_node_numbers_of_both_the
-----absolute_path_in_/proc/net/unix_and_name_via_successful
-----stat(2)_system_calls. --When those conditions are met,
-----_ls of will be able to search for the_UNIX_domain_socket
-----when some_path_to_it is specified in_name. --Thus, for
-----example, if the_path is /dev/log, and an _ls of search is
-----initiated when the_working_directory is /dev, then_name
-----could be ./log.

-----If a_name is none_of_the_above, _ls of will list any_open
-----files_whose_device_and_inode_match_that_of_the_specified
-----path_name.

-----If you have also specified the_-b_option, the_only_names
-----you may safely specify are_file_systems_for_which_your
-----mount_table_supplies_alternate_device_numbers. --See the
-----AVOIDING_KERNEL_BLOCKS_and_ALTERNATE_DEVICE_NUMBERS
-----sections for more information.

-----Multiple_file_names_are_joined_in_a_single_ORed_set
-----before_participating_in_AND_option_selection.
AFS
-----_Ls of supports the_recognition_of_AFS_files_for_these_dialects
----- (and_AFS_versions):
-----
-----AIX_4.1.4_(AFS_3.4a)
-----HP-UX_9.0.5_(AFS_3.4a)
-----Linux_1.2.13_(AFS_3.3)
-----Solaris_2.[56]_(AFS_3.4a)

-----It may recognize_AFS_files_on_other_versions_of_these_dialects,
-----but has not been tested there. --Depending on how_AFS_is
-----implemented, _ls of may recognize_AFS_files_in_other_dialects, or
-----may have difficulties recognizing_AFS_files_in_the_supported
-----dialects.

-----_Ls of may have trouble identifying all_aspects_of_AFS_files_in
-----supported_dialects_when_AFS_kernel_support_is_implemented_via
-----dynamic_modules_whose_addresses_do_not_appear_in_the_kernel's
-----variable_name_list. In that case, _ls of may have to guess at the
-----identity of_AFS_files, and might not be able to obtain volume
-----information from the kernel that is needed for calculating AFS
-----volume node numbers. When _ls of can't compute volume_node
-----numbers, it_reports_blank_in_the_NODE_column.

-----The_-A_option is available in some_dialect_implementations_of
-----_ls of_for_specifying_the_name_list_file_where_dynamic_module
-----kernel_addresses_may_be_found. --When this_option_is_available, it
-----will be listed in the _ls of_help_output, presented in response to
-----the_-h_or_-.

-----See the _ls of_FAQ_(The_FAQ_section_gives_its_location.) --for more
-----information_about_dynamic_modules, their_symbols, and how they
-----affect _ls of_options.

-----Because_AFS_path_lookups_don't seem to participate in the
-----kernel's_name_cache_operations, _ls of can't identify path name
-----components for AFS files.
SECURITY
-----_Ls of has three features that may cause security concerns. First,
-----its default compilation mode allows anyone to list all open files
-----with it. Second, by default it creates a user-readable and
-----user-writable device cache file in the home directory of the real
-----user ID that executes _ls of. (The list_all_open_files_and_device
-----cache features may be disabled when _ls of is compiled.) Third,
-----its_-k_and_-m_options name alternate kernel name list or memory
-----files.

-----Restricting the listing of all open files is controlled by the
-----compile-time HASSECURITY and HASNOSOCKSECURITY options. When
-----HASSECURITY is defined, _ls of will allow only the root user to
-----list all open files. The non-root user may list only open files
-----of processes with the same user IDENTIFICATION number as the real

```

user ID number of the lsOF process (the one that its user logged on with).

However, if HASSECURITY and HASNOSOCKSECURITY are both defined, anyone may list open socket files, provided they are selected with the -i option.

When HASSECURITY is not defined, anyone may list all open files.

Help output, presented in response to the -h or -? option, gives the status of the HASSECURITY and HASNOSOCKSECURITY definitions.

See the Security section of the 00README file of the lsOF distribution for information on building lsOF with the HASSECURITY and HASNOSOCKSECURITY options enabled.

Creation and use of a user-readable and user-writable device cache file is controlled by the compile-time HASDCACHE option. See the DEVICE CACHE FILE section and the sections that follow it for details on how its path is formed. For security considerations it is important to note that in the default lsOF distribution, if the real user ID under which lsOF is executed is root, the device cache file will be written in root's home directory -- e.g., ~/ or /root. When HASDCACHE is not defined, lsOF does not write or attempt to read a device cache file.

When HASDCACHE is defined, the lsOF help output, presented in response to the -h, -D?, or -? options, will provide device cache file handling information. When HASDCACHE is not defined, the -h or -? output will have no -D option description.

Before you decide to disable the device cache file feature -- enabling it improves the performance of lsOF by reducing the startup overhead of examining all the nodes in /dev (or /devices) -- read the discussion of it in the 00DCACHE file of the lsOF distribution and the lsOF FAQ (The FAQ section gives its location.)

When in doubt, you can temporarily disable the use of the device cache file with the -Di option.

When lsOF user declares alternate kernel name list or memory files with the -k and -m options, lsOF checks the user's authority to read them with access(2). This is intended to prevent whatever special power lsOF's modes might confer on it from letting it read files not normally accessible via the authority of the real user ID.

OUTPUT

This section describes the information lsOF lists for each open file. See the OUTPUT FOR OTHER PROGRAMS section for additional information on output that can be processed by another program.

LsOF only outputs printable (declared so by isprint(3)) 8-bit characters. Non-printable characters are printed in one of three forms: the C '\bfrnt]' form; the control character '-' form (e.g., '@'); or hexadecimal leading '\x' form (e.g., '\xab'). Space is non-printable in the COMMAND column ('\x20') and printable elsewhere.

For some dialects -- if HASSETLOCALE is defined in the dialect's machine.h header file -- lsOF will print the extended 8-bit characters of a language locale. The lsOF process must be supplied a language locale environment variable (e.g., LANG) whose value represents a known language locale in which the extended characters are considered printable by isprint(3). Otherwise lsOF considers the extended characters non-printable and prints them according to its rules for non-printable characters, stated above. Consult your dialect's setlocale(3) man page for the names of other environment variables that may be used in place of LANG -- e.g., LC_ALL, LC_CTYPE, etc.

LsOF's language locale support for a dialect also covers wide characters -- e.g., UTF-8 -- when HASSETLOCALE and HASWIDECHAR are defined in the dialect's machine.h header file, and when a suitable language locale has been defined in the appropriate environment variable for the lsOF process. Wide characters are printable under those conditions if iswprint(3) reports them to be. If HASSETLOCALE, HASWIDECHAR and a suitable language locale aren't defined, or if iswprint(3) reports wide characters that aren't printable, lsOF considers the wide characters non-printable and prints each of their 8 bits according to its rules for non-printable characters, stated above.

Consult the answers to the "Language locale support" questions in the lsOF FAQ (The FAQ section gives its location.) for more information.

LsOF dynamically sizes the output columns each time it runs, guaranteeing that each column is a minimum size. It also guarantees that each column is separated from its predecessor by at least one space.

COMMAND

contains the first nine characters of the name of the UNIX **command** associated with the process. If a non-zero *w* value is specified to the *+c w* option, the column contains the first *w* characters of the name of the UNIX **command** associated with the process up to the limit of characters supplied to *ls* of by the UNIX dialect. (See the description of the *+c w* **command** or the *ls* of FAQ for more information. The FAQ section gives its location.)

If *w* is less than the length of the column title, 'COMMAND', it will be raised to that length.

If a zero *w* value is specified to the *+c w* option, the column contains all the characters of the name of the UNIX **command** associated with the process.

All **command** name characters maintained by the kernel in its structures are displayed in field output when the **command** name descriptor ('c') is specified. See the OUTPUT_FOR_OTHER_COMMANDS section for information on selecting field output and the associated command name descriptor.

-----PID----- is the Process IDentification number of the process.

-----TID----- is the task (thread) IDentification number, if task (thread) reporting is supported by the dialect and a task (thread) is being listed. (If help output --i.e., the output of the -h or -? options-- shows this option, then task (thread) reporting is supported by the dialect.)

-----A blank TID column in Linux indicates a process --i.e., a non-task.

-----TASKCMD----- is the task command name. Generally this will be the same as the process named in the COMMAND column, but some task implementations (e.g., Linux) permit a task to change its command name.

-----The TASKCMD column width is subject to the same size limitation as the COMMAND column.

-----ZONE----- is the Solaris 10 and higher zone name. This column must be selected with the -Z option.

-----SECURITY-CONTEXT----- is the SELinux security context. This column must be selected with the -Z option. Note that the -Z option is inhibited when SELinux is disabled in the running Linux kernel.

-----PPID----- is the Parent Process IDentification number of the process. It is only displayed when the -R option has been specified.

-----PGID----- is the process group IDentification number associated with the process. It is only displayed when the -g option has been specified.

-----USER----- is the user ID number or login name of the user to whom the process belongs, usually the same as reported by *ps(1)*. However, on Linux USER is the user ID number or login that owns the directory in */proc* where *ls* of finds information about the process. Usually that is the same value reported by *ps(1)*, but may differ when the process has changed its effective user ID. (See the -l option description for information on when a user ID number or login name is displayed.)

-----FD----- is the File Descriptor number of the file or:

-----c wd----- current working directory;
-----Lnn----- library references (AIX);
-----err FD----- information error (see NAME column);
-----jld----- jail directory (FreeBSD);
-----ltx----- shared library text (code and data);
-----Mxx----- hex memory-mapped type number xx;
-----m86----- DOS Merge mapped file;
-----mem----- memory-mapped file;
-----mmap----- memory-mapped device;
-----pd----- parent directory;
-----rtd----- root directory;
-----tr----- kernel trace file (OpenBSD);
-----txt----- program text (code and data);
-----v86----- VP/ix mapped file;

-----FD is followed by one of these characters, describing the mode under which the file is open:

-----r----- for read access;

```

-----w_for_write_access;
-----u_for_read_and_write_access;
-----space_if_mode_unknown_and_no_lock
-----character_follows;
-----'_' if mode unknown and lock
-----character follows.

```

The mode character is followed by one of these lock characters, describing the **type** of lock applied to the file:

```

    N for a Solaris NFS lock of unknown type;
    r for read lock on part of the file;
    R for a read lock on the entire file;
    w for a write lock on part of the file;
    W for a write lock on the entire file;
    u for a read and write lock of any length;
    U for a lock of unknown type;
    x for an SCO OpenServer Xenix lock on part of
the file;
    X for an SCO OpenServer Xenix lock on the entire
file;
    space if there is no lock.

```

See the LOCKS section for more information on the lock information character.

The FD column contents constitutes a single field for parsing in post-processing scripts.

TYPE is the **type** of the node associated with the file — e.g., GDIR, GREG, VDIR, VREG, etc.

```

or 'IPv4' for an IPv4 socket;
or 'IPv6' for an open IPv6 network file — even if its
address is IPv4, mapped in an IPv6 address;
or 'ax25' for a Linux AX.25 socket;
or 'inet' for an Internet domain socket;
or 'lla' for a HP-UX link level access file;
or 'rte' for an AF_ROUTE socket;
or 'sock' for a socket of unknown domain;
or 'unix' for a UNIX domain socket;
or 'x.25' for an HP-UX x.25 socket;
or 'BLK' for a block special file;
or 'CHR' for a character special file;
or 'DEL' for a Linux map file that has been deleted;
or 'DIR' for a directory;
or 'DOOR' for a VDOOR file;
or 'FIFO' for a FIFO special file;
or 'KQUEUE' for a BSD style kernel event queue file;
or 'LINK' for a symbolic link file;
or 'MPB' for a multiplexed block file;
or 'MPC' for a multiplexed character file;
or 'NOFD' for a Linux /proc/<PID>/fd directory that
can't be opened — the directory path appears in the NAME
-----column, followed by an error message;
-----or 'PAS' for a_/proc/as_file;
-----or 'PAXV' for a_/proc/auxv_file;
-----or 'PCRE' for a_/proc/cred_file;
-----or 'PCTL' for a_/proc/control_file;
-----or 'PCUR' for the current_/proc/process;
-----or 'PCWD' for a_/proc/current_working_directory;
-----or 'PDIR' for a_/proc/directory;
-----or 'PETY' for a_/proc/executable_type_(etype);

```

```

-----or_ 'PFD' '_for_a_/proc_file_descriptor;
-----or_ 'PFDR' '_for_a_/proc_file_descriptor_directory;
-----or_ 'PFIL' '_for_an_executable_/proc_file;
-----or_ 'PFPR' '_for_a_/proc_FP_register_set;
-----or_ 'PGD' '_for_a_/proc/pagedata_file;
-----or_ 'PGID' '_for_a_/proc_group_notifier_file;
-----or_ 'PIPE' '_for_pipes;
-----or_ 'PLC' '_for_a_/proc/lwpctl_file;
-----or_ 'PLDR' '_for_a_/proc/lpw_directory;
-----or_ 'PLDT' '_for_a_/proc/ldt_file;
-----or_ 'PLPI' '_for_a_/proc/lpsinfo_file;
-----or_ 'PLST' '_for_a_/proc/lstatus_file;
-----or_ 'PLU' '_for_a_/proc/lusage_file;
-----or_ 'PLWG' '_for_a_/proc/gwindows_file;
-----or_ 'PLWI' '_for_a_/proc/lwpsinfo_file;
-----or_ 'PLWS' '_for_a_/proc/lwpstatus_file;
-----or_ 'PLWU' '_for_a_/proc/lwpusage_file;
-----or_ 'PLWX' '_for_a_/proc/xregs_file;
-----or_ 'PMAP' '_for_a_/proc_map_file_(map);
-----or_ 'PMEM' '_for_a_/proc_memory_image_file;
-----or_ 'PNTF' '_for_a_/proc_process_notifier_file;
-----or_ 'POBJ' '_for_a_/proc/object_file;
-----or_ 'PODR' '_for_a_/proc/object_directory;
-----or_ 'POLP' '_for_an_old_format_/proc_light_weight_process
-----file;
-----or_ 'POPF' '_for_an_old_format_/proc_PID_file;
-----or_ 'POPG' '_for_an_old_format_/proc_page_data_file;
-----or_ 'PORT' '_for_a_SYSV_named_pipe;
-----or_ 'PREG' '_for_a_/proc_register_file;
-----or_ 'PRMP' '_for_a_/proc/rmap_file;
-----or_ 'PRTD' '_for_a_/proc_root_directory;
-----or_ 'PSGA' '_for_a_/proc/sigact_file;
-----or_ 'PSIN' '_for_a_/proc/psinfo_file;
-----or_ 'PSTA' '_for_a_/proc_status_file;
-----or_ 'PSXSEM' '_for_a_POSIX_semaphore_file;
-----or_ 'PSXSHM' '_for_a_POSIX_shared_memory_file;
-----or_ 'PTS' '_for_a_/dev/pts_file;
-----or_ 'PUSG' '_for_a_/proc/usage_file;
-----or_ 'PW' '_for_a_/proc/watch_file;
-----or_ 'PXMP' '_for_a_/proc/xmap_file;
-----or_ 'REG' '_for_a_regular_file;
-----or_ 'SMT' '_for_a_shared_memory_transport_file;
-----or_ 'STSO' '_for_a_stream_socket;
-----or_ 'UNNM' '_for_an_unnamed_type_file;
-----or_ 'XNAM' '_for_an_OpenServer_Xenix_special_file_of
-----unknown_type;
-----or_ 'XSEM' '_for_an_OpenServer_Xenix_semaphore_file;

```



```

-----or 'XSD' for an OpenServer_Xenix_shared_data_file;

-----or the four_type_number_octets if the corresponding_name
-----isn't known.

FILE-ADDR
    contains the kernel file structure address when f has been
    specified to +f;

FCT
    contains the file reference count from the kernel file
    structure when c has been specified to +f;

FILE-FLAG
    when g or G has been specified to +f, this field contains
    the contents of the f_flag[s] member of the kernel file
    structure and the kernel's_per-process_open_file_flags_(if
    available); 'G' causes them to be displayed in
    hexadecimal; 'g' causes them to be displayed in
    short-hand names; two_lists_may_be
    displayed_with_entries_separated_by_commas_the_lists
    separated_by_a_semicolon(';'); the first list may contain
    short-hand names for f_flag[s] values from the following
    table:

AIO      asynchronous I/O (e.g., FAIO)
AP        append
ASYN      asynchronous I/O (e.g., FASYNC)
BAS       block, test, and set in use
BKIU      block if in use
BL        use block offsets
BSK       block seek
CA        copy avoid
CIO       concurrent I/O
CLON      clone
CLRD      CL read
CR        create
DF        defer
DFI       defer IND
DFLU      data flush
DIR       direct
DLY       delay
DOCL      do clone
DSYN      data-only integrity
DTY       must be a directory
EVO       event only
EX        open for exec
EXCL      exclusive open
FSYN      synchronous writes
GCDF      defer during unp_gc() (AIX)
GCMK      mark during unp_gc() (AIX)
GTTY      accessed via /dev/tty
HUP       HUP in progress
KERN      kernel
KIOC      kernel-issued ioctl
LCK       has lock
LG        large file
MBLK      stream message block
MK        mark
MNT       mount
MSYN      multiplex synchronization
NATM      don't update atime
-----NB-----non-blocking_I/O
-----NBDR-----no_BDRM_check
-----NBIO-----SYSV_non-blocking_I/O
-----NBF-----n-buffering_in_effect
-----NC-----no_cache
-----ND-----no_delay
-----NDSY-----no_data_synchronization
-----NET-----network
-----NFLK-----don't follow links
NMFS      NM file system
NOTO      disable background stop
NSH       no share
NTTY      no controlling TTY
OLRM      OLR mirror
PAIO      POSIX asynchronous I/O
PP        POSIX pipe
R         read
RC        file and record locking cache
REV       revoked
RSH       shared read
RSYN      read synchronization
RW        read and write access
SL        shared lock
SNAP      cooked snapshot
SOCK      socket
SQSH      Sequent shared set on open
SQSV      Sequent SVM set on open
SQR       Sequent set repair on open
SQS1      Sequent full shared open
SQS2      Sequent partial shared open
STPI      stop I/O
SWR       synchronous read

```

SYN	file integrity while writing
TCPM	avoid TCP collision
TR	truncate
W	write
WKUP	parallel I/O synchronization
WTG	parallel I/O synchronization
VH	vhangup pending
VTXT	virtual text
XL	exclusive lock

this list of names was derived from F* *#define* 's in dialect header files <fcntl.h>, <linux/fs.h>, <sys/fcntl.c>, <sys/fcntlcom.h>, and <sys/file.h>; see the lsof.h header file for a list showing the correspondence between the above short-hand names and the header file definitions;

the second list (after the semicolon) may contain short-hand names for kernel per-process open file flags from this table:

ALLC	allocated
BR	the file has been read
BHUP	activity stopped by SIGHUP
BW	the file has been written
CLSG	closing
CX	close-on-exec (see fcntl(F.SETFD))
LCK	lock was applied
MP	memory-mapped
OPIP	open pending - in progress
RSVW	reserved wait
SHMT	UF_SHMAT set (AIX)
USE	in use (multi-threaded)

NODE-ID

(or INODE-ADDR for some dialects) contains a unique identifier for the file node (usually the kernel vnode or inode address, but also occasionally a concatenation of device and node number) when n has been specified to +f;

DEVICE

contains the device numbers, separated by commas, for a character special, block special, regular, directory or NFS file;

or 'memory' for a memory file system node under Tru64 UNIX;

or the address of the private data area of a Solaris socket stream;

or a kernel reference address that identifies the file (The kernel reference address may be used for FIFO's, for example.);

-----or the base address or device name of a Linux AX.25 socket device.

-----Usually only the lower thirty-two bits of Tru64 UNIX kernel addresses are displayed.

-----SIZE, _SIZE/OFF, or _OFFSET is the size of the file or the file offset in bytes. _A value is displayed in this column only if it is available. Lsof displays whatever value _size or _offset _is appropriate for the type of the file and the version of lsof.

-----On some UNIX dialects lsof can't obtain accurate or consistent file offset information from its kernel data sources, sometimes just for particular kinds of files (e.g., socket files.) In other cases, files don't have true sizes -- e.g., sockets, FIFOs, pipes -- so lsof displays for their sizes the content amounts it finds in their kernel buffer descriptors (e.g., socket buffer size counts or TCP/IP window sizes.) Consult the lsof FAQ (The FAQ section gives its location.) for more information.

-----The file _size is displayed in decimal; the offset is normally displayed in decimal with a leading '0t' if it contains 8 digits or less; in hexadecimal with a leading '0x' if it is longer than 8 digits. (Consult the -o.o option description for information on when 8 might default to some other value.)

-----Thus the leading '0t' and '0x' identify an offset when the column may contain both a size and an offset (i.e., its title is SIZE/OFF).

-----If the -o.o option is specified, lsof always displays the file offset (or nothing if no offset is available) and labels the column _OFFSET. The offset always begins with '0t' or '0x' as described above.

```

-----The -lsof user can control the switch from '-0t' to '-0x'
-----with the -o-o option. Consult its description for more
-----information.

-----If the -s option is specified, -lsof always displays the
-----file size (or nothing if no size is available) and labels
-----the column SIZE. The -o and -s options are mutually
-----exclusive; they can't both be specified.

For files that don't have a fixed size - e.g., don't
reside on a disk device - lsof will display appropriate
information about the current size or position of the file
if it is available in the kernel structures that define
the file.

NLINK contains the file link count when +L has been specified;

NODE is the node number of a local file;

or the inode number of an NFS file in the server host;

or the Internet protocol type - e.g., 'TCP';

or 'STR' for a stream;

or 'CCITT' for an HP-UX x.25 socket;

or the IRQ or inode number of a Linux AX.25 socket device.

NAME is the name of the mount point and file system on which
the file resides;

or the name of a file specified in the names option (after
any symbolic links have been resolved);

or the name of a character special or block special
device;

or the local and remote Internet addresses of a network
file; the local host name or IP number is followed by a
colon (':'), the port, '->', and the two-part remote
address; IP addresses may be reported as numbers or names,
depending on the +|-M, -n, and -P options; colon-separated
IPv6 numbers are enclosed in square brackets; IPv4
INADDR_ANY and IPv6 IN6_IS_ADDR_UNSPECIFIED addresses, and
zero port numbers are represented by an asterisk ('*'); a
UDP destination address may be followed by the amount of
time elapsed since the last packet was sent to the
destination; TCP, UDP and UDPLITE remote addresses may be
followed by TCP/TPI information in parentheses - state
(e.g., '(ESTABLISHED)', '(Unbound)'), queue sizes, and
window sizes (not all dialects) - in a fashion similar to
what netstat(1) reports; see the -T option description or
the description of the TCP/TPI field in OUTPUT FOR OTHER
PROGRAMS for more information on state, queue size, and
window size;

or the address or name of a UNIX domain socket, possibly
including a stream clone device name, a file system
object's_path_name, local and foreign kernel addresses,
socket_pair_information, and a_bound_vnode_address;

-----or the local and remote mount point names of an NFS file;

-----or 'STR', followed by the stream name;

-----or a stream character device name, followed by '->' and
-----the stream name or a list of stream module names,
-----separated by '->';

-----or 'STR:' followed by the SCO OpenServer stream device
-----and module names, separated by '->';

-----or system directory name, '- - -', and as many components
-----of the path name as -lsof can find in the kernel's name
-----cache for selected dialects (See the KERNEL NAME CACHE
-----section for more information.);

or 'PIPE->', followed by a Solaris kernel pipe
destination address;

or 'COMMON:', followed by the vnode device information
structure's_device_name, for a Solaris common vnode;

-----or the address family, followed by a slash ('/'), followed
-----by fourteen comma-separated bytes of a non-Internet raw
-----socket address;

or the HP-UX x.25 local address, followed by the virtual
connection number (if any), followed by the remote address
(if any);

```

```

or '(dead)' for disassociated Tru64 UNIX files —
typically terminal files that have been flagged with the
TIOCNOTTY ioctl and closed by daemons;

or 'rd=<offset>' and 'wr=<offset>' for the values of
the read and write offsets of a FIFO;

or 'clone n:/dev/event' for SCO OpenServer file clones
of the /dev/event device, where n is the minor device
number of the file;

or '(socketpair: n)' for a Solaris 2.6, 8, 9 or 10 UNIX
domain socket, created by the socketpair(3N) network
function;

or 'no PCB' for socket files that do not have a protocol
block associated with them, optionally followed by ',',
CANTSENDMORE' if sending on the socket has been disabled,
or ', CANTRCVMORE' if receiving on the socket has been
disabled (e.g., by the shutdown(2) function);

or the local and remote addresses of a Linux IPX socket
file in the form <net>[:<node>]:<port>, followed in
parentheses by the transmit and receive queue sizes, and
the connection state;

or 'dgram' or 'stream' for the type UnixWare 7.1.1 and
above in-kernel UNIX domain sockets, followed by a colon
(':') and the local path name when available, followed by
'-'>' and the remote path name or kernel socket address
in hexadecimal when available;

or the association value, association index, endpoint
value, local address, local port, remote address and
remote port for Linux SCTP sockets;

or 'protocol: ' followed by the Linux socket's protocol
attribute.

-----For dialects that support a 'namefs' file system, allowing one
file to be attached to another with fattach(3C), lsof will add
' (FA:<address1><direction><address2>)' to the NAME column.
<address1> and <address2> are hexadecimal vnode addresses.
<direction> will be '<->' if <address2> has been fattach'ed to
this vnode whose address is <address1>; and '<->' if <address1>,
the vnode address of this vnode, has been fattach'ed to
<address2>. <address1> may be omitted if it already appears in
the DEVICE column.

-----Lsof may add two parenthetical notes to the NAME column for open
Solaris 10 files: - '(?)' if lsof considers the path name of
questionable accuracy; and - '(deleted)' if the -X option has
been specified and lsof detects the open file's path name has
been deleted. Consult the lsof FAQ (The FAQ section gives its
location.) for more information on these NAME column additions.

LOCKS
Lsof can't adequately report the wide variety of UNIX dialect
file locks in a single character. What it reports in a single
character is a compromise between the information it finds in the
kernel and the limitations of the reporting format.

-----Moreover, when a process holds several byte-level locks on a
file, lsof only reports the status of the first lock it
encounters. If it is a byte-level lock, then the lock character
will be reported in lower case — i.e., 'r', 'w', or 'x' — rather
than the upper case equivalent reported for a full file lock.

Generally lsof can only report on locks held by local processes
on local files. When a local process sets a lock on a remotely
mounted (e.g., NFS) file, the remote server host usually records
the lock state. One exception is Solaris — at some patch levels
of 2.3, and in all versions above 2.4, the Solaris kernel records
information on remote locks in local structures.

Lsof has trouble reporting locks for some UNIX dialects. Consult
the BUGS section of this manual page or the lsof FAQ (The FAQ
section gives its location.) for more information.

OUTPUT FOR OTHER PROGRAMS
When the -F option is specified, lsof produces output that is
suitable for processing by another program — e.g., an awk or Perl
script, or a C program.

Each unit of information is output in a field that is identified
with a leading character and terminated by a NL (012) (or a NUL
(000) if the 0 (zero) field identifier character is specified.)
The data of the field follows immediately after the field
identification character and extends to the field terminator.

It is possible to think of field output as process and file sets.
A process set begins with a field whose identifier is 'p' (for
-----process Identifier (PID)). It extends to the beginning of the

```

```

-----next_PID_field_or_the_beginning_of_the_first_file_set_of_the
-----process, whichever comes first. Included in the process set are
-----fields that identify the command, the process group
-----IDentification (PGID) number, the task (thread) ID (TID), and the
-----user ID (UID) number or login name.

-----A file set begins with a field whose identifier is 'f' (for file
-----descriptor). It is followed by lines that describe the file's
-----access mode, lock state, type, device, size, offset, inode,
-----protocol, name and stream module names. It extends to the
-----beginning of the next file or process set, whichever comes first.

-----When the NUL (000) field terminator has been selected with the 0
----- (zero) field identifier character, lsof ends each process and
----- file set with an NL (012) character.

-----Lsof always produces one field, the PID ('p') field. All other
----- fields may be declared optionally in the field identifier
----- character list that follows the -F option. When a field
----- selection character identifies an item lsof does not normally
----- list - e.g., PPID, selected with -R - specification of the field
----- character - e.g., '-FR' - also selects the listing of the item.

It is entirely possible to select a set of fields that cannot
easily be parsed - e.g., if the field descriptor field is not
selected, it may be difficult to identify file sets. To help you
avoid this difficulty, lsof supports the -F option; it selects
the output of all fields with NL terminators (the -F0 option pair
selects the output of all fields with NUL terminators). For
compatibility reasons neither -F nor -F0 select the raw device
field.

These are the fields that lsof will produce. The single
character listed first is the field identifier.

a file access mode
c process command name (all characters from proc or
  user structure)
C file structure share count
d file 's_device character code
-----D----- file 's major/minor device number (0x<hexadecimal>)
f file descriptor (always selected)
F file structure address (0x<hexadecimal>)
G file flaGs (0x<hexadecimal>; names if +fg follows)
g process group ID
i file 's_inode_number
-----K----- task_ID
-----k----- link_count
-----l----- file 's lock status
L process login name
m marker between repeated output
M the task comMand name
n file name, comment, Internet address
N node identifier (ox<hexadecimal>)
o file 's_offset (decimal)
-----p----- process_ID (always selected)
-----P----- protocol_name
-----r----- raw_device_number (0x<hexadecimal>)
-----R----- parent_process_ID
-----s----- file 's size (decimal)
-----S----- file 's_stream_identification
-----t----- file 's type
-----T----- TCP/TPI information, identified by prefixes (the
----- 'is' is part of the prefix):
-----QR<read_queue_size>
-----QS<send_queue_size>
-----SO<socket_options_and_values> (not all dialects)
-----SS<socket_states> (not all dialects)
-----ST<connection_state>
-----TF<TCP_flags_and_values> (not all dialects)
-----WR<window_read_size> (not all dialects)
-----WW<window_write_size> (not all dialects)
----- (TCP/TPI information isn't reported for all supported
----- UNIX dialects. The -h or -? help output for the
----- -T option will show what TCP/TPI reporting can be
----- requested.)
u process user ID
z Solaris 10 and higher zone name
Z SELinux security context (inhibited when SELinux is disabled)
0 use NUL field terminator character in place of NL
1-9 dialect-specific field identifiers (The output
  of -F? identifies the information to be found
  in dialect-specific fields.)

You can get on-line help information on these characters and
their descriptions by specifying the -F? option pair. (Escape
the '?' character as your shell requires.) Additional
information on field content can be found in the OUTPUT section.

-----As an example, '-F-pcfn' will select the process_ID ('p'),
----- command name ('c'), file_descriptor ('f') and file name ('n')
----- fields with an NL field terminator character; '-F-pcfn0'

```

```

-----selects the same output with a NUL (000) field terminator
-----character.

-----Lsof doesn't produce all fields for every process or file set,
only those that are available. Some fields are mutually
exclusive: file device characters and file major/minor device
numbers; file inode number and protocol name; file name and
stream identification; file size and offset. One or the other
member of these mutually exclusive sets will appear in field
output, but not both.

Normally lsof ends each field with a NL (012) character. The 0
(zero) field identifier character may be specified to change the
field terminator character to a NUL (000). A NUL terminator may
be easier to process with xargs(1), for example, or with programs
whose quoting mechanisms may not easily cope with the range of
characters in the field output. When the NUL field terminator is
in use, lsof ends each process and file set with a NL (012).

Three aids to producing programs that can process lsof field
output are included in the lsof distribution. The first is a C
header file, lsof.fields.h, that contains symbols for the field
identification characters, indexes for storing them in a table,
and explanation strings that may be compiled into programs. Lsof
uses this header file.

The second aid is a set of sample scripts that process field
output, written in awk, Perl 4, and Perl 5. They're located in
-----the scripts subdirectory of the lsof distribution.

-----The third aid is the C library used for the lsof test suite. The
test suite is written in C and uses field output to validate the
correct operation of lsof. The library can be found in the
tests/LTlib.c file of the lsof distribution. The library uses
-----the first aid, the lsof.fields.h header file.
BLOCKS AND TIMEOUTS
-----Lsof can be blocked by some kernel functions that it uses --
lstat(2), readlink(2), and stat(2). These functions are stalled
-----in the kernel, for example, when the hosts where mounted NFS file
systems reside become inaccessible.

-----Lsof attempts to break these blocks with timers and child
processes, but the techniques are not wholly reliable. When lsof
-----does manage to break a block, it will report the break with an
error message. The messages may be suppressed with the -t and -w
-----options.

-----The default timeout value may be displayed with the -h or -?
option, and it may be changed with the -S[t] option. The
-----minimum for t is two seconds, but you should avoid small values,
since slow system responsiveness can cause short timeouts to
-----expire unexpectedly and perhaps stop lsof before it can produce
any output.

-----When lsof has to break a block during its access of mounted file
system information, it normally continues, although with less
-----information available to display about open files.

-----Lsof can also be directed to avoid the protection of timers and
child processes when using the kernel functions that might block
-----by specifying the -O option. While this will allow lsof to start
up with less overhead, it exposes lsof completely to the kernel
-----situations that might block it. Use this option cautiously.
AVOIDING KERNEL BLOCKS
-----You can use the -b option to tell lsof to avoid using kernel
functions that would block. Some cautions apply.

-----First, using this option usually requires that your system supply
alternate device numbers in place of the device numbers that lsof
-----would normally obtain with the lstat(2) and stat(2) kernel
functions. See the ALTERNATE_DEVICE_NUMBERS section for more
-----information on alternate device numbers.

-----Second, you can't specify names for lsof to locate unless they're
file system names. This is because lsof needs to know the device
-----and inode numbers of files listed with names in the lsof options,
and the -b option prevents lsof from obtaining them. Moreover,
-----since lsof only has device numbers for the file systems that have
alternates, its ability to locate files on file systems depends
-----completely on the availability and accuracy of the alternates.
If no alternates are available, or if they're incorrect, lsof
-----won't be able to locate files on the named file systems.

-----Third, if the names of your file system directories that lsof
obtains from your system's mount table are symbolic links, lsof
-----won't be able to resolve the links. This is because the -b
option causes lsof to avoid the kernel readlink(2) function it
-----uses to resolve symbolic links.

-----Finally, using the -b option causes lsof to issue warning
messages when it needs to use the kernel functions that the -b
-----option directs it to avoid. You can suppress these messages by

```

```

specifying the -w option, but if you do, you won't see the
    alternate device numbers reported in the warning messages.
ALTERNATE DEVICE NUMBERS
    On some dialects, when lsuf has to break a block because it can't
    get information about a mounted file system via the lstat(2) and
    stat(2) kernel functions, or because you specified the -b option,
    lsuf can obtain some of the information it needs -- the device
    number and possibly the file system type -- from the system mount
    table. When that is possible, lsuf will report the device number
    it obtained. (You can suppress the report by specifying the -w
    option.)

You can assist this process if your mount table is supported with
an /etc/mtab or /etc/mnttab file that contains an options field
by adding a "dev=xxxx" field for mount points that do not have
one in their options strings. Note: you must be able to edit the
file -- i.e., some mount tables like recent Solaris /etc/mnttab or
Linux /proc/mounts are read-only and can't be modified.

    You may also be able to supply device numbers using the +m and +n
    options, provided they are supported by your dialect. Check
    the output of lsuf's -h or -? options to see if the +m and +n
    options are available.

The "xxxx" portion of the field is the hexadecimal value of the
file system's device number. (Consult the st_dev field of the
output of the lstat(2) and stat(2) functions for the appropriate
values for your file systems.) Here's an example from a Sun
Solaris 2.6 /etc/mnttab for a file system remotely mounted via
NFS:

nfs ignore, noquota, dev=2a40001

There's an advantage to having "dev=xxxx" entries in your mount
table file, especially for file systems that are mounted from
remote NFS servers. When a remote server crashes and you want
to identify its users by running lsuf on one of its clients, lsuf
probably won't be able to get output from the lstat(2) and
stat(2) functions for the file system. If it can obtain the file
system's device number from the mount table, it will be able to
display the files open on the crashed NFS server.

Some dialects that do not use an ASCII /etc/mtab or /etc/mnttab
file for the mount table may still provide an alternative device
number in their internal mount tables. This includes AIX, Apple
Darwin, FreeBSD, NetBSD, OpenBSD, and Tru64 UNIX. Lsuf knows how
to obtain the alternative device number for these dialects and
uses it when its attempt to lstat(2) or stat(2) the file system
is blocked.

    If you're not sure your dialect supplies alternate device numbers
    for file systems from its mount table, use this lsuf incantation
    to see if it reports any alternate device numbers:

lsuf -b

Look for standard error file warning messages that begin
"assuming dev=xxxx" from "...".
KERNEL_NAME_CACHE
    Lsuf is able to examine the kernel's name cache or use other
    kernel facilities (e.g., the ADVFS 4.x tag.to.path() function
    under Tru64 UNIX) on some dialects for most file system types,
    excluding AFS, and extract recently used path name components
    from it. (AFS file system path lookups don't use the kernel's
    name cache; some Solaris VxFS file system operations apparently
    don't use it, either.)

Lsuf reports the complete paths it finds in the NAME column. If
lsuf can't report all components in a path, it reports in the
NAME column the file system name, followed by a space, two '-'
characters, another space, and the name components it has
located, separated by the '/' character.

    When lsuf is run in repeat mode -- i.e., with the -r option
    specified -- the extent to which it can report path name
    components for the same file may vary from cycle to cycle.
    That's because other running processes can cause the kernel to
    remove entries from its name cache and replace them with others.

Lsuf's use of the kernel name cache to identify the paths of
files can lead it to report incorrect components under some
circumstances. This can happen when the kernel name cache uses
device and node number as a key (e.g., SCO OpenServer) and a key
on a rapidly changing file system is reused. If the UNIX
dialect's kernel doesn't purge the name cache entry for a file
when it is unlinked, lsuf may find a reference to the wrong entry
in the cache. The lsuf FAQ (The FAQ section gives its location.)
has more information on this situation.

Lsuf can report path name components for these dialects:

FreeBSD

```

```

        HP-UX
        Linux
        NetBSD
        NEXTSTEP
        OpenBSD
        OPENSTEP
        SCO OpenServer
        SCO|Caldera UnixWare
        Solaris
        Tru64 UNIX

Lsof can't report path name components for these dialects:

-----AIX

-----If you want to know why lsof can't report path name components
for some dialects, see the lsof FAQ (The FAQ section gives its
location.)
DEVICE CACHE FILE
    Examining all members of the /dev (or /devices) node tree with
    stat(2) functions can be time consuming. What's more, the
    information that lsof needs -- device number, inode number, and
    path -- rarely changes.

-----Consequently, lsof normally maintains an ASCII text file of
cached /dev (or /devices) information (exception: the /proc-based
Linux lsof where it's not needed.) The local system
administrator who builds lsof can control the way the device
cache file path is formed, selecting from these options:

    Path from the -D option;
    Path from an environment variable;
    System-wide path;
    Personal path (the default);
    Personal path, modified by an environment variable.

Consult the output of the -h, -D?, or -? help options for the
current state of device cache support. The help output lists the
default read-mode device cache file path that is in effect for
the current invocation of lsof. The -D? option output lists the
read-only and write device cache file paths, the names of any
applicable environment variables, and the personal device cache
path format.

Lsof can detect that the current device cache file has been
accidentally or maliciously modified by integrity checks,
including the computation and verification of a sixteen bit
Cyclic Redundancy Check (CRC) sum on the file's contents. When
lsof senses something wrong with the file, it issues a warning
and attempts to remove the current cache file and create a new
copy, but only to a path that the process can legitimately write.

-----The path from which a lsof process may attempt to read a device
cache file may not be the same as the path to which it can
legitimately write. Thus when lsof senses that it needs to
update the device cache file, it may choose a different path for
writing it from the path from which it read an incorrect or
outdated version.

-----If available, the -Dr option will inhibit the writing of a new
device cache file. (It's always available when specified without
a path name argument.)

When a new device is added to the system, the device cache file
may need to be recreated. Since lsof compares the mtime of the
device cache file with the mtime and ctime of the /dev (or
/devices) directory, it usually detects that a new device has
been added; in that case lsof issues a warning message and
attempts to rebuild the device cache file.

Whenever lsof writes a device cache file, it sets its ownership
to the real UID of the executing process, and its permission
modes to 0600, this restricting its reading and writing to the
file's owner.
LSOF_PERMISSIONS_THAT_AFFECT_DEVICE_CACHE_FILE_ACCESS
-----Two permissions of the lsof executable affect its ability to
access device cache files. The permissions are set by the local
system administrator when lsof is installed.

-----The first and rarer permission is setuid-root. It comes into
effect when lsof is executed; its effective UID is then root,
while its real (i.e., that of the logged-on user) UID is not.
The lsof distribution recommends that versions for these dialects
run setuid-root.

-----HP-UX 11.11 and 11.23
-----Linux

-----The second and more common permission is setgid. It comes into
effect when the effective group identification number (GID) of
the lsof process is set to one that can access kernel memory
devices -- e.g., 'kmem', 'sys', or 'system'.

```



```

-----An_lsof_process_that_has_setgid_permission_usually_surrenders_the
-----permission_after_it_has_accessed_the_kernel_memory_devices.  When
-----it_does_that,_lsof_can_allow_more_liberal_device_cache_path
-----information.  The_lsof_distribution_recommends_that_versions_for
-----these_dialects_run_setgid_and_be_allowed_to_surrender_setgid
-----permission.

-----AIX_5.[12]_and_5.3-ML1
-----Apple_Darwin_7.x_Power_Macintosh_systems
-----FreeBSD_4.x,_4.1x,_5.x_and_[6789].x_for_x86-based_systems
-----FreeBSD_5.x,_[6789].x_and_1[012].8_for_Alpha,_AMD64_and_Sparc64
-----based_systems
-----HP-UX_11.00
-----NetBSD_1.[456],_2.x_and_3.x_for_Alpha,_x86,_and_SPARC-based
-----systems
-----NEXTSTEP_3.[13]_for_NEXTSTEP_architectures
-----OpenBSD_2.[89]_and_3.[0-9]_for_x86-based_systems
-----OPENSTEP_4.x
-----SCO_OpenServer_Release_5.0.6_for_x86-based_systems
-----SCO_Caldera_UnixWare_7.1.4_for_x86-based_systems
-----Solaris_2.6,_8,_9_and_10
-----Tru64_UNIX_5.1

----- (Note:_lsof_for_AIX_5L_and_above_needs_setuid-root_permission_if
----- its_-X_option_is_used.)

-----Lsof_for_these_dialects_does_not_support_a_device_cache,_so_the
----- permissions_given_to_the_executable_don't_apply_to_the_device
----- cache_file.

Linux
DEVICE CACHE FILE PATH FROM THE -D OPTION
The -D option provides limited means for specifying the device
cache file path. Its ? function will report the read-only and
write device cache file paths that lsof will use.

When the -D b, r, and u functions are available, you can use them
to request that the cache file be built in a specific location
(b[path]); read but not rebuilt (r[path]); or read and rebuilt
(u[path]). The b, r, and u functions are restricted under some
conditions. They are restricted when the lsof process is
setuid-root. The path specified with the r function is always
read-only, even when it is available.

The b, r, and u functions are also restricted when the lsof
process runs setgid and lsof doesn't surrender the setgid
permission. (See the LSOF_PERMISSIONS_THAT_AFFECT_DEVICE_CACHE
FILE_ACCESS section for a list of implementations that normally
don't surrender their setgid permission.)

A further -D function, i (for ignore), is always available.

When available, the b function tells lsof to read device
information from the kernel with the stat(2) function and build a
device cache file at the indicated path.

When available, the r function tells lsof to read the device
cache file, but not update it. When a path argument accompanies
-r, it names the device cache file path. The r function is
always available when it is specified without a path name
argument. If lsof is not running setuid-root and surrenders its
setgid permission, a path name argument may accompany the r
function.

When available, the u function tells lsof to attempt to read and
use the device cache file. If it can't read the file, or if it
finds the contents of the file incorrect or outdated, it will
read information from the kernel, and attempt to write an updated
version of the device cache file, but only to a path it considers
legitimate for the lsof process effective and real UIDs.
DEVICE_CACHE_PATH_FROM_AN_ENVIRONMENT_VARIABLE
-----Lsof's second choice for the device cache file is the contents of
the LSOFDEVCACHE environment variable. It avoids this choice if
the lsof process is setuid-root, or the real UID of the process
is root.

A further restriction applies to a device cache file path taken
from the LSOFDEVCACHE environment variable: lsof will not write a
device cache file to the path if the lsof process doesn't
surrender its setgid permission. (See the LSOF_PERMISSIONS_THAT
AFFECT_DEVICE_CACHE_FILE_ACCESS section for information on
implementations that don't surrender their setgid permission.)

The local system administrator can disable the use of the
LSOFDEVCACHE environment variable or change its name when
building lsof. Consult the output of -D? for the environment
variable's name.
SYSTEM-WIDE_DEVICE_CACHE_PATH
-----The_local_system_administrator_may_choose_to_have_a_system-wide
-----device_cache_file_when_building_lsof. That_file_will_generally
-----be_constructed_by_a_special_system_administration_procedure_when

```

```

-----the_system_is_booted_or_when_the_contents_of_/dev_or_/devices)
-----changes. If defined, it is lsof's third device cache file path
choice.

You can tell that a system-wide device cache file is in effect
for your local installation by examining the lsof help option
output - i.e., the output from the -h or -? option.

Lsof will never write to the system-wide device cache file path
by default. It must be explicitly named with a -D function in a
root-owned procedure. Once the file has been written, the
procedure must change its permission modes to 0644 (owner-read
and owner-write, group-read, and other-read).
PERSONAL DEVICE CACHE PATH (DEFAULT)
The default device cache file path of the lsof distribution is
one recorded in the home directory of the real UID that executes
lsof. Added to the home directory is a second path component of
the form .lsof_hostname.

This is lsof's fourth device cache file path choice, and is
usually the default. If a system-wide device cache file path was
defined when lsof was built, this fourth choice will be applied
when lsof can't find the system-wide device cache file. This is
the only time lsof uses two paths when reading the device cache
file.

The hostname part of the second component is the base name of the
executing host, as returned by gethostname(2). The base name is
defined to be the characters preceding the first '.' in the
-----gethostname(2) output, or all the gethostname(2) output if it
-----contains no '.'.

The device cache file belongs to the user ID and is readable and
writable by the user ID alone - i.e., its modes are 0600. Each
distinct real user ID on a given host that executes lsof has a
distinct device cache file. The hostname part of the path
distinguishes device cache files in an NFS-mounted home directory
into which device cache files are written from several different
hosts.

The personal device cache file path formed by this method
represents a device cache file that lsof will attempt to read,
and will attempt to write should it not exist or should its
contents be incorrect or outdated.

The -Dr option without a path name argument will inhibit the
writing of a new device cache file.

The -D? option will list the format specification for
constructing the personal device cache file. The conversions
used in the format specification are described in the 00DCACHE
file of the lsof distribution.
MODIFIED PERSONAL DEVICE CACHE PATH
If this option is defined by the local system administrator when
lsof is built, the LSOFPERSDCPATH environment variable contents
may be used to add a component of the personal device cache file
path.

The LSOFPERSDCPATH variable contents are inserted in the path at
the place marked by the local system administrator with the
"%p" conversion in the HASPERSDC format specification of the
dialect's machine.h header file. (It's placed right after the
home directory in the default lsof distribution.)

Thus, for example, if LSOFPERSDCPATH contains "LSOF", the home
directory is "/Homes/abe", the host name is
"lsof.itap.purdue.edu", and the HASPERSDC format is the default
("(%h/%p.lsof.%L")", the modified personal device cache file path
is:

/Homes/abe/LSOF/.lsof.vic

The LSOFPERSDCPATH environment variable is ignored when the lsof
process is setuid-root or when the real UID of the process is
root.

Lsof will not write to a modified personal device cache file path
if the lsof process doesn't surrender setgid permission. (See
-----the _LSOF_PERMISSIONS_THAT_AFFECT_DEVICE_CACHE_FILE_ACCESS_ section
-----for a list of implementations that normally don't surrender their
setgid permission.)

If, for example, you want to create a sub-directory of personal
device cache file paths by using the LSOFPERSDCPATH environment
variable to name it, and lsof doesn't surrender its setgid
-----permission, you will have to allow lsof to create device cache
-----files at the standard personal path and move them to your
-----subdirectory with shell commands.

-----The local system administrator may: disable this option when lsof
-----is built; change the name of the environment variable from
-----LSOFPERSDCPATH to something else; change the HASPERSDC format to

```

-----include_the_personal_path_component_in_another_place;or_exclude
 -----the_personal_path_component_entirely.--Consult_the_output_of_the
 -----D?_option_for_the_environment_variable's_name_and_the_HASPERSDC
 format_specification.

DIAGNOSTICS

Errors are identified with messages on the standard error file.

Lsof returns a one (1) if any error was detected, including the failure to locate **command** names, file names, Internet addresses or files, **login** names, NFS files, PIDs, PGIDs, or UIDs it was asked to list. If the -V option is specified, lsof will indicate the search items it failed to list.

It returns a zero (0) if no errors were detected and if it was able to list some information about all the specified search arguments.

When lsof cannot open access to /dev (or /devices) or one of its subdirectories, or get information on a file in them with stat(2), it issues a warning message and continues. That lsof will issue warning messages about inaccessible files in /dev (or /devices) is indicated in its **help** output - requested with the -h or >B -? options - with the message:

Inaccessible /dev warnings are enabled.

The warning message may be suppressed with the -w option. It may also have been suppressed by the system administrator when lsof was compiled by the setting of the WARNDVACCESS definition. In this case, the output from the **help** options will include the message:

Inaccessible /dev warnings are disabled.

Inaccessible device warning messages usually disappear after lsof has created a working device cache file.

EXAMPLES

For a more extensive set of examples, documented more fully, see the 00QUICKSTART file of the lsof distribution.

To list all open files, use:

```
lsof
```

To list all open Internet, x.25 (HP-UX), and UNIX domain files, use:

```
lsof -i -U
```

To list all open IPv4 network files in use by the process whose PID is 1234, use:

```
lsof -i 4 -a -p 1234
```

Presuming the UNIX dialect supports IPv6, to list only open IPv6 network files, use:

```
lsof -i 6
```

To list all files using any protocol on ports 513, 514, or 515 of host wonderland.cc.purdue.edu, use:

```
lsof -i @wonderland.cc.purdue.edu:513-515
```

To list all files using any protocol on any port of mace.cc.purdue.edu (cc.purdue.edu is the default domain), use:

```
lsof -i @mace
```

To list all open files for login name 'abe', or user ID 1234, or process 456, or process 123, or process 789, use:

```
lsof -p 456,123,789 -u 1234,abe
```

To list all open files on device /dev/hd4, use:

```
lsof /dev/hd4
```

To find the process that has /u/abe/foo open, use:

```
lsof /u/abe/foo
```

To send a SIGHUP to the processes that have /u/abe/bar open, use:

```
kill -HUP `lsof -t /u/abe/bar`
```

To find any open file, including an open UNIX domain socket file, with the name /dev/log, use:

```
lsof /dev/log
```

To find processes with open files on the NFS file system named

/nfs/mount/point whose server is inaccessible, and presuming your mount table supplies the device number for /nfs/mount/point, use:

```
lsuf -b /nfs/mount/point
```

To do the preceding search with warning messages suppressed, use:

```
lsuf -bw /nfs/mount/point
```

To ignore the device cache file, use:

```
lsuf -Di
```

To obtain PID and command name field output for each process, file descriptor, file device number, and file inode number for each file of each process, use:

```
lsuf -FpcfDi
```

To list the files at descriptors 1 and 3 of every process running the lsuf command for login ID 'abe' every 10 seconds, use:

```
lsuf -c lsuf -a -d 1 -d 3 -u abe -r10
```

To list the current working directory of processes running a command that is exactly four characters long and has an 'o' or 'O' in character three, use this regular expression form of the -c c option:

```
lsuf -c /^..o.$/i -a -d cwd
```

To find an IP version 4 socket file by its associated numeric dot-form address, use:

```
lsuf -i@128.210.15.17
```

To find an IP version 6 socket file (when the UNIX dialect supports IPv6) by its associated numeric colon-form address, use:

```
lsuf -i@[0:1:2:3:4:5:6:7]
```

To find an IP version 6 socket file (when the UNIX dialect supports IPv6) by an associated numeric colon-form address that has a run of zeroes in it - e.g., the loop-back address - use:

```
lsuf -i@[::1]
```

To obtain a repeat mode marker line that contains the current time, use:

```
lsuf -rm====%T====
```

To add spaces to the previous marker line, use:

```
lsuf -r "m====%T===="
```

BUGS

Since lsuf reads kernel memory in its search for open files, rapid changes in kernel memory may produce unpredictable results.

When a file has multiple record locks, the lock status character (following the file descriptor) is derived from a test of the first lock structure, not from any combination of the individual record locks that might be described by multiple lock structures.

~~~~~Lsuf can't search for files with restrictive access permissions by name unless it is installed with root set-UID permission. ~~~~~Otherwise it is limited to searching for files to which its user ~~~~~or its set-GID group (if any) has access permission. ~~~~~

~~~~~The display of the destination address of a raw socket (e.g., for ping) depends on the UNIX operating system. ~Some dialects store ~~~~~the destination address in the raw socket's protocol control block, some do not. ~~~~~

~~~~~Lsuf can't always represent Solaris device numbers in the same way that ls(1) does. ~For example, the major and minor device ~~~~~numbers that the lstat(2) and stat(2) functions report for the ~~~~~directory on which CD-ROM files are mounted (typically /cdrom) ~~~~~are not the same as the ones that it reports for the device on ~~~~~which CD-ROM files are mounted (typically /dev/sr0). ~~~~~(Lsuf ~~~~~reports the directory numbers.) ~~~~~

~~~~~The support for /proc file systems is available only for BSD and ~~~~~Tru64 UNIX dialects, Linux, and dialects derived from SYSV\_R4 ~~~~~e.g., FreeBSD, NetBSD, OpenBSD, Solaris, UnixWare. ~~~~~

~~~~~Some /proc file items - device number, inode number, and file ~~~~~size - are unavailable in some dialects. ~Searching for file ~~~~~a /proc file system may require that the full path name be ~~~~~specified. ~~~~~

~~~~~No text (txt) file descriptors are displayed for Linux processes. ~~~~~

```

-----All entries for files other than the current working directory,
-----the root directory, and numerical file descriptors are labeled
-----mem descriptors.

-----Lsof can't search for Tru64 UNIX named pipes by name, because
-----their kernel implementation of lstat(2) returns an improper
-----device number for a named pipe.

-----Lsof can't report fully or correctly on HP-UX 9.01, 10.20, and
-----11.00 locks because of insufficient access to kernel data or
-----errors in the kernel data. See the lsof FAQ (The FAQ section
-----gives its location.) for details.

-----The AIX SMT file type is a fabrication. It's made up for file
-----structures whose type (15) isn't defined in the AIX
-----/usr/include/sys/file.h header file. One way to create such file
-----structures is to run X clients with the DISPLAY variable set to
-----':0.0'.

-----The +|-f[cfgn] option is not supported under /proc-based Linux
-----lsof, because it doesn't read kernel structures from kernel
-----memory.

```

ENVIRONMENT

Lsof may access these environment variables.

LANG defines a language locale. See setlocale(3) for the names of other variables that can be used in place of LANG — e.g., LC_ALL, LC_TYPE, etc.

LSOFDEVCACHE

defines the path to a device cache file. See the DEVICE CACHE PATH FROM AN ENVIRONMENT VARIABLE section for more information.

LSOFPERSDCPATH

defines the middle component of a modified personal device cache file path. See the MODIFIED PERSONAL DEVICE CACHE PATH section for more information.

FAQ

Frequently-asked questions and their answers (an FAQ) are available in the 00FAQ file of the lsof distribution.

That file is also available via anonymous ftp from lsof.itap.purdue.edu at pub/tools/unix/lsof/FAQ. The URL is:

ftp://lsof.itap.purdue.edu/pub/tools/unix/lsof/FAQ

FILES

/dev/kmem
kernel virtual memory device

/dev/mem
physical memory device

/dev/swap
system paging device

.lsof_hostname

lsof's device cache file (The suffix, _hostname, is the first component of the host's name returned by gethostname(2).)

AUTHORS

Lsof was written by Victor A. Abell <abe@purdue.edu> of Purdue University. Many others have contributed to lsof. They're listed in the 00CREDITS file of the lsof distribution.

DISTRIBUTION

The latest distribution of lsof is available via anonymous ftp from the host lsof.itap.purdue.edu. You'll find the lsof distribution in the pub/tools/unix/lsof directory.

You can also use this URL:

ftp://lsof.itap.purdue.edu/pub/tools/unix/lsof

Lsof is also mirrored elsewhere. When you access lsof.itap.purdue.edu and change to its pub/tools/unix/lsof directory, you'll be given a list of some mirror sites. The pub/tools/unix/lsof directory also contains a more complete list in its mirrors file. Use mirrors with caution — not all mirrors always have the latest lsof revision.

Some pre-compiled Lsof executables are available on lsof.itap.purdue.edu, but their use is discouraged — it's better that you build your own from the sources. If you feel you must use a pre-compiled executable, please read the cautions that appear in the README files of the pub/tools/unix/lsof/binaries subdirectories and in the 00* files of the distribution.

More information on the lsof distribution can be found in its README.lsof.<version> file. If you intend to get the lsof distribution and build it, please read README.lsof.<version> and the other 00* files of the distribution before sending questions to the author.

SEE ALSO

Not all the following manual pages may exist in every UNIX dialect to which lsof has been ported.

access(2), **awk**(1), crash(1), fattach(3C), ff(1), fstat(8), fuser(1), gethostname(2), isprint(3), **kill**(1), localtime(3), lstat(2), modload(8), mount(8), netstat(1), ofiles(8L), perl(1), **ps**(1), readlink(2), setlocale(3), stat(2), strftime(3), time(2), uname(1).

COLOPHON

This page is part of the lsof (LiSt Open Files) project. Information about the project can be found at <http://people.freebsd.org/~abe/>. If you have a bug report for this manual page, send it to abe@purdue.edu. This page was obtained from the tarball lsof_4.91.src.tar fetched from <ftp://ftp.fu-berlin.de/pub/unix/tools/lsof/lsof.tar.gz> on 2024-06-14. If you discover any rendering problems in this HTML version of the page, or you believe there is a better or more up-to-date **source** for the page, or you have corrections or improvements to the information in this COLOPHON (which is not part of the original manual page), send a mail to man-pages@man7.org

Revision -4.91

LSOF(8)

3.6 objdump: Display Information From Object Files

NAME

objdump - display information from object files

SYNOPSIS

```
objdump [-a|--archive-headers]
[-b bfdname|--target=bfdname]
[-C|--demangle[=style] ]
[-d|--disassemble[=symbol]]
[-D|--disassemble-all]
[-z|--disassemble-zeroes]
[-EB|--EL|--endian={big | little }]
[-f|--file-headers]
[-F|--file-offsets]
--file-start-context]
[-g|--debugging]
[-e|--debugging-tags]
[-h|--section-headers|--headers]
[-i|--info]
[-j section|--section=section]
[-l|--line-numbers]
[-S|--source]
--source-comment[=text ]]
[-m machine|--architecture=machine]
[-M options|--disassembler-options=options]
[-p|--private-headers]
[-P options|--private=options]
[-r|--reloc]
[-R|--dynamic-reloc]
[-s|--full-contents]
[-Z|--decompress]
[-W{ILiaprmiFsoORtUuTgAck} ]
--dwarf[=rawline,=decodedline,=info,=abbrev,=pubnames,=aranges,=macro,=frames,=frames-interp,=str,=str-off]
[-WK|--dwarf=follow-links]
[-WN|--dwarf=no-follow-links]
[-wD|--dwarf=use-debuginfod]
[-wE|--dwarf=do-not-use-debuginfod]
[-L|--process-links]
--ctf=section]
--sframe=section]
[-G|--stabs]
[-t|--syms]
[-T|--dynamic-syms]
[-x|--all-headers]
[-w|--wide]
--start-address=address]
--stop-address=address]
--no-addresses]
--prefix-addresses]
--[no-]show-raw-insn]
--adjust-vma=offset]
--show-all-symbols]
--dwarf-depth=n]
--dwarf-start=n]
--ctf-parent=section]
--no-recurse-limit|--recurse-limit]
--special-syms]
--prefix=prefix]
--prefix-strip=level]
--insn-width=width]
--visualize-jumps[=color|extended-color|off]
--disassembler-color=[off|terminal|on|extended]
-U method] [--unicode=method]
```

```
[-V|--version]
[-H|--help]
objfile...
```

DESCRIPTION

objdump displays information about one or more object files. The options control what particular information to display. This information is mostly useful to programmers who are working on the compilation tools, as opposed to programmers who just want their program to compile and work.

objfile... are the object files to be examined. When you specify archives, objdump shows information on each of the member object files.

OPTIONS

The long and short forms of options, shown here as alternatives, are equivalent. At least one option from the list -a, -d, -D, -e, -f, -g, -G, -h, -H, -p, -P, -r, -R, -s, -S, -t, -T, -V, -x must be given.

-a
--archive-header
If any of the objfile files are archives, display the archive header information (in a format similar to ls -l). Besides the information you could list with ar tv, objdump -a shows the object file format of each archive member.

--adjust-vma=offset
When dumping information, first add offset to all the section addresses. This is useful if the section addresses do not correspond to the symbol table, which can happen when putting sections at particular addresses when using a format which can not represent section addresses, such as a.out.

-b bfdname
--target=bfdname
Specify that the object-code format for the object files is bfdname. This option may not be necessary; objdump can automatically recognize many formats.

For example,

```
objdump -b oasys -m vax -h fu.o
```

displays summary information from the section headers (-h) of fu.o, which is explicitly identified (-m) as a VAX object file in the format produced by Oasys compilers. You can list the formats available with the -i option.

-C
--demangle[=style]
Decode (demangle) low-level symbol names into user-level names. Besides removing any initial underscore prepended by the system, this makes C++ function names readable. Different compilers have different mangling styles. The optional demangling style argument can be used to choose an appropriate demangling style for your compiler.

--recurse-limit
--no-recurse-limit
--recursion-limit
--no-recursion-limit
Enables or disables a limit on the amount of recursion performed whilst demangling strings. Since the name mangling formats allow for an infinite level of recursion it is possible to create strings whose decoding will exhaust the amount of stack space available on the host machine, triggering a memory fault. The limit tries to prevent this from happening by restricting recursion to 2048 levels of nesting.

The default is for this limit to be enabled, but disabling it may be necessary in order to demangle truly complicated names. Note however that if the recursion limit is disabled then stack exhaustion is possible and any bug reports about such an event will be rejected.

-g
--debugging
Display debugging information. This attempts to parse STABS debugging format information stored in the file and print it out using a C like syntax. If no STABS debugging was found this option falls back on the -W option to print any DWARF information in the file.

-e
--debugging-tags
Like -g, but the information is generated in a format compatible with ctags tool.

-d
--disassemble
--disassemble=symbol

Display the assembler mnemonics for the machine instructions from the input file. This option only disassembles those sections which are expected to contain instructions. If the optional symbol argument is given, then display the assembler mnemonics starting at symbol. If symbol is a function name then disassembly will stop at the end of the function, otherwise it will stop when the next symbol is encountered. If there are no matches for symbol then nothing will be displayed.

Note if the `—dwarf=follow-links` option is enabled then any symbol tables in linked debug info files will be read in and used when disassembling.

`—D`

`—disassemble-all`

Like `—d`, but disassemble the contents of all non-empty non-bss sections, not just those expected to contain instructions. `—j` may be used to select specific sections.

This option also has a subtle effect on the disassembly of instructions in code sections. When option `—d` is in effect objdump will assume that any symbols present in a code section occur on the boundary between instructions and it will refuse to disassemble across such a boundary. When option `—D` is in effect however this assumption is suppressed. This means that it is possible for the output of `—d` and `—D` to differ if, for example, data is stored in code sections.

If the target is an ARM architecture this switch also has the effect of forcing the disassembler to decode pieces of data found in code sections as if they were instructions.

Note if the `—dwarf=follow-links` option is enabled then any symbol tables in linked debug info files will be read in and used when disassembling.

`—no-addresses`

When disassembling, don't print addresses on each line or for symbols and relocation offsets. In combination with `—no-show-raw-insn` this may be useful for comparing compiler output.

`—prefix-addresses`

When disassembling, print the complete address on each line. This is the older disassembly format.

`—EB`

`—EL`

`—endian={big|little}`
Specify the endianness of the object files. This only affects disassembly. This can be useful when disassembling a file format which does not describe endianness information, such as S-records.

`—f`

`—file-headers`

Display summary information from the overall header of each of the objfile files.

`—F`

`—file-offsets`

When disassembling sections, whenever a symbol is displayed, also display the file offset of the region of data that is about to be dumped. If zeroes are being skipped, then when disassembly resumes, tell the user how many zeroes were skipped and the file offset of the location from where the disassembly resumes. When dumping sections, display the file offset of the location from where the dump starts.

`—file-start-context`

Specify that when displaying interlisted source code/disassembly (assumes `—S`) from a file that has not yet been displayed, extend the context to the start of the file.

`—h`

`—section-headers`

`—headers`

Display summary information from the section headers of the object file.

File segments may be relocated to nonstandard addresses, for example by using the `—Ttext`, `—Tdata`, or `—Tbss` options to `ld`. However, some object file formats, such as `a.out`, do not store the starting address of the file segments. In those situations, although `ld` relocates the sections correctly, using `objdump -h` to list the file section headers cannot show the correct addresses. Instead, it shows the usual addresses, which are implicit for the target.

Note, in some cases it is possible for a section to have both the `—READONLY` and the `—NOREAD` attributes set. In such cases


```

-----the _NOREAD_ attribute takes precedence, but _objdump_ will
-----report both since the exact setting of the _flag_bits_ might be
-----important.

-----H
-----help
-----Print a summary of the options to _objdump_ and exit.

-----i
-----info
-----Display a list showing all architectures and object formats
-----available for specification with _-b_ or _-m_.

-----j _name
-----section=_name
-----Display information for section _name_. This option may be
-----specified multiple times.

-----L
-----process=_links
-----Display the contents of non-debug sections found in separate
-----debuginfo files that are linked to the main file. This
-----option automatically implies the _-WK_ option, and only
-----sections requested by other command line options will be
-----displayed.

-----l
-----line=_numbers
-----Label the display (using debugging information) with the
-----filename and source line numbers corresponding to the object
-----code or relocations shown. Only useful with _-d_, _-D_, or _-r_.

-----m _machine
-----architecture=_machine
-----Specify the architecture to use when disassembling object
-----files. This can be useful when disassembling object files
-----which do not describe architecture information, such as
-----S-relocs. You can list the available architectures with the
-----_i_ option.

-----For most architectures it is possible to supply an
-----architecture name and a machine name, separated by a colon.
-----For example _foo:bar_ would refer to the _bar_ machine type in
-----the _foo_ architecture. This can be helpful if _objdump_ has
-----been configured to support multiple architectures.

-----If the target is an ARM architecture then this switch has an
-----additional effect. It restricts the disassembly to only
-----those instructions supported by the architecture specified by
-----machine. If it is necessary to use this switch because the
-----input file does not contain any architecture information, but
-----it is also desired to disassemble all the instructions use
-----_marm_.

-----M _options
-----disassembler=_options
-----Pass target specific information to the disassembler. Only
-----supported on some targets. If it is necessary to specify
-----more than one disassembler option then multiple _-M_ options
-----can be used or can be placed together into a comma-separated
-----list.

-----For ARC, _dsp_ controls the printing of DSP instructions, _spfp_
-----selects the printing of FPX single precision FP instructions,
-----_dpfp_ selects the printing of FPX double precision FP
-----instructions, _quarkse.em_ selects the printing of special
-----QuarkSE-EM instructions, _fpuda_ selects the printing of double
-----precision assist instructions, _fpus_ selects the printing of
-----FPU single precision FP instructions, while _fpud_ selects the
-----printing of FPU double precision FP instructions.
-----Additionally, one can choose to have all the immediates
-----printed in hexadecimal using _hex_. By default, the short
-----immediates are printed using the decimal representation,
-----while the long immediate values are printed as hexadecimal.

-----cpu=... allows one to enforce a particular ISA when
-----disassembling instructions, overriding the _-m_ value or
-----whatever is in the ELF file. This might be useful to select
-----ARC_EM or HS ISA, because architecture is same for those and
-----disassembler relies on private ELF header data to decide if
-----code is for EM or HS. This option might be specified
-----multiple times -- only the latest value will be used. Valid
-----values are same as for the assembler _-mcpu=..._ option.

-----If the target is an ARM architecture then this switch can be
-----used to select which register name set is used during
-----disassembly. Specifying _-M reg-names=std_ (the default) will
-----select the register names as used in ARM's instruction set
-----documentation, but with register 13 called 'sp', register 14
-----called 'lr' and register 15 called 'pc'. Specifying _-M reg-
-----names=apcs_ will select the name set used by the ARM Procedure
-----Call Standard, whilst specifying _-M reg-names=raw_ will just

```

use `r` followed by the register number.

There are also two variants on the APCS register naming scheme enabled by `-M reg-names=atpcs` and `-M reg-names=special-atpcs` which use the ARM/Thumb Procedure Call Standard naming conventions. (Either with the normal register names or the special register names).

This option can also be used for ARM architectures to force the disassembler to interpret all instructions as Thumb instructions by using the switch `--disassembler-options=force-thumb`. This can be useful when attempting to disassemble thumb code produced by other compilers.

For AArch64 targets this switch can be used to set whether instructions are disassembled as the most general instruction using the `-M no-aliases` option or whether instruction notes should be generated as comments in the disassembly using `-M notes`.

For the x86, some of the options duplicate functions of the `-m` switch, but allow finer grained control.

"x86-64"

"i386"

"i8086"

Select disassembly for the given architecture.

"intel"

"att"

Select between intel syntax mode and AT&T syntax mode.

"amd64"

"intel64"

Select between AMD64 ISA and Intel64 ISA.

"intel-mnemonic"

"att-mnemonic"

Select between intel mnemonic mode and AT&T mnemonic mode. Note: "intel-mnemonic" implies "intel" and "att-mnemonic" implies "att".

"addr64"

"addr32"

"addr16"

"data32"

"data16"

Specify the default address size and operand size. These five options will be overridden if "x86-64", "i386" or "i8086" appear later in the option string.

"suffix"

When in AT&T mode and also for a limited set of instructions when in Intel mode, instructs the disassembler to print a mnemonic suffix even when the suffix could be inferred by the operands or, for certain instructions, the execution mode's defaults.

-----For PowerPC, the `-M argument raw` selects disassembly of hardware insns rather than aliases. For example, you will see "rlwinm" rather than "crlwi", and "addi" rather than "li". All of the `-m arguments` for `-gas` that select a CPU are supported. These are: 403, 405, 440, 464, 476, 601, 603, 604, 620, 7400, 7410, 7450, 7455, 750cl, 821, 850, 860, a2, booke, booke32, cell, com, e200z2, e200z4, e300, e500, e500mc, e500mc64, e500x2, e5500, e6500, efs, power4, power5, power6, power7, power8, power9, power10, ppc, ppc32, ppc64, ppc64bridge, ppcps, pwr, pwr2, pwr4, pwr5, pwr5x, pwr6, pwr7, pwr8, pwr9, pwr10, pwrX, titan, vle, and future. 32 and 64 modify the default or a prior CPU selection, disabling and enabling 64-bit insns respectively. In addition, `-altivec`, `-any`, `-lsp`, `-htm`, `-vsx`, `-spe` and `-spe2` add capabilities to a previous or later CPU selection. Any will disassemble any opcode known to binutils, but in cases where an opcode has two different meanings or different arguments, you may not see the disassembly you expect. If you disassemble without giving a CPU selection, a default will be chosen from information gleaned by BFD from the object files headers, but the result again may not be as you expect.

-----For MIPS, this option controls the printing of instruction mnemonic names and register names in disassembled instructions. Multiple selections from the following may be specified as a comma-separated string, and invalid options are ignored:

-----"no-aliases"

Print the 'raw' instruction mnemonic instead of some pseudo instruction mnemonic. I.e., print 'daddu_or' or 'instead_of_move', 'sll' instead of 'nop', etc.

```

-----"msa"
-----Disassemble_MSA_instructions.

-----"virt"
-----Disassemble_the_virtualization_ASE_instructions.

-----"xpa"
-----Disassemble_the_eXtended_Physical_Address_(XPA)_ASE
-----instructions.

-----"gpr-names=ABI"
-----Print_GPR_(general-purpose_register)_names_as_appropriate
-----for_the_specified_ABI. By default, _GPR_names_are
-----selected_according_to_the_ABI_of_the_binary_being
-----disassembled.

-----"fpr-names=ABI"
-----Print_FPR_(floating-point_register)_names_as_appropriate
-----for_the_specified_ABI. By default, _FPR_numbers_are
-----printed_rather_than_names.

-----"cp0-names=ARCH"
-----Print_CP0_(system_control_coprocessor;_coprocessor_0)
-----register_names_as_appropriate_for_the_CPU_or_architecture
-----specified_by_ARCH. By default, _CP0_register_names_are
-----selected_according_to_the_architecture_and_CPU_of_the
-----binary_being_disassembled.

-----"hwr-names=ARCH"
-----Print_HWR_(hardware_register,_used_by_the_"rdhwr"
-----instruction)_names_as_appropriate_for_the_CPU_or
-----architecture_specified_by_ARCH. By default, _HWR_names
-----are_selected_according_to_the_architecture_and_CPU_of_the
-----binary_being_disassembled.

-----"reg-names=ABI"
-----Print_GPR_and_FPR_names_as_appropriate_for_the_selected
-----ABI.

-----"reg-names=ARCH"
-----Print_CPU-specific_register_names_(CP0_register_and_HWR
-----names)_as_appropriate_for_the_selected_CPU_or
-----architecture.

-----For any of the options listed above, _ABI or _ARCH may be
-----specified as numeric to have numbers printed rather than
-----names, for the selected types of registers. You can list the
-----available values of _ABI and _ARCH using the --help option.

-----For VAX, you can specify function entry addresses with -M
-----entry:0xf00ba. You can use this multiple times to properly
-----disassemble VAX binary files that don't contain symbol tables
----- (like ROM dumps). In these cases, the function entry mask
----- would otherwise be decoded as VAX instructions, which would
----- probably lead the rest of the function being wrongly
----- disassembled.

-p
--private=headers
Print information that is specific to the object file format.
The exact information printed depends upon the object file
format. For some object file formats, no additional
information is printed.

-P options
--private=options
Print information that is specific to the object file format.
The argument options is a comma separated list that depends
on the format (the lists of options is displayed with the
help).

For XCOFF, the available options are:

"header"
"aout"
"sections"
"syms"
"relocs"
"lineno,"
"loader"
"except"
"typchk"
"traceback"
"toc"
"ldinfo"

For PE, the available options are:

"header"
"sections"

Not all object formats support this option. In particular

```

the ELF format does not use it.

-r
—reloc
 Print the relocation entries of the file. If used with **-d** or **-D**, the relocations are printed interspersed with the disassembly.

-R
—dynamic-reloc
 Print the dynamic relocation entries of the file. This is only meaningful **for** dynamic objects, such as certain types of shared libraries. As **for** **-r**, if used with **-d** or **-D**, the relocations are printed interspersed with the disassembly.

-s
—full-contents
 Display the full contents of sections, often used in combination with **-j** to request specific sections. By default all non-empty non-bss sections are displayed. By default any compressed section will be displayed in its compressed form. In order to see the contents in a decompressed form add the **-Z** option to the **command** line.

-S
—source
 Display **source** code intermixed with disassembly, if possible. Implies **-d**.

—show-all-symbols
 When disassembling, show all the symbols that match a given address, not just the first one.

—source-comment[=txt]
 Like the **-S** option, but all **source** code lines are displayed with a prefix of **txt**. Typically **txt** will be a comment string which can be used to distinguish the assembler code from the **source** code. If **txt** is not provided then a default string of **"#_"** (hash followed by a space), will be used.

—prefix=prefix
 Specify prefix to add to the absolute paths when used with **-S**.

—prefix-strip=level
 Indicate how many initial directory names to strip off the hardwired absolute paths. It has no effect without **—prefix=prefix**.

—show-raw-insn
 When disassembling instructions, print the instruction in hex as well as in symbolic form. This is the default except when **—prefix-addresses** is used.

—no-show-raw-insn
 When disassembling instructions, **do** not print the instruction bytes. This is the default when **—prefix-addresses** is used.

—insn-width=width
 Display width bytes on a single line when disassembling instructions.

—visualize-jumps[=color|extended-color|off]
 Visualize jumps that stay inside a **function** by drawing ASCII art between the start and target addresses. The optional **=color** argument adds color to the output using simple terminal colors. Alternatively the **=extended-color** argument will add color using 8bit colors, but these might not work on all terminals.

If it is necessary to disable the **visualize-jumps** option after it has previously been enabled then use **visualize-jumps=off**.

—disassembler-color=off
—disassembler-color=terminal
—disassembler-color=on|color|colour
—disassembler-color=extended|extended-color|extended-colour
 Enables or disables the use of colored syntax highlighting in disassembly output. The default behaviour is determined via a configure time option. Note, not all architectures support colored syntax highlighting, and depending upon the terminal used, colored output may not actually be legible.

The **on** argument adds colors using simple terminal colors.

The **terminal** argument does the same, but only if the output device is a terminal.

The **extended-color** argument is similar to the **on** argument, but it uses 8-bit colors. These may not work on all terminals.

The `off` argument disables colored disassembly.

`-W[LIaprmfFsoORtUuTgAckK]`

`—dwarf[=rawline,=decodedline,=info,=abbrev,=pubnames,=aranges,=macro,=frames,=frames-interp,=str,=str-offsets,=loc]`

Displays the contents of the DWARF debug sections in the file, if any are present. Compressed debug sections are automatically decompressed (temporarily) before they are displayed. If one or more of the optional letters or words follows the switch **then** only those **type(s)** of data will be dumped. The letters and words refer to the following information:

- "a"
"=abbrev"
Displays the contents of the `.debug-abbrev` section.
- "A"
"=addr"
Displays the contents of the `.debug-addr` section.
- "c"
"=cu-index"
Displays the contents of the `.debug-cu-index` and/or `.debug-tu-index` sections.
- "f"
"=frames"
Display the raw contents of a `.debug-frame` section.
- "F"
"=frames-interp"
Display the interpreted contents of a `.debug-frame` section.
- "g"
"=gdb-index"
Displays the contents of the `.gdb-index` and/or `.debug-names` sections.
- "i"
"=info"
Displays the contents of the `.debug-info` section. Note: the output from this option can also be restricted by the use of the `—dwarf-depth` and `—dwarf-start` options.
- "k"
"=links"
Displays the contents of the `.gnu-debuglink`, `.gnu-debugaltlink` and `.debug-sup` sections, if any of them are present. Also displays any links to separate dwarf object files (dwo), if they are specified by the `DW_AT_GNU_dwo_name` or `DW_AT_dwo_name` attributes in the `.debug-info` section.
- "K"
"=follow-links"
Display the contents of any selected debug sections that are found in linked, separate debug info file(s). This can result in multiple versions of the same debug section being displayed if it exists in more than one file.

In addition, when displaying DWARF attributes, if a form is found that references the separate debug info file, **then** the referenced contents will also be displayed.

Note – in some distributions this option is enabled by default. It can be disabled via the `N` debug option. The default can be chosen when configuring the binutils via the `—enable-follow-debug-links=yes` or `—enable-follow-debug-links=no` options. If these are not used **then** the default is to **enable** the following of debug links.

Note – if support for the debuginfod protocol was enabled when the binutils were built **then** this option will also include an attempt to contact any debuginfod servers mentioned in the `DEBUGINFOD_URLS` environment variable. This could take some time to resolve. This behaviour can be disabled via the `=do-not-use-debuginfod` debug option.

- "N"
"=no-follow-links"
Disables the following of links to separate debug info files.
- "D"
"=use-debuginfod"
Enables contacting debuginfod servers if there is a need to follow debug links. This is the default behaviour.
- "E"

```

"=do-not-use-debuginfod"
    Disables contacting debuginfod servers when there is a
    need to follow debug links.

"I"
"=rawline"
    Displays the contents of the .debug_line section in a raw
    format.

"L"
"=decodedline"
    Displays the interpreted contents of the .debug_line
    section.

"m"
"=macro"
    Displays the contents of the .debug-macro and/or
    .debug-macinfo sections.

"o"
"=loc"
    Displays the contents of the .debug-loc and/or
    .debug-loclists sections.

"O"
"=str-offsets"
    Displays the contents of the .debug-str-offsets section.

"p"
"=pubnames"
    Displays the contents of the .debug-pubnames and/or
    .debug-gnu-pubnames sections.

"r"
"=aranges"
    Displays the contents of the .debug-aranges section.

"R"
"=Ranges"
    Displays the contents of the .debug-ranges and/or
    .debug-rnglists sections.

"s"
"=str"
    Displays the contents of the .debug-str, .debug_line-str
    and/or .debug-str-offsets sections.

"t"
"=pubtype"
    Displays the contents of the .debug-pubtypes and/or
    .debug-gnu-pubtypes sections.

"T"
"=trace-aranges"
    Displays the contents of the .trace-aranges section.

"u"
"=trace-abbrev"
    Displays the contents of the .trace-abbrev section.

"U"
"=trace-info"
    Displays the contents of the .trace-info section.

Note: displaying the contents of .debug-static-funcs,
.debug-static-vars and debug-weaknames sections is not
currently supported.

--dwarf-depth=n
    Limit the dump of the ".debug_info" section to n children.
    This is only useful with --debug-dump=info. The default is
    to print all DIEs; the special value 0 for n will also have
    this effect.

    With a non-zero value for n, DIEs at or deeper than n levels
    will not be printed. The range for n is zero-based.

--dwarf-start=n
    Print only DIEs beginning with the DIE numbered n. This is
    only useful with --debug-dump=info.

    If specified, this option will suppress printing of any
    header information and all DIEs before the DIE numbered n.
    Only siblings and children of the specified DIE will be
    printed.

    This can be used in conjunction with --dwarf-depth.

--dwarf-check
    Enable additional checks for consistency of Dwarf
    information.

```

```

--ctf[=section]
    Display the contents of the specified CTF section. CTF
    sections themselves contain many subsections, all of which
    are displayed in order.

    By default, display the name of the section named .ctf, which
    is the name emitted by ld.

--ctf-parent=member
    If the CTF section contains ambiguously-defined types, it
    will consist of an archive of many CTF dictionaries, all
    inheriting from one dictionary containing unambiguous types.
    This member is by default named .ctf, like the section
    containing it, but it is possible to change this name using
    the "ctf.link.set_memb_name_changer" function at link time.
    When looking at CTF archives that have been created by a
    linker that uses the name changer to rename the parent
    archive member, --ctf-parent can be used to specify the name
    used for the parent.

--sframe[=section]
    Display the contents of the specified SFrame section.

    By default, display the name of the section named .sframe,
    which is the name emitted by ld.

-G
--stabs
    Display the full contents of any sections requested. Display
    the contents of the .stab and .stab.index and .stab.excl
    sections from an ELF file. This is only useful on systems
    (such as Solaris 2.0) in which ".stab" debugging symbol-table
    entries are carried in an ELF section. In most other file
    formats, debugging symbol-table entries are interleaved with
    linkage symbols, and are visible in the --syms output.

--start-address=address
    Start displaying data at the specified address. This affects
    the output of the -d, -r and -s options.

--stop-address=address
    Stop displaying data at the specified address. This affects
    the output of the -d, -r and -s options.

-t
--syms
    Print the symbol table entries of the file. This is similar
    to the information provided by the nm program, although the
    display format is different. The format of the output
    depends upon the format of the file being dumped, but there
    are two main types. One looks like this:

        [ 4](sec 3)(fl 0x00)(ty 0)(scl 3) (nx 1) 0x00000000 .bss
        [ 6](sec 1)(fl 0x00)(ty 0)(scl 2) (nx 0) 0x00000000 fred

    where the number inside the square brackets is the number of
    the entry in the symbol table, the sec number is the section
    number, the fl value are the symbol's flag-bits, the ty
    number is the symbol's type, the scl number is the symbol's
    storage class and the nx value is the number of auxiliary
    entries associated with the symbol. The last two fields are
    the symbol's value and its name.

    The other common output format, usually seen with ELF based
    files, looks like this:

        00000000 l d .bss 00000000 .bss
        00000000 g .text 00000000 fred

    Here the first number is the symbol's value (sometimes
    referred to as its address). The next field is actually a
    set of characters and spaces indicating the flag-bits that
    are set on the symbol. These characters are described below.
    Next is the section with which the symbol is associated or
    ABS if the section is absolute (ie not connected with any
    section), or *UND* if the section is referenced in the file
    being dumped, but not defined there.

    After the section name comes another field, a number, which
    for common symbols is the alignment and for other symbols is
    the size. Finally the symbol's name is displayed.

    The flag characters are divided into 7 groups as follows:

    "l"
    "g"
    "u"
    "!" The symbol is a local (l), global (g), unique global (u),
    neither global nor local (a space) or both global and
    local (!). A symbol can be neither local or global for a
    variety of reasons, e.g., because it is used for
    debugging, but it is probably an indication of a bug if

```

```

it is ever both local and global. Unique global symbols
are a GNU extension to the standard set of ELF symbol
bindings. For such a symbol the dynamic linker will make
sure that in the entire process there is just one symbol
with this name and type in use.

"w" The symbol is weak (w) or strong (a space).

"C" The symbol denotes a constructor (C) or an ordinary
symbol (a space).

"W" The symbol is a warning (W) or a normal symbol (a space).
A warning symbol's name is a message to be displayed if
the symbol following the warning symbol is ever
referenced.

-----"I"
-----"i" The symbol is an indirect reference to another symbol
-----      (I), a function to be evaluated during reloc processing
-----      (i) or a normal symbol (a space).

-----"d"
-----"D" The symbol is a debugging symbol (d) or a dynamic symbol
-----      (D) or a normal symbol (a space).

-----"p"
-----"f"
-----"O" The symbol is the name of a function (F) or a file (f) or
-----      an object (O) or just a normal symbol (a space).

-----T
-----dynamic-syms
-----Print the dynamic symbol table entries of the file. This is
-----only meaningful for dynamic objects, such as certain types of
-----shared libraries. This is similar to the information
-----provided by the nm program when given the -D (dynamic)
-----option.

-----The output format is similar to that produced by the --syms
-----option, except that an extra field is inserted before the
-----symbol's name, giving the version information associated with
-----the symbol. If the version is the default version to be used
-----when resolving unversioned references to the symbol then it's
-----displayed as is, otherwise it's put into parentheses.

--special-syms
When displaying symbols include those which the target
considers to be special in some way and which would not
normally be of interest to the user.

-U [d|i|l|e|x|h]
--unicode=[default|invalid|locale|escape|hex|highlight]
Controls the display of UTF-8 encoded multibyte characters in
strings. The default (--unicode=default) is to give them no
special treatment. The --unicode=locale option displays the
sequence in the current locale, which may or may not support
them. The options --unicode=hex and --unicode=invalid
display them as hex byte sequences enclosed by either angle
brackets or curly braces.

The --unicode=escape option displays them as escape sequences
(\uxxxx) and the --unicode=highlight option displays them as
escape sequences highlighted in red (if supported by the
output device). The colouring is intended to draw attention
to the presence of unicode sequences where they might not be
expected.

-V
--version
Print the version number of objdump and exit.

-x
--all-headers
Display all available header information, including the
symbol table and relocation entries. Using -x is equivalent
to specifying all of -a -f -h -p -r -t.

-w
--wide
Format some lines for output devices that have more than 80
columns. Also do not truncate symbol names when they are
displayed.

-z
--disassemble-zeroes
Normally the disassembly output will skip blocks of zeroes.
This option directs the disassembler to disassemble those
blocks, just like any other data.

-Z
--decompress
The -Z option is meant to be used in conjunction with the -s

```


option. It instructs objdump to decompress any compressed sections before displaying their contents.

@file
Read **command**-line options from file. The options **read** are inserted in place of the original **@file** option. If file does not exist, or cannot be **read**, **then** the option will be treated literally, and not removed.

Options in file are separated by whitespace. A whitespace character may be included in an option by surrounding the entire option in either single or double quotes. Any character (including a backslash) may be included by prefixing the character to be included with a backslash. The file may itself contain additional **@file** options; any such options will be processed recursively.

SEE ALSO **nm(1)**, **readelf(1)**, and the Info entries for **binutils**.

COPYRIGHT
Copyright (c) 1991–2024 Free Software Foundation, Inc.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with no Invariant Sections, with no Front-Cover Texts, and with no Back-Cover Texts. A copy of the license is included in the section entitled "GNU Free Documentation License".

COLOPHON
This page is part of the binutils (a collection of tools for working with executable binaries) project. Information about the project can be found at <http://www.gnu.org/software/binutils/>. If you have a bug report for this manual page, see http://sourceware.org/bugzilla/enter_bug.cgi?product=binutils. This page was obtained from the tarball **binutils-2.42.tar.gz** fetched from <https://ftp.gnu.org/gnu/binutils/> on 2024-06-14. If you discover any rendering problems in this HTML version of the page, or you believe there is a better or more up-to-date source for the page, or you have corrections or improvements to the information in this COLOPHON (which is not part of the original manual page), send a mail to man-pages@man7.org

binutils-2.42 2024-06-14 OBJDUMP(1)

3.7 readelf: Display Information On ELF Files

NAME
readelf — display information about ELF files

SYNOPSIS
readelf [-a|--all]
 [-h|--file-header]
 [-l|--program-headers|--segments]
 [-S|--section-headers|--sections]
 [-g|--section-groups]
 [-t|--section-details]
 [-e|--headers]
 [-s|--syms|--symbols]
 [--dyn-syms|--lto-syms]
 [--sym-base={0|8|10|16}]
 [--demangle=style|--no-demangle]
 [--quiet]
 [--recurse-limit|--no-recurse-limit]
 [-U method|--unicode=method]
 [-X|--extra-sym-info|--no-extra-sym-info]
 [-n|--notes]
 [-r|--relocs]
 [-u|--unwind]
 [-d|--dynamic]
 [-V|--version-info]
 [-A|--arch-specific]
 [-D|--use-dynamic]
 [-L|--lint|--enable-checks]
 [-x <number or name>|--hex-dump=<number or name>]
 [-p <number or name>|--string-dump=<number or name>]
 [-R <number or name>|--relocated-dump=<number or name>]
 [-z|--decompress]
 [-c|--archive-index]
 [-w [LIaprmFFsoORtUuTgAck]]
 --debug-dump[={rawline,=decodedline,=info,=abbrev,=pubnames,=ranges,=macro,=frames,=frames-interp,=str,=st...]
 [-wK|--debug-dump=follow-links]
 [-wN|--debug-dump=no-follow-links]
 [-wD|--debug-dump=use-debuginfod]
 [-wE|--debug-dump=do-not-use-debuginfod]
 [-P|--process-links]
 --dwarf-depth=n
 --dwarf-start=n
 --ctf=section
 --ctf-parent=section
 --ctf-symbols=section]


```

-----symbol sizes in octal. --sym-base=10 will always give the
-----symbol sizes in decimal. --sym-base=16 will always give the
-----symbol sizes in hexadecimal with a_0x_ prefix.

-----C
-----demangle[=style]
-----Decode_(demangle)_low-level_symbol_names_into_user-level
-----names. --This makes C++ function names readable. --Different
-----compilers have different mangling styles. --The optional
-----demangling_style argument can be used to choose an
-----appropriate demangling style for your compiler.

-----no-demangle
-----Do not demangle low-level symbol names. --This is the default.

-----recurse-limit
-----no-recurse-limit
-----recursion-limit
-----no-recursion-limit
-----Enables or disables a limit on the amount of recursion
-----performed whilst demangling strings. --Since the name mangling
-----formats allow for an infinite level of recursion it is
-----possible to create strings whose decoding will exhaust the
-----amount of stack space available on the host machine,
-----triggering a memory fault. --The limit tries to prevent this
-----from happening by restricting recursion to 2048 levels of
-----nesting.

-----The default is for this limit to be enabled, but disabling it
-----may be necessary in order to demangle truly complicated
-----names. --Note however that if the recursion limit is disabled
-----then stack exhaustion is possible and any bug reports about
-----such an event will be rejected.

-----U_[d|i|l|e|x|h]
-----unicode=[default|invalid|locale|escape|hex|highlight]
-----Controls the display of non-ASCII characters in identifier
-----names. --The default (--unicode=locale or --unicode=default)
-----is to treat them as multibyte characters and display them in
-----the current locale. --All other versions of this option treat
-----the bytes as UTF-8 encoded values and attempt to interpret
-----them. --If they cannot be interpreted or if the
-----unicode=invalid option is used then they are displayed as a
-----sequence of hex bytes, enclosed in curly parentheses
-----characters.

-----Using the --unicode=escape option will display the characters
-----as unicode escape sequences (\uxxxx). --Using the
-----unicode=hex will display the characters as hex byte
-----sequences enclosed between angle brackets.

-----Using the --unicode=highlight will display the characters as
-----unicode escape sequences but it will also highlight them in
-----red, assuming that colouring is supported by the output
-----device. --The colouring is intended to draw attention to the
-----presence of unicode sequences when they might not be
-----expected.

-----X
-----extra-sym-info
-----When displaying details of symbols, include extra information
-----not normally presented. --Currently this just adds the name of
-----the section referenced by the symbol's index field, if there
-----is one. In the future more information may be displayed when
-----this option is enabled.

-----Enabling this option effectively enables the --wide option as
-----well, at least when displaying symbol information.

-----no-extra-sym-info
-----Disables the effect of the --extra-sym-info option. This is
-----the default.

-----e
-----headers
-----Display all the headers in the file. Equivalent to -h -l -S.

-----n
-----notes
-----Displays the contents of the NOTE segments and/or sections,
-----if any.

-----r
-----relocs
-----Displays the contents of the file's relocation section, if it
-----has one.

-----u
-----unwind
-----Displays the contents of the file's unwind section, if it has
-----one. Only the unwind sections for IA64 ELF files, as well as
-----ARM unwind tables ("ARM.exidx" / "ARM.extab") are currently

```

```

supported. If support is not yet implemented for your
architecture you could try dumping the contents of the
.debug_frames section using the --debug-dump=frames or
--debug-dump=frames=interp options.

-d
--dynamic
    Displays the contents of the file's .dynamic section, if it
    has one.

-----V
-----version-info
    Displays the contents of the .version sections in the file, if
    they exist.

-----A
-----arch-specific
    Displays architecture-specific information in the file, if
    there is any.

-----D
-----use-dynamic
    When displaying symbols, this option makes readelf use the
    symbol hash tables in the file's .dynamic section, rather than
    the symbol table sections.

    When displaying relocations, this option makes readelf
    display the dynamic relocations rather than the static
    relocations.

-L
--lint
--enable-checks
    Displays warning messages about possible problems with the
    file(s) being examined. If used on its own then all of the
    contents of the file(s) will be examined. If used with one
    of the dumping options then the warning messages will only be
    produced for the things being displayed.

-x <number or name>
--hex-dump=<number or name>
    Displays the contents of the indicated section as a
    hexadecimal bytes. A number identifies a particular section
    by index in the section table; any other string identifies
    all sections with that name in the object file.

-R <number or name>
--relocated-dump=<number or name>
    Displays the contents of the indicated section as a
    hexadecimal bytes. A number identifies a particular section
    by index in the section table; any other string identifies
    all sections with that name in the object file. The contents
    of the section will be relocated before they are displayed.

-p <number or name>
--string-dump=<number or name>
    Displays the contents of the indicated section as printable
    strings. A number identifies a particular section by index
    in the section table; any other string identifies all
    sections with that name in the object file.

-z
--decompress
    Requests that the section(s) being dumped by x, R or p
    options are decompressed before being displayed. If the
    section(s) are not compressed then they are displayed as is.

-c
--archive-index
    Displays the file symbol index information contained in the
    header part of binary archives. Performs the same function
    as the t command to ar, but without using the BFD library.

-w[LIaprmfFsOoRtUuTgAckK]
--debug-dump[=rawline,=decodedline,=info,=abbrev,=pubnames,=anges,=macro,=frames,=frames=interp,=str,=str=offsets]
    Displays the contents of the DWARF debug sections in the
    file, if any are present. Compressed debug sections are
    automatically decompressed (temporarily) before they are
    displayed. If one or more of the optional letters or words
    follows the switch then only those type(s) of data will be
    dumped. The letters and words refer to the following
    information:

    "a"
    "=abbrev"
        Displays the contents of the .debug_abbrev section.

    "A"
    "=addr"
        Displays the contents of the .debug_addr section.

    "c"

```

```

"=cu-index"
    Displays the contents of the .debug-cu-index and/or
    .debug-tu-index sections.

"f"
"=frames"
    Display the raw contents of a .debug-frame section.

"F"
"=frames-interp"
    Display the interpreted contents of a .debug-frame
    section.

"g"
"=gdb-index"
    Displays the contents of the .gdb-index and/or
    .debug-names sections.

"i"
"=info"
    Displays the contents of the .debug-info section. Note:
    the output from this option can also be restricted by the
    use of the --dwarf-depth and --dwarf-start options.

"k"
"=links"
    Displays the contents of the .gnu_debuglink,
    .gnu_debugaltlink and .debug_sup sections, if any of them
    are present. Also displays any links to separate dwarf
    object files (dwo), if they are specified by the
    DW_AT_GNU_dwo_name or DW_AT_dwo_name attributes in the
    .debug-info section.

"K"
"=follow-links"
    Display the contents of any selected debug sections that
    are found in linked, separate debug info file(s). This
    can result in multiple versions of the same debug section
    being displayed if it exists in more than one file.

    In addition, when displaying DWARF attributes, if a form
    is found that references the separate debug info file,
    then the referenced contents will also be displayed.

    Note - in some distributions this option is enabled by
    default. It can be disabled via the N debug option. The
    default can be chosen when configuring the binutils via
    the --enable-follow-debug-links=yes or
    --enable-follow-debug-links=no options. If these are not
    used then the default is to enable the following of debug
    links.

    Note - if support for the debuginfod protocol was enabled
    when the binutils were built then this option will also
    include an attempt to contact any debuginfod servers
    mentioned in the DEBUGINFOD_URLS environment variable.
    This could take some time to resolve. This behaviour can
    be disabled via the =do-not-use-debuginfod debug option.

"N"
"=no-follow-links"
    Disables the following of links to separate debug info
    files.

"D"
"=use-debuginfod"
    Enables contacting debuginfod servers if there is a need
    to follow debug links. This is the default behaviour.

"E"
"=do-not-use-debuginfod"
    Disables contacting debuginfod servers when there is a
    need to follow debug links.

"l"
"=rawline"
    Displays the contents of the .debug_line section in a raw
    format.

"L"
"=decodedline"
    Displays the interpreted contents of the .debug_line
    section.

"m"
"=macro"
    Displays the contents of the .debug-macro and/or
    .debug-macinfo sections.

"o"
"=loc"
    Displays the contents of the .debug-loc and/or

```

```

    .debug_loclists sections.

"O"
"=str-offsets"
    Displays the contents of the .debug_str_offsets section.

"P"
"=pubnames"
    Displays the contents of the .debug_pubnames and/or
    .debug_gnu_pubnames sections.

"r"
"=ranges"
    Displays the contents of the .debug_ranges section.

"R"
"=Ranges"
    Displays the contents of the .debug_ranges and/or
    .debug_rnglists sections.

"s"
"=str"
    Displays the contents of the .debug_str, .debug_line_str
    and/or .debug_str_offsets sections.

"t"
"=pubtype"
    Displays the contents of the .debug_pubtypes and/or
    .debug_gnu_pubtypes sections.

"T"
"=trace_aranges"
    Displays the contents of the .trace_aranges section.

"u"
"=trace_abbrev"
    Displays the contents of the .trace_abbrev section.

"U"
"=trace_info"
    Displays the contents of the .trace_info section.

Note: displaying the contents of .debug_static_funcs,
.debug_static_vars and debug_weaknames sections is not
currently supported.

--dwarf-depth=n
    Limit the dump of the ".debug_info" section to n children.
    This is only useful with --debug-dump=info. The default is
    to print all DIEs; the special value 0 for n will also have
    this effect.

    With a non-zero value for n, DIEs at or deeper than n levels
    will not be printed. The range for n is zero-based.

--dwarf-start=n
    Print only DIEs beginning with the DIE numbered n. This is
    only useful with --debug-dump=info.

    If specified, this option will suppress printing of any
    header information and all DIEs before the DIE numbered n.
    Only siblings and children of the specified DIE will be
    printed.

    This can be used in conjunction with --dwarf-depth.

-P
--process-links
    Display the contents of non-debug sections found in separate
    debuginfo files that are linked to the main file. This
    option automatically implies the -wK option, and only
    sections requested by other command line options will be
    displayed.

--ctf[=section]
    Display the contents of the specified CTF section. CTF
    sections themselves contain many subsections, all of which
    are displayed in order.

    By default, display the name of the section named .ctf, which
    is the name emitted by ld.

--ctf-parent=member
    If the CTF section contains ambiguously-defined types, it
    will consist of an archive of many CTF dictionaries, all
    inheriting from one dictionary containing unambiguous types.
    This member is by default named .ctf, like the section
    containing it, but it is possible to change this name using
    the "ctf_link_set_memb_name_changer" function at link time.
    When looking at CTF archives that have been created by a
    linker that uses the name changer to rename the parent
    archive member, --ctf-parent can be used to specify the name

```

```

        used for the parent.

--ctf=symbols=section
--ctf=strings=section
    Specify the name of another section from which the CTF file
    can inherit strings and symbols. By default, the ".symtab"
    and its linked string table are used.

    If either of --ctf=symbols or --ctf=strings is specified, the
    other must be specified as well.

-I
--histogram
    Display a histogram of bucket list lengths when displaying
    the contents of the symbol tables.

-v
--version
    Display the version number of readelf.

-W
--wide
    Don't break output lines to fit into 80 columns. By default
    readelf breaks section header and segment listing lines for
    64-bit ELF files, so that they fit into 80 columns. This
    option causes readelf to print each section header resp. each
    segment one a single line, which is far more readable on
    terminals wider than 80 columns.

-T
--silent-truncation
    Normally when readelf is displaying a symbol name, and it has
    to truncate the name to fit into an 80 column display, it
    will add a suffix of "[...]" to the name. This command line
    option disables this behaviour, allowing 5 more characters of
    the name to be displayed and restoring the old behaviour of
    readelf (prior to release 2.35).

-H
--help
    Display the command-line options understood by readelf.

@file
    Read command-line options from file. The options read are
    inserted in place of the original @file option. If file does
    not exist, or cannot be read, then the option will be treated
    literally, and not removed.

    Options in file are separated by whitespace. A whitespace
    character may be included in an option by surrounding the
    entire option in either single or double quotes. Any
    character (including a backslash) may be included by
    prefixing the character to be included with a backslash. The
    file may itself contain additional @file options; any such
    options will be processed recursively.
SEE_ALSO
    objdump(1), and the Info entries for binutils.
COPYRIGHT
    Copyright (c) 1991-2024 Free Software Foundation, Inc.

    Permission is granted to copy, distribute and/or modify this
    document under the terms of the GNU Free Documentation License,
    Version 1.3 or any later version published by the Free Software
    Foundation; with no Invariant Sections, with no Front-Cover
    Texts, and with no Back-Cover Texts. A copy of the license is
    included in the section entitled "GNU Free Documentation
    License".
COLOPHON
    This page is part of the binutils (a collection of tools for
    working with executable binaries) project. Information about the
    project can be found at http://www.gnu.org/software/binutils/.
    If you have a bug report for this manual page, see
    http://sourceware.org/bugzilla/enter\_bug.cgi?product=binutils.
    This page was obtained from the tarball binutils-2.42.tar.gz
    fetched from https://ftp.gnu.org/gnu/binutils/ on 2024-06-14.
    If you discover any rendering problems in this HTML version of
    the page, or you believe there is a better or more up-to-date
    source for the page, or you have corrections or improvements to
    the information in this COLOPHON (which is not part of the
    original manual page), send a mail to man-pages@man7.org

binutils-2.42-----2024-06-14-----READ ELF(1)

```

3.8 nm: List Symbols From Object Files

```

NAME
    nm - list symbols from object files
SYNOPSIS

```

```

nm [-A|-o|--print-file-name]
   [-a|--debug-syms]
   [-B|--format=bsd]
   [-C|--demangle[=style]]
   [-D|--dynamic]
   [-fformat|--format=format]
   [-g|--extern-only]
   [-h|--help]
   [--ifunc-chars=CHARS]
   [-j|--format=just-symbols]
   [-l|--line-numbers] [--inlines]
   [-n|-v|--numeric-sort]
   [-P|--portability]
   [-p|--no-sort]
   [-r|--reverse-sort]
   [-S|--print-size]
   [-s|--print-arnmap]
   [-t radix|--radix=radix]
   [-u|--undefined-only]
   [-U|--defined-only]
   [-V|--version]
   [-W|--no-weak]
   [-X 32_64]
   [--no-demangle]
   [--no-recurse-limit|--recurse-limit]]
   [--plugin name]
   [--size-sort]
   [--special-syms]
   [--synthetic]
   [--target=bfdname]
   [--unicode=method]
   [--with-symbol-versions]
   [--without-symbol-versions]
   [objfile ...]

```

DESCRIPTION

GNU nm lists the symbols from object files objfile.... If no object files are listed as arguments, nm assumes the file a.out.

For each symbol, nm shows:

- * The symbol value, in the radix selected by options (see below), or hexadecimal by default.
- * The symbol **type**. At least the following types are used; others are, as well, depending on the object file format. If lowercase, the symbol is usually **local**; if uppercase, the symbol is global (external). There are however a few lowercase symbols that are shown for special global symbols ("u", "v" and "w").

"A" The symbol's value is absolute, and will not be changed by further linking.

"B"
"b" The symbol is in the BSS data section. This section typically contains zero-initialized or uninitialized data, although the exact behavior is system-dependent.

"C"
"c" The symbol is common. Common symbols are uninitialized data. When linking, multiple common symbols may appear with the same name. If the symbol is defined anywhere, the common symbols are treated as undefined references. The lower case c character is used when the symbol is in a special section for small commons.

"D"
"d" The symbol is in the initialized data section.

"G"
"g" The symbol is in an initialized data section for small objects. Some object file formats permit more efficient access to small data objects, such as a global int variable as opposed to a large global array.

"i" For PE-format files this indicates that the symbol is in a section specific to the implementation of DLLs.

For ELF-format files this indicates that the symbol is an indirect function. This is a GNU extension to the standard set of ELF symbol types. It indicates a symbol which, if referenced by a relocation, does not evaluate to its address, but instead must be invoked at runtime. The runtime execution will then return the value to be used in the relocation.

Note -- the actual symbols display for GNU indirect symbols is controlled by the --ifunc-chars command line option. If this option has been provided then the first character in the string will be used for global indirect function symbols. If the string contains a second character then that will be used for local indirect


```

-----function symbols.

-----"I" The symbol is an indirect reference to another symbol.

-----"N" The symbol is a debugging symbol.

-----"n" The symbol is in a non-data, non-code, non-debug-read-
-----only section.

-----"p" The symbol is in a stack unwind section.

-----"R"
-----"r" The symbol is in a read-only data section.

-----"S"
-----"s" The symbol is in an uninitialized or zero-initialized
-----data section for small objects.

-----"T"
-----"t" The symbol is in the text (code) section.

-----"U" The symbol is undefined.

-----"u" The symbol is a unique global symbol. This is a GNU
-----extension to the standard set of ELF symbol bindings.
-----For such a symbol the dynamic linker will make sure that
-----in the entire process there is just one symbol with this
-----name and type in use.

-----"v"
-----"v" The symbol is a weak object. When a weak defined symbol
-----is linked with a normal defined symbol, the normal
-----defined symbol is used with no error. When a weak
-----undefined symbol is linked and the symbol is not defined,
-----the value of the weak symbol becomes zero with no error.
-----On some systems, uppercase indicates that a default value
-----has been specified.

-----"W"
-----"w" The symbol is a weak symbol that has not been
-----specifically tagged as a weak object symbol. When a weak
-----defined symbol is linked with a normal defined symbol,
-----the normal defined symbol is used with no error. When a
-----weak undefined symbol is linked and the symbol is not
-----defined, the value of the symbol is determined in a
-----system-specific manner without error. On some systems,
-----uppercase indicates that a default value has been
-----specified.

-----"-" The symbol is a stabs symbol in an a.out object file. In
-----this case, the next values printed are the stabs other
-----field, the stabs desc field, and the stab type. Stabs
-----symbols are used to hold debugging information.

-----"?" The symbol type is unknown, or object file format
-----specific.

-----* The symbol name. If a symbol has version information
-----associated with it, then the version information is displayed
-----as well. If the versioned symbol is undefined or hidden from
-----linker, the version string is displayed as a suffix to the
-----symbol name, preceded by an @ character. For example
-----foo@VER.1. If the version is the default version to be used
-----when resolving unversioned references to the symbol, then it
-----is displayed as a suffix preceded by two @ characters. For
-----example foo@@VER.2.
OPTIONS
-----The long and short forms of options, shown here as alternatives,
-----are equivalent.

-----A
-----o
-----print-file-name
-----Precede each symbol by the name of the input file (or archive
-----member) in which it was found, rather than identifying the
-----input file once only, before all of its symbols.

-----a
-----debug-syms
-----Display all symbols, even debugger-only symbols; normally
-----these are not listed.

-----B The same as --format=bsd (for compatibility with the MIPS
-----nm).

-----C
-----demangle[=style]
-----Decode (demangle) low-level symbol names into user-level
-----names. Besides removing any initial underscore prepended by
-----the system, this makes C++ function names readable. Different
-----compilers have different mangling styles. The optional
-----demangling style argument can be used to choose an

```

```

-----appropriate_demangling_style_for_your_compiler.

-----no-demangle
-----Do not demangle low-level symbol names.  This is the default.

-----recurse-limit
-----no-recurse-limit
-----recursion-limit
-----no-recursion-limit
-----Enables or disables a limit on the amount of recursion
-----performed whilst demangling strings.  Since the name mangling
-----formats allow for an infinite level of recursion it is
-----possible to create strings whose decoding will exhaust the
-----amount of stack space available on the host machine,
-----triggering a memory fault.  The limit tries to prevent this
-----from happening by restricting recursion to 2048 levels of
-----nesting.

-----The default is for this limit to be enabled, but disabling it
-----may be necessary in order to demangle truly complicated
-----names.  Note however that if the recursion limit is disabled
-----then stack exhaustion is possible and any bug reports about
-----such an event will be rejected.

-----D
-----dynamic
-----Display the dynamic symbols rather than the normal symbols.
-----This is only meaningful for dynamic objects, such as certain
-----types of shared libraries.

-----f-format
-----format=format
-----Use the output format format, which can be "bsd", "sysv",
-----"posix" or "just-symbols".  The default is "bsd".  Only the
-----first character of format is significant; it can be either
-----upper or lower case.

-----g
-----extern-only
-----Display only external symbols.

-----h
-----help
-----Show a summary of the options to nm and exit.

-----ifunc-chars=CHARS
-----When display_GNU_indirect_function_symbols nm will default to
-----using the "i" character for both local indirect functions and
-----global indirect functions.  The ifunc-chars option allows
-----the user to specify a string containing one or two
-----characters.  The first character will be used for global
-----indirect function symbols and the second character, if
-----present, will be used for local indirect function symbols.

-----j---The same as ---format=just-symbols.

-----l
-----line-numbers
-----For each symbol, use debugging information to try to find a
-----filename and line number.  For a defined symbol, look for the
-----line number of the address of the symbol.  For an undefined
-----symbol, look for the line number of a relocation entry which
-----refers to the symbol.  If line number information can be
-----found, print it after the other symbol information.

-----inlines
-----When option -l is active, if the address belongs to a
-----function that was inlined, then this option causes the source
-----information for all enclosing scopes back to the first non-
-----inlined function to be printed as well.  For example, if
-----"main" inlines "callee1" which inlines "callee2", and address
-----is from "callee2", the source information for "callee1" and
-----"main" will also be printed.

-----n
-----v
-----numeric-sort
-----Sort symbols numerically by their addresses, rather than
-----alphabetically by their names.

-----p
-----no-sort
-----Do not bother to sort the symbols in any order; print them in
-----the order encountered.

-----P
-----portability
-----Use the POSIX.2 standard output format instead of the default
-----format.  Equivalent to -f-posix.

-----r
-----reverse-sort

```

```

Reverse the order of the sort (whether numeric or
alphabetic); let the last come first.

-S
Print size
Print both value and size of defined symbols for the "bsd"
output style. This option has no effect for object formats
that do not record symbol sizes, unless size-sort is also
used in which case a calculated size is displayed.

-s
Print armap
When listing symbols from archive members, include the index:
a mapping (stored in the archive by ar or ranlib) of which
modules contain definitions for which names.

-t radix
radix=radix
Use radix as the radix for printing the symbol values. It
must be d for decimal, o for octal, or x for hexadecimal.

-u
undefined-only
Display only undefined symbols (those external to each object
file). By default both defined and undefined symbols are
displayed.

-U
defined-only
Display only defined symbols for each object file. By
default both defined and undefined symbols are displayed.

-V
version
Show the version number of nm and exit.

-X This option is ignored for compatibility with the AIX version
of nm. It takes one parameter which must be the string
32_64. The default mode of AIX nm corresponds to -X-32,
which is not supported by GNU nm.

plugin-name
Load the plugin called name to add support for extra target
types. This option is only available if the toolchain has
been built with plugin support enabled.

If plugin is not provided, but plugin support has been
enabled then nm iterates over the files in
$(libdir)/bfd-plugins in alphabetic order and the first
plugin that claims the object in question is used.

Please note that this plugin search directory is not the one
used by ld's -plugin option. In order to make nm use the
linker plugin it must be copied into the
$(libdir)/bfd-plugins directory. For GCC based compilations
the linker plugin is called liblto.plugin.so.0.0.0. For
Clang based compilations it is called LLVMgold.so. The GCC
plugin is always backwards compatible with earlier versions,
so it is sufficient to just copy the newest one.

-size-sort
Sort symbols by size. For ELF objects symbol sizes are read
from the ELF, for other object types the symbol sizes are
computed as the difference between the value of the symbol
and the value of the symbol with the next higher value. If
the "bsd" output format is used the size of the symbol is
printed, rather than the value, and -S must be used in order
both size and value to be printed.

Note - this option does not work if --undefined-only has been
enabled as undefined symbols have no size.

-special-syms
Display symbols which have a target-specific special meaning.
These symbols are usually used by the target for some special
processing and are not normally helpful when included in the
normal symbol lists. For example for ARM targets this option
would skip the mapping symbols used to mark transitions
between ARM code, THUMB code and data.

-synthetic
Include synthetic symbols in the output. These are special
symbols created by the linker for various purposes. They are
not shown by default since they are not part of the binary's
original source code.

unicode=[default|invalid|locale|escape|hex|highlight]
Controls the display of UTF-8 encoded multibyte characters in
strings. The default (--unicode=default) is to give them no
special treatment. The --unicode=locale option displays the
sequence in the current locale, which may or may not support
them. The options --unicode=hex and --unicode=invalid

```

```

display them as hex byte sequences enclosed by either angle
brackets or curly braces.

The unicode=escape option displays them as escape sequences
(\uxxxx) and the unicode=highlight option displays them as
escape sequences highlighted in red (if supported by the
output device). The colouring is intended to draw attention
to the presence of unicode sequences where they might not be
expected.

-W
no-weak
Do not display weak symbols.

with-symbol-versions
without-symbol-versions
Enables or disables the display of symbol version
information. The version string is displayed as a suffix to
the symbol name, preceded by an @ character. For example
foo@VER.1. If the version is the default version to be used
when resolving unversioned references to the symbol then it
is displayed as a suffix preceded by two @ characters. For
example foo@@VER.2. By default, symbol version information
is displayed.

target=bfdname
Specify an object code format other than your system's
default format.

@file
Read command-line options from file. The options read are
inserted in place of the original @file option. If file does
not exist, or cannot be read, then the option will be treated
literally, and not removed.

Options in file are separated by whitespace. A whitespace
character may be included in an option by surrounding the
entire option in either single or double quotes. Any
character (including a backslash) may be included by
prefixing the character to be included with a backslash. The
file may itself contain additional @file options; any such
options will be processed recursively.

```

SEE ALSO

ar(1), objdump(1), ranlib(1), and the Info entries for binutils.

COPYRIGHT

Copyright (c) 1991–2024 Free Software Foundation, Inc.

Permission is granted to copy, distribute and/or modify this
document under the terms of the GNU Free Documentation License,
Version 1.3 or any later version published by the Free Software
Foundation; with no Invariant Sections, with no Front-Cover
Texts, and with no Back-Cover Texts. A copy of the license is
included in the section entitled "GNU Free Documentation
License".

COLOPHON

This page is part of the binutils (a collection of tools for
working with executable binaries) project. Information about the
project can be found at <http://www.gnu.org/software/binutils/>.
If you have a bug report for this manual page, see
http://sourceware.org/bugzilla/enter_bug.cgi?product=binutils.
This page was obtained from the tarball binutils-2.42.tar.gz
fetched from <https://ftp.gnu.org/gnu/binutils/> on 2024-06-14.
If you discover any rendering problems in this HTML version of
the page, or you believe there is a better or more up-to-date
source for the page, or you have corrections or improvements to
the information in this COLOPHON (which is not part of the
original manual page), send a mail to man-pages@man7.org

binutils-2.42 2024-06-14 NM(1)

3.9 strace: Trace System Calls and Signals

NAME

strace – trace system calls and signals

SYNOPSIS

```

strace [-ACdfhikkqrrttTvVwxyyYzZ] [-a column] [-b execve]
[-e expr]... [-I n] [-o file] [-O overhead] [-p pid]...
[-P path]... [-s strsize] [-S sortby] [-U columns]
[-X format] [--seccomp-bpf]
[--stack-trace-frame-limit=limit] [--syscall-limit=limit]
[--seccomp-context=format] [--tips=format] { -p pid | [-DDD]
[-E var=val]}... [-u username] command [args]

strace -c [-dfwzZ] [-b execve] [-e expr]... [-I n] [-O overhead]
[-p pid]... [-P path]... [-S sortby] [-U columns]
[--seccomp-bpf] [--syscall-limit=limit] [--tips=format]
{ -p pid | [-DDD] [-E var=val]}... [-u username] command
[args]

```

strace —tips[=format]
DESCRIPTION

In the simplest case strace runs the specified command until it exits. It intercepts and records the system calls which are called by a process and the signals which are received by a process. The name of each system call, its arguments and its return value are printed on standard error or to the file specified with the -o option.

strace is a useful diagnostic, instructional, and debugging tool. System administrators, diagnosticians and trouble-shooters will find it invaluable for solving problems with programs for which the source is not readily available since they do not need to be recompiled in order to trace them. Students, hackers and the overly-curious will find that a great deal can be learned about a system and its system calls by tracing even ordinary programs. And programmers will find that since system calls and signals are events that happen at the user/kernel interface, a close examination of this boundary is very useful for bug isolation, sanity checking and attempting to capture race conditions.

Each line in the trace contains the system call name, followed by its arguments in parentheses and its return value. An example from tracing the command "cat /dev/null" is:

```
open("/dev/null", O_RDONLY) = 3
```

Errors (typically a return value of -1) have the errno symbol and error string appended.

```
open("/foo/bar", O_RDONLY) = -1 ENOENT (No such file or directory)
```

Signals are printed as signal symbol and decoded signifo structure. An excerpt from stracing and interrupting the command "sleep_666" is:

```
sigsuspend([ <unfinished ...>
— SIGINT {si_signo=SIGINT, si_code=SIUSER, si_pid=...} —
+++ killed by SIGINT +++
```

If a system call is being executed and meanwhile another one is being called from a different thread/process then strace will try to preserve the order of those events and mark the ongoing call as being unfinished. When the call returns it will be marked as resumed.

```
[pid 28772] select(4, [3], NULL, NULL, NULL <unfinished ...>
[pid 28779] clock_gettime(CLOCK_REALTIME, {tv_sec=1130322148, tv_nsec=3977000}) = 0
[pid 28772] <... select resumed> ) = 1 (in [3])
```

Interruption of a (restartable) system call by a signal delivery is processed differently as kernel terminates the system call and also arranges its immediate reexecution after the signal handler completes.

```
read(0, 0x7ffff72cf5cf, 1) = ? ERESTARTSYS (To be restarted)
— SIGALRM {si_signo=SIGALRM, si_code=SI_KERNEL} —
rt_sigreturn({mask=[]}) = 0
read(0, "", 1) = 0
```

Arguments are printed in symbolic form with passion. This example shows the shell performing ">>xyzy" output redirection:

```
open("xyzy", O_WRONLY|O_APPEND|O_CREAT, 0666) = 3
```

Here, the second and the third argument of open(2) are decoded by breaking down the flag argument into its three bitwise-OR constituents and printing the mode value in octal by tradition. Where the traditional or native usage differs from ANSI or POSIX, the latter forms are preferred. In some cases, strace output is proven to be more readable than the source.

Structure pointers are dereferenced and the members are displayed as appropriate. In most cases, arguments are formatted in the most C-like fashion possible. For example, the essence of the command "ls -l /dev/null" is captured as:

```
lstat("/dev/null", {st_mode=S_IFCHR|0666, st_rdev=makedev(0x1, 0x3), ...}) = 0
```

Notice how the 'struct_stat' argument is dereferenced and how each member is displayed symbolically. In particular, observe how the st_mode member is carefully decoded into a bitwise-OR of symbolic and numeric values. Also notice in this example that the first argument to lstat(2) is an input to the system call and the second argument is an output. Since output arguments are not modified if the system call fails, arguments may not always be dereferenced. For example, retrying the "ls -l" example with a non-existent file produces the following line:

```
lstat("/foo/bar", 0xb004) = -1 ENOENT (No such file or directory)
```

```

In this case the porch light is on but nobody is home.

Syscalls unknown to strace are printed raw, with the unknown
system call number printed in hexadecimal form and prefixed with
"syscall.":

syscall-0xbad(0x1, 0x2, 0x3, 0x4, 0x5, 0x6) = -1 ENOSYS (Function not implemented)

Character pointers are dereferenced and printed as C strings.
Non-printing characters in strings are normally represented by
ordinary C escape codes. Only the first strsize (32 by default)
bytes of strings are printed; longer strings have an ellipsis
appended following the closing quote. Here is a line from "ls
-----l" where the getpwuid(3) library routine is reading the password
file:

read(3, "root::0:0:System-Administrator:/"... , 1024) = 422

While structures are annotated using curly braces, pointers to
basic types and arrays are printed using square brackets with
commas separating the elements. Here is an example from the
command id(1) on a system with supplementary group ids:

getgroups(32, [100, 0]) = 2

On the other hand, bit-sets are also shown using square brackets,
but set elements are separated only by a space. Here is the
shell, preparing to execute an external command:

sigprocmask(SIG_BLOCK, [CHLD TTOU], []) = 0

Here, the second argument is a bit-set of two signals, SIGCHLD
and SIGTTOU. In some cases, the bit-set is so full that printing
out the unset elements is more valuable. In that case, the bit-
set is prefixed by a tilde like this:

sigprocmask(SIG_UNBLOCK, ~[], NULL) = 0

Here, the second argument represents the full set of all signals.

```

OPTIONS

General

```

-e expr
    A qualifying expression which modifies which events to
    trace or how to trace them. The format of the expression
    is:

        [qualifier=][!]value[,value]...

    where qualifier is one of trace (or t), trace-fds (or
    trace-fd or fd or fds), abbrev (or a), verbose (or v), raw
    (or x), signal (or signals or s), read (or reads or r),
    write (or writes or w), fault, inject, status, quiet (or
    silent or silence or q), secontext, decode-fds (or
    decode-fd), decode-pids (or decode-pid), or kvm, and value
    is a qualifier-dependent symbol or number. The default
    qualifier is trace. Using an exclamation mark negates the
    set of values. For example, -e open means literally
    -e trace=open which in turn means trace only the open
    system call. By contrast, -e trace!=open means to trace
    every system call except open. In addition, the special
    values all and none have the obvious meanings.

    Note that some shells use the exclamation point for
    history expansion even inside quoted arguments. If so,
    you must escape the exclamation point with a backslash.

```

Startup

```

-E var=val
  --env=var=val
    Run command with var=val in its list of environment
    variables.

-E var
  --env=var
    Remove var from the inherited list of environment
    variables before passing it on to the command.

-p pid
  --attach=pid
    Attach to the process with the process ID pid and begin
    tracing. The trace may be terminated at any time by a
    keyboard interrupt signal (CTRL-C). strace will respond
    by detaching itself from the traced process(es) leaving it
    (them) to continue running. Multiple -p options can be
    used to attach to many processes in addition to command
    (which is optional if at least one -p option is given).
    Multiple process IDs, separated by either comma (",") or
    space (" "), _tab_, or _newline_ character, _can_ be _provided_ as
    an argument to a single -p option, so, for example, -p
    "$(pidof PROG)" and -p "$(pgrep PROG)" _syntaxes_ are
    supported.

```

```

-----u_username
-----user=username
-----Run command with the user ID, group ID, and supplementary
-----groups of username. This option is only useful when
-----running as root and enables the correct execution of
-----setuid and/or setgid binaries. Unless this option is used
-----setuid and setgid programs are executed without effective
-----privileges.
-----u_UID:GID
-----user=UID:GID
-----Alternative syntax where the program is started with
-----exactly the given user and group IDs, and an empty list of
-----supplementary groups. In this case, user and group name
-----lookups are not performed.

-----argv0=name
-----Set argv[0] of the command being executed to name. Useful
-----for tracing multi-call executables which interpret
-----argv[0], such as busybox or kmod.

---Tracing
-----b_syscall
-----detach-on=syscall
-----If specified syscall is reached, detach from traced
-----process. Currently, only execve(2) syscall is supported.
-----This option is useful if you want to trace multi-threaded
-----process and therefore require -f, but don't want to trace
-----its (potentially very complex) children.

-----D
-----daemonize
-----daemonize=grandchild
-----Run tracer process as a grandchild, not as the parent of
-----the tracee. This reduces the visible effect of strace by
-----keeping the tracee a direct child of the calling process.

-----DD
-----daemonize=pgroup
-----daemonize=grp
-----Run tracer process as tracee's grandchild in a separate
-----process group. In addition to reduction of the visible
-----effect of strace, it also avoids killing of strace with
-----kill(2) issued to the whole process group.

-----DDD
-----daemonize=session
-----Run tracer process as tracee's grandchild in a separate
-----session ("true daemonisation"). In addition to reduction
-----of the visible effect of strace, it also avoids killing of
-----strace upon session termination.

-----f
-----follow-forks
-----Trace child processes as they are created by currently
-----traced processes as a result of the fork(2), vfork(2) and
-----clone(2) system calls. Note that -pPID-f will attach
-----all threads of process PID if it is multi-threaded, not
-----only thread with thread.id==PID.

-----output-separately
-----If the --output=filename option is in effect, each
-----process trace is written to filename.pid where pid is
-----the numeric process id of each process.

-----ff
-----follow-forks --output-separately
-----Combine the effects of --follow-forks and
-----output-separately options. This is incompatible with
-----c, since no per-process counts are kept.

-----One might want to consider using strace-log-merge(1) to
-----obtain a combined strace-log-view.

-----I_interruptible
-----interruptible=interruptible
-----When strace can be interrupted by signals (such as
-----pressing CTRL-C).

-----1, anywhere
-----no signals are blocked;
-----2, waiting
-----fatal signals are blocked while decoding syscall
----- (default);
-----3, never
-----fatal signals are always blocked (default if -o
-----FILE_PROG);
-----4, never tstp
-----fatal signals and SIGTSTP (CTRL-Z) are always
-----blocked (useful to make strace -o FILE_PROG not
-----stop on CTRL-Z, default if -D).

-----syscall-limit=limit

```

```

-----Detach all tracees when limit number of syscalls have been
-----captured. Syscalls filtered out via --trace, --trace-path
-----or --status options are not considered when keeping track
-----of the number of syscalls that are captured.

-----kill -on-exit
-----Apply PTRACE_O_EXITKILL ptrace option to all tracee
-----processes (which sends a SIGKILL signal to the tracee if
-----the tracer exits) and do not detach them on cleanup so
-----they will not be left running after the tracer exits.
-----kill -on-exit is not compatible with -p/--attach options.

---Filtering
-----e_trace=syscall_set
-----e_t=syscall_set
-----trace=syscall_set
-----Trace only the specified set of system calls. --syscall-set
-----is defined as [!] value[, value], and value can be one of
-----the following:

-----syscall
-----Trace specific syscall, specified by its name (see
-----syscalls(2) for a reference, but also see NOTES).

-----?value Question mark before the syscall qualification
-----allows suppression of error in case no syscalls
-----matched the qualification provided.

-----value@64
-----Limit the syscall specification described by value
-----to 64-bit personality.

-----value@32
-----Limit the syscall specification described by value
-----to 32-bit personality.

-----value@x32
-----Limit the syscall specification described by value
-----to x32 personality.

-----all ---Trace all system calls.

-----/regex Trace only those system calls that match the regex.
-----You can use POSIX Extended Regular Expression
-----syntax (see regex(7)).

-----%file
-----file ---Trace all system calls which take a file name as an
-----argument. You can think of this as an abbreviation
-----for --e_trace=open, stat, chmod, unlink, ... which is
-----useful to seeing what files the process is
-----referencing. Furthermore, using the abbreviation
-----will ensure that you don't accidentally forget to
-----include a call like lstat(2) in the list. --Betchya
-----woulda forgot that one. The syntax without a
-----preceding percent sign ("e trace=file") is
-----deprecated.

-----%process
-----process ---Trace system calls associated with process
-----lifecycle (creation, exec, termination). The
-----syntax without a preceding percent sign ("e
-----trace=process") is deprecated.

-----%net
-----%network
-----network ---Trace all the network-related system calls. The
-----syntax without a preceding percent sign ("e
-----trace=network") is deprecated.

-----%signal
-----signal ---Trace all signal-related system calls. The syntax
-----without a preceding percent sign ("e
-----trace=signal") is deprecated.

-----%ipc
-----ipc ---Trace all IPC-related system calls. The syntax
-----without a preceding percent sign ("e trace=ipc")
-----is deprecated.

-----%desc
-----desc ---Trace all file descriptor-related system calls.
-----The syntax without a preceding percent sign ("e
-----trace=desc") is deprecated.

-----%memory
-----memory ---Trace all memory-mapping-related system calls. The
-----syntax without a preceding percent sign ("e
-----trace=memory") is deprecated.

```



```

-----%creds_Trace_system_calls_that_read_or_modify_user_and
-----group_identifiers_or_capability_sets.

-----%stat_Trace_stat_syscall_variants.

-----%lstat_Trace_lstat_syscall_variants.

-----%fstat_Trace_fstat, _fstatat, _and_statx_syscall_variants.

-----%%stat_Trace_syscalls_used_for_requesting_file_status
----- (stat, _lstat, _fstat, _fstatat, _statx, _and_their
----- variants).

-----%statfs
-----Trace_statfs, _statfs64, _fstatvfs, _osf_statfs, _and
-----osf_statfs64_system_calls. _The_same_effect_can_be
-----achieved_with--e_trace=/^(.*)?statvfs_regular
-----expression.

-----%fstatfs
-----Trace_fstatfs, _fstatfs64, _fstatvfs, _osf_fstatfs,
-----and_osf_fstatfs64_system_calls. _The_same_effect
-----can_be_achieved_with--e_trace=/fstatvfs_regular
-----expression.

-----%%statfs
-----Trace_syscalls_related_to_file_system_statistics
----- (statfs-like, _fstatfs-like, _and_ustat). _The_same
-----effect_can_be_achieved_with
-----e_trace=/statvfs|fsstat|ustat_regular_expression.

-----%clock_Trace_system_calls_that_read_or_modify_system
-----clocks.

-----%pure_Trace_syscalls_that_always_succeed_and_have_no
-----arguments. _Currently, _this_list_includes
-----arc_gettls(2), _getdtablesize(2), _getegid(2),
-----getegid32(2), _geteuid(2), _geteuid32(2), _getgid(2),
-----getgid32(2), _getpagesize(2), _getpgid(2), _getpid(2),
-----getppid(2), _get_thread_area(2) (on_architectures
-----other_than_x86), _gettid(2), _gettls(2), _getuid(2),
-----getuid32(2), _getxgid(2), _getxpid(2), _getxuid(2),
-----kern_features(2), _and_metag_gettls(2)_syscalls.

-----The--c_option_is_useful_for_determining_which_system_calls
-----might_be_useful_to_trace. _For_example,
-----trace=open, close, read, write_means_to_only_trace_those_four
-----system_calls. _Be_careful_when_making_inferences_about_the
-----user/kernel_boundary_if_only_a_subset_of_system_calls_are
-----being_monitored. _The_default_is_trace=all.

-----e_trace-fd=set
-----e_trace-fds=set
-----e_fd=set
-----e_fds=set
-----trace-fds=set
-----Trace_only_the_syscalls_that_operate_on_the_specified
-----subset_of_(non-negative)_file_descriptors. _Note_that
-----usage_of_this_option_also_filters_out_all_the_syscalls
-----that_do_not_operate_on_file_descriptors_at_all. _Applies
-----in_(inclusive)_disjunction_with_the--trace-path_option.

-----e_signal=set
-----e_signals=set
-----e_s=set
-----signal=set
-----Trace_only_the_specified_subset_of_signals. _The_default
-----is_signal=all. _For_example, _signal=!SIGIO_(or_signal=!io)
-----causes_SIGIO_signals_not_to_be_traced.

-----e_status=set
-----status=set
-----Print_only_system_calls_with_the_specified_return_status.
-----The_default_is_status=all. _When_using_the_status
-----qualifier, _because_trace_waits_for_system_calls_to_return
-----before_deciding_whether_they_should_be_printed_or_not, _the
-----traditional_order_of_events_may_not_be_preserved_anymore.
-----If_two_system_calls_are_executed_by_concurrent_threads,
-----trace_will_first_print_both_the_entry_and_exit_of_the
-----first_system_call_to_exit, _regardless_of_their_respective
-----entry_time. _The_entry_and_exit_of_the_second_system_call
-----to_exit_will_be_printed_afterwards. _Here_is_an_example
-----when_select(2)_is_called, _but_a_different_thread_calls
-----clock_gettime(2)_before_select(2)_finishes:

-----[pid-28779]_1130322148.939977_clock_gettime(CLOCK_REALTIME,_{1130322148,_939977000})_=_0
-----[pid-28772]_1130322148.438139_select(4,_{3},_NULL,_NULL,_NULL)_=_1_(in_{3})

-----set_can_include_the_following_elements:

-----successful
-----Trace_system_calls_that_returned_without_an_error

```

```

-----code. The -z option has the effect of
-----status=successful.
-----failed Trace system calls that returned with an error
-----code. The -Z option has the effect of
-----status=failed.
-----unfinished
-----Trace system calls that did not return. This might
-----happen, for example, due to an execve call in a
-----neighbour thread.
-----unavailable
-----Trace system calls that returned but strace failed
-----to fetch the error status.
-----detached
-----Trace system calls for which strace detached before
-----the return.

-----P_path
-----trace-path=path
-----Trace only system calls accessing path. Multiple -P
-----options can be used to specify several paths. Applies in
-----{inclusive} disjunction with the --trace-fds option.

-----z
-----successful-only
-----Print only syscalls that returned without an error code.

-----Z
-----failed-only
-----Print only syscalls that returned with an error code.

___Output_format
-----a_column
-----columns=column
-----Align return values in a specific column (default column
-----40).

-----e_abbrev=syscall_set
-----e_a=syscall_set
-----abbrev=syscall_set
-----Abbreviate the output from printing each member of large
-----structures. The syntax of the syscall_set specification
-----is the same as in the -e trace option. The default is
-----abbrev=all. The -v option has the effect of abbrev=none.

-----e_verbose=syscall_set
-----e_v=syscall_set
-----verbose=syscall_set
-----Dereference structures for the specified set of system
-----calls. The syntax of the syscall_set specification is the
-----same as in the -e trace option. The default is
-----verbose=all.

-----e_raw=syscall_set
-----e_x=syscall_set
-----raw=syscall_set
-----Print raw, undecoded arguments for the specified set of
-----system calls. The syntax of the syscall_set specification
-----is the same as in the -e trace option. This option has
-----the effect of causing all arguments to be printed in
-----hexadecimal. This is mostly useful if you don't trust the
-----decoding or you need to know the actual numeric value of
-----an argument. See also -X raw option.

-----e_read=set
-----e_reads=set
-----e_r=set
-----read=set
-----Perform a full hexadecimal and ASCII dump of all the data
-----read from file descriptors listed in the specified set.
-----For example, to see all input activity on file descriptors
-----3 and 5 use -e_read=3,5. Note that this is independent
-----from the normal tracing of the read(2) system call which
-----is controlled by the option -e trace=read.

-----e_write=set
-----e_writes=set
-----e_w=set
-----write=set
-----Perform a full hexadecimal and ASCII dump of all the data
-----written to file descriptors listed in the specified set.
-----For example, to see all output activity on file
-----descriptors 3 and 5 use -e_write=3,5. Note that this is
-----independent from the normal tracing of the write(2) system
-----call which is controlled by the option -e trace=write.

-----e_quiet=set
-----e_silent=set
-----e_silence=set
-----e_q=set
-----quiet=set
-----silent=set
-----silence=set

```

```

-----Suppress various information messages. --The default is
-----quiet=none. --set can include the following elements:

-----attach Suppress messages about attaching and detaching (" [
        Process NNNN attached ]", "[ Process NNNN detached
        ]").
-----exit Suppress messages about process exits ("+++ exited
        with SSS +++").
-----path-resolution
        Suppress messages about resolution of paths
        provided via the -P option ("Requested path "... "
        resolved into "...").
-----personality
        Suppress messages about process personality changes
        (" [ Process PID=NNNN runs in PPP mode. ]").
-----thread-execute
        Suppress messages about process being superseded by
        execve(2) in another thread ("+++ superseded by
        execve in pid NNNN +++").

-----e_decode-fds=set
        decode-fds=set
        Decode various information associated with file
        descriptors. --The default is decode-fds=none. --set can
        include the following elements:

        -----path Print file paths. --Also enables printing of
        -----tracee's current working directory when _AT_FDCWD
        -----constant is used.
        -----socket Print socket protocol-specific information,
        -----dev Print character/block device numbers.
        -----pidfd Print PIDs associated with pidfd file
        -----descriptors.
        -----signalfd Print signal masks associated with signalfd file
        -----descriptors.

-----e_decode-pids=set
        decode-pids=set
        Decode various information associated with process IDs
        (and also thread IDs, process group IDs, and session IDs).
        The default is decode-pids=none. --set can include the
        following elements:

        -----comm Print command names associated with thread or
        -----process IDs.
        -----pidns Print thread, process, process group, and session
        -----IDs in _tracee's PID namespace if the tracee is in
        -----a different PID namespace.

-----e_kvm=vcpu
        kvm=vcpu
        Print the exit reason of kvm-vcpu. --Requires Linux kernel
        version 4.16.0 or higher.

-----i
        instruction-pointer
        Print the instruction pointer at the time of the system
        call.

-----n
        syscall-number
        Print the syscall number.

-----k
        stack-trace[=symbol]
        Print the execution stack trace of the traced processes
        after each system call.

-----kk
        stack-trace=source
        Print the execution stack trace and source code
        information of the traced processes after each system
        call. --This option expects the target program is compiled
        with appropriate debug options: "-g" (gcc), or "-g
        -gdwarf-aranges" (clang).

-----stack-trace-frame-limit=limit
        Print no more than this amount of stack trace frames when
        backtracing a system call (the default is 256). --Use this
        option with the --stack-trace (or -k) option.

-----o_filename
        output=filename
        Write the trace output to the file filename rather than to
        stderr. --filename.pid form is used if -ff option is
        supplied. --If the argument begins with '|' or '!', the
        rest of the argument is treated as a command and all
        output is piped to it. --This is convenient for piping the
        debugging output to a program without affecting the
        redirections of executed programs. --The latter is not
        compatible with -ff option currently.

```

```

-----A
-----output--append--mode
-----Open the file provided in the -o option in append mode.

-----q
-----quiet
-----quiet=attach, personality
-----Suppress messages about attaching, detaching, and
-----personality changes. This happens automatically when
-----output is redirected to a file and the command is run
-----directly instead of attaching.

-----qq
-----quiet=attach, personality, exit
-----Suppress messages attaching, detaching, personality
-----changes, and about process exit status.

-----qqq
-----quiet=all
-----Suppress all suppressible messages (please refer to the -e
-----quiet option description for the full list of suppressible
-----messages).

-----r
-----relative--timestamps[=precision]
-----Print a relative timestamp upon entry to each system call.
-----This records the time difference between the beginning of
-----successive system calls. precision can be one of s (for
-----seconds), ms (milliseconds), us (microseconds), or ns
----- (nanoseconds), and allows setting the precision of time
-----value being printed. Default is us (microseconds). Note
-----that since -r option uses the monotonic clock time for
-----measuring time difference and not the wall clock time, its
-----measurements can differ from the difference in time
-----reported by the -t option.

-----s--strsize
-----string--limit=strsize
-----Specify the maximum string size to print (the default is
-----32). Note that filenames are not considered strings and
-----are always printed in full.

-----absolute--timestamps[=[format:]format],[precision:]precision]]
-----timestamps[=[format:]format],[precision:]precision]]
-----Prefix each line of the trace with the wall clock time in
-----the specified format with the specified precision. format
-----can be one of the following:

-----none No time stamp is printed. Can be used to override
-----the previous setting.
-----time Wall clock time (strftime(3) format string is %T).
-----unix Number of seconds since the epoch (strftime(3)
-----format string is %s).

-----precision can be one of s (for seconds), ms
----- (milliseconds), us (microseconds), or ns (nanoseconds).
-----Default arguments for the option are
-----format:time, precision:s.

-----t
-----absolute--timestamps
-----Prefix each line of the trace with the wall clock time.

-----tt
-----absolute--timestamps=precision:us
-----If given twice, the time printed will include the
-----microseconds.

-----ttt
-----absolute--timestamps=format:unix, precision:us
-----If given thrice, the time printed will include the
-----microseconds and the leading portion will be printed as
-----the number of seconds since the epoch.

-----T
-----syscall--times[=precision]
-----Show the time spent in system calls. This records the
-----time difference between the beginning and the end of each
-----system call. precision can be one of s (for seconds), ms
----- (milliseconds), us (microseconds), or ns (nanoseconds),
-----and allows setting the precision of time value being
-----printed. Default is us (microseconds).

-----v
-----no--abbrev
-----Print unabbreviated versions of environment, stat,
-----termios, etc. calls. These structures are very common in
-----calls and so the default behavior displays a reasonable
-----subset of structure members. Use this option to get all
-----of the gory details.

```

```

-----strings-in-hex[=option]
-----Control usage of escape sequences with hexadecimal numbers
-----in the printed strings. Normally (when no
-----strings-in-hex-or-x option is supplied), escape
-----sequences are used to print non-printable and non-ASCII
-----characters (that is, characters with a character code less
-----than 32 or greater than 127), or to disambiguate the
-----output (so, for quotes and other characters that encase
-----the printed string, for example, angle brackets, in case
-----of file descriptor path output); for the former use case,
-----unless it is a white space character that has a symbolic
-----escape sequence defined in the C standard (that is, "\textbackslash t"
-----for a horizontal tab, "\textbackslash n" for a newline, "\textbackslash v" for a
-----vertical tab, "\textbackslash f" for a form feed, page break, and "\textbackslash r"
-----for a carriage return) are printed using escape sequences
-----with numbers that correspond to their byte values, with
-----octal number format being the default. Option can be one
-----of the following:

-----none Hexadecimal numbers are not used in the output at
-----all. When there is a need to emit an escape
-----sequence, octal numbers are used.
-----non-ascii-chars
-----Hexadecimal numbers are used instead of octal in
-----the escape sequences.
-----non-ascii
-----Strings that contain non-ASCII characters are
-----printed using escape sequences with hexadecimal
-----numbers.
-----all All strings are printed using escape sequences with
-----hexadecimal numbers.

-----When the option is supplied without an argument, all is
-----assumed.

-----x
-----strings-in-hex=non-ascii
-----Print all non-ASCII strings in hexadecimal string format.

-----xx
-----strings-in-hex[=all]
-----Print all strings in hexadecimal string format.

-----X,format
-----const-print-style=format
-----Set the format for printing of named constants and flags.
-----Supported format values are:

-----raw Raw number output, without decoding.
-----abbrev Output a named constant or a set of flags instead
-----of the raw number if they are found. This is the
-----default strace behaviour.
-----verbose
-----Output both the raw value and the decoded string
----- (as a comment).

-----Y
-----decode-fds
-----decode-fds=path
-----Print paths associated with file descriptor arguments and
-----with the _AT_FDCWD constant.

-----yy
-----decode-fds=all
-----Print all available information associated with file
-----descriptors: protocol-specific information associated with
-----socket file descriptors, block/character device number
-----associated with device file descriptors, and PIDs
-----associated with pidfd file descriptors.

-----pidns-translation
-----decode-pids=pidns
-----If strace and tracee are in different PID namespaces,
-----print PIDs in strace's namespace, too.

-----Y
-----decode-pids=comm
-----Print command names for PIDs.

-----secontext[=format]
-----e_secontext=format
-----When SELinux is available and is not disabled, print in
-----square brackets SELinux contexts of processes, files, and
-----descriptors. The format argument is a comma-separated
-----list of items being one of the following:

-----full Print the full context (user, role, type
-----level and category).
-----mismatch Also print the context recorded by the
-----SELinux database in case the current
-----context differs. The latter is printed
-----after two exclamation marks (!!).

```

```

-----The default value for --secontext is !full, mismatch which
-----prints only the type instead of full context and doesn't
-----check for context mismatches.

-----always-show-pid
-----Show PID prefix also for the process started by strace.
-----Implied when -f and -o are both specified.

___Statistics
-----c
-----summary-only
-----Count time, calls, and errors for each system call and
-----report a summary on program exit, suppressing the regular
-----output. This attempts to show system time (CPU time spent
-----running in the kernel) independent of wall clock time. If
-----c is used with -f, only aggregate totals for all traced
-----processes are kept.

-----C
-----summary
-----Like -c but also print regular output while processes are
-----running.

-----O_overhead
-----summary-syscall-overhead=overhead
-----Set the overhead for tracing system calls to overhead.
-----This is useful for overriding the default heuristic for
-----guessing how much time is spent in mere measuring when
-----timing system calls using the -c option. The accuracy of
-----the heuristic can be gauged by timing a given program run
-----without tracing (using time(1)) and comparing the
-----accumulated system call time to the total produced using
-----c.

-----The format of overhead specification is described in
-----section Time specification format description.

-----S_sortby
-----summary-sort-by=sortby
-----Sort the output of the histogram printed by the -c option
-----by the specified criterion. Legal values are time (or
-----time-percent or time-total or total-time), min-time (or
-----shortest or time-min), max-time (or longest or time-max),
-----avg-time (or time-avg), calls (or count), errors (or
-----error), name (or syscall or syscall-name), and nothing (or
-----none); default is time.

-----U_columns
-----summary-columns=columns
-----Configure a set (and order) of columns being shown in the
-----call summary. The columns argument is a comma-separated
-----list with items being one of the following:

-----time-percent (or time)
-----Percentage of cumulative time consumed by a
-----specific system call.
-----total-time (or time-total)
-----Total system (or wall clock, if -w option is
-----provided) time consumed by a specific system call.
-----min-time (or shortest or time-min)
-----Minimum observed call duration.
-----max-time (or longest or time-max)
-----Maximum observed call duration.
-----avg-time (or time-avg)
-----Average call duration.
-----calls (or count)
-----Call count.
-----errors (or error)
-----Error count.
-----name (or syscall or syscall-name)
-----Syscall name.

-----The default value is
-----time-percent, total-time, avg-time, calls, errors, name. If
-----the name field is not supplied explicitly, it is added as
-----the last column.

-----w
-----summary-wall-clock
-----Summarise the time difference between the beginning and
-----end of each system call. The default is to summarise the
-----system time.

___Tampering
-----e inject=syscall.set[: error=errno[: retval=value][: signal=sig]
-----[: syscall=syscall][: delay_enter=delay][: delay_exit=delay]
-----[: poke_enter=@argN=DATAN, @argM=DATAM...][:
-----[: poke_exit=@argN=DATAN, @argM=DATAM...][: when=expr]
-----inject=syscall.set[: error=errno[: retval=value][: signal=sig]
-----[: syscall=syscall][: delay_enter=delay][: delay_exit=delay]
-----[: poke_enter=@argN=DATAN, @argM=DATAM...]
```

```

[[: poke_exit=@argN=DATAN,@argM=DATAM... ][: when=expr ]
Perform __syscall__tampering__for__the__specified__set__of
__syscalls__. __The__syntax__of__the__syscall__set__specification__is
the__same__as__in__the__e__trace__option__.

At__least__one__of__error__,__retval__,__signal__,__delay_enter__,
__delay_exit__,__poke_enter__,__or__poke_exit__options__has__to__be
specified. __error__and__retval__are__mutually__exclusive__.

If__error=errno__option__is__specified__,__a__fault__is__injected
into__a__syscall__invocation__:__the__syscall__number__is__replaced
by__1__which__corresponds__to__an__invalid__syscall__(unless__a
syscall__is__specified__with__:__syscall=__option__),__and__the__error
code__is__specified__using__a__symbolic__errno__value__like__ENOSYS
or__a__numeric__value__within__1..4095__range__.

If__retval=value__option__is__specified__,__success__injection__is
performed__:__the__syscall__number__is__replaced__by__1__,__but__a
bogus__success__value__is__returned__to__the__callee__.

If__signal=sig__option__is__specified__with__either__a__symbolic
value__like__SIGSEGV__or__a__numeric__value__within__1..SIGRTMAX
range__,__that__signal__is__delivered__on__entering__every__syscall
specified__by__the__set__.

If__delay_enter=delay__or__:__delay_exit=delay__options__are
specified__,__delay__injection__is__performed__:__the__tracee__is
delayed__by__time__period__specified__by__delay__on__entering__or
exiting__the__syscall__,__respectively__. __The__format__of__delay
specification__is__described__in__section__Time__specification
format__description__.

If_____:__poke_enter=@argN=DATAN,@argM=DATAM..._____:__or
_____:__poke_exit=@argN=DATAN,@argM=DATAM..._____:__options_____:__are
specified__,__tracee's__memory__at__locations__,__pointed__to__by
system__call__arguments__argN__and__argM__(going__from__arg1__to
arg7)__is__overwritten__by__data__DATAN__and__DATAM__(specified__in
hexadecimal__format__);_____:__for_____:__example
_____:__poke_enter=@arg1=0000DEAD0000BEEF)_____:__poke_enter__modifies
memory__on__syscall__enter__,__and__:__poke_exit__on__exit__.

If__signal=sig__option__is__specified__without__:__error=errno__,
__retval=value__,__or__:__delay_{enter,exit}=usecs__options__,__then
only__a__signal__sig__is__delivered__without__a__syscall__fault__or
delay__injection__. __Conversely__,_____:__error=errno__,__or
__retval=value__,__option__without_____:__delay_enter=delay__,
__delay_exit=delay__,__or__:__signal=sig__options__injects__a__fault
without__delivering__a__signal__or__injecting__a__delay__,__etc__.

If__signal=sig__option_____:__is_____:__specified__together_____:__with
_____:__error=errno__,__or__:__retval=value__,__then__both__injection__of__a
fault__or__success__and__signal__delivery__are__performed__.

if__syscall=syscall__option__is__specified__,__the__corresponding
syscall__with__no__side__effects__is__injected__instead__of__1__.
Currently__,__only__"pure"__(see__e__trace=%pure__description)
syscalls__can__be__specified__there__.

Unless__a__:__when=expr__subexpression__is_____:__specified__,__an
injection__is__being__made__into__every__invocation__of__each
syscall__from__the__set__.

The__format__of__the__subexpression__is__:

first [.. last ][+ [step ]]

Number__first__stands__for__the__first__invocation__number__in__the
range__,__number__last__stands__for__the__last__invocation__number
in__the__range__,__and__step__stands__for__the__step__between__two
consecutive__invocations__. __The__following__combinations__are
useful__:

first__For__every__syscall__from__the__set__,__perform__an
injection__for__the__syscall__invocation__number__first
only.
first .. last__For__every__syscall__from__the__set__,__perform__an
injection__for__the__syscall__invocation__number__first
and__all__subsequent__invocations__until__the__invocation
number__last__(inclusive).
first+__For__every__syscall__from__the__set__,__perform__injections
for__the__syscall__invocation__number__first__and__all
subsequent__invocations.
first .. last+__For__every__syscall__from__the__set__,__perform__injections
for__the__syscall__invocation__number__first__and__all
subsequent__invocations__until__the__invocation__number
last__(inclusive).
first+step__For__every__syscall__from__the__set__,__perform__injections
for__syscall__invocations__number__first__,__first+step__,
first+step+step__,__and__so__on.

```

```

-----first .. last+step
-----Same as the previous, but consider only syscall
-----invocations with numbers up to last (inclusive).

-----For example, to fail each third and subsequent chdir
-----syscalls with ENOENT, use
-----e_inject=chdir: error=ENOENT: when=3+.

-----The valid range for numbers first and step is 1..65535,
-----and for number last is 1..65534.

-----An injection expression can contain only one error= or
-----retval= specification, and only one signal= specification.
-----If an injection expression contains multiple when=
-----specifications, the last one takes precedence.

-----Accounting of syscalls that are subject to injection is
-----done per syscall and per tracee.

-----Specification of syscall injection can be combined with
-----other syscall filtering options, for example, -P
-----/dev/urandom -e_inject=file: error=ENOENT.

-----e_fault=syscall.set[: error=errno][: when=expr]
-----fault=syscall.set[: error=errno][: when=expr]
-----Perform syscall fault injection for the specified set of
-----syscalls.

-----This is equivalent to more generic -e_inject=expression
-----with default value of errno option set to ENOSYS.

___Miscellaneous
-----d
-----debug
-----Show some debugging output of strace itself on the
-----standard error.

-----F-----This option is deprecated. It is retained for backward
-----compatibility only and may be removed in future releases.
-----Usage of multiple instances of -F option is still
-----equivalent to a single -f, and it is ignored at all if
-----used along with one or more instances of -f option.

-----h
-----help Print the help summary.

-----seccomp-bpf
-----Try to enable use of seccomp-bpf (see seccomp(2)) to have
-----ptrace(2) stops only when system calls that are being
-----traced occur in the traced processes.

-----This option has no effect unless -f/--follow-forks is also
-----specified. --seccomp-bpf is not compatible with
-----syscall-limit and -b/--detach-on options. It is also
-----not applicable to processes attached using -p/--attach
-----option.

-----An attempt to enable system calls filtering using seccomp-
-----bpf may fail for various reasons, e.g. there are too many
-----system calls to filter, the seccomp API is not available,
-----or strace itself is being traced. In cases when seccomp-
-----bpf filter setup failed, strace proceeds as usual and
-----stops traced processes on every system call.

-----When --seccomp-bpf is activated and -p/--attach option is
-----not used, --kill-on-exit option is activated as well.

-----Note that in cases when the tracee has another seccomp
-----filter that returns an action value with a precedence
-----greater than SECCOMP_RET_TRACE, strace --seccomp-bpf will
-----not be notified. That is, if another seccomp filter, for
-----example, disables the syscall or kills the tracee, then
-----strace --seccomp-bpf will not be aware of that syscall
-----invocation at all.

-----tips [=[: id:],[: format:]]
-----Show strace tips, tricks, and tweaks before exit. id can
-----be a non-negative integer number, which enables printing
-----of specific tip, trick, or tweak (these ID are not
-----guaranteed to be stable), or random (the default), in
-----which case a random tip is printed. format can be one of
-----the following:

-----none No tip is printed. Can be used to override the
-----previous setting
-----compact Print the tip just big enough to contain all the
-----text
-----full Print the tip in its full glory.

-----Default is id: random, format: compact.

-----V

```



```

-----version
-----Print the version number of strace. Multiple instances of
-----the option beyond specific threshold tend to increase
-----Strauss awareness.

---Time specification format description
-----Time values can be specified as a decimal floating point number
----- (in a format accepted by strtod(3)), optionally followed by one
----- of the following suffices that specify the unit of time: _s
----- (seconds), _ms (milliseconds), _us (microseconds), or _ns
----- (nanoseconds). If no suffix is specified, the value is
----- interpreted as microseconds.

-----The described format is used for -O, -e inject=delay_enter, and
----- -e inject=delay_exit options.
DIAGNOSTICS
-----When command exits, strace exits with the same exit status. If
----- command is terminated by a signal, strace terminates itself with
----- the same signal, so that strace can be used as a wrapper process
----- transparent to the invoking parent process. Note that parent-
----- child relationship (signal stop notifications, getppid(2) value,
----- etc) between traced process and its parent are not preserved
----- unless -D is used.

-----When using -p without a command, the exit status of strace is
----- zero unless no processes has been attached or there was an
----- unexpected error in doing the tracing.
SETUID INSTALLATION
-----If strace is installed setuid to root then the invoking user will
----- be able to attach to and trace processes owned by any user. In
----- addition setuid and setgid programs will be executed and traced
----- with the correct effective privileges. Since only users trusted
----- with full root privileges should be allowed to do these things,
----- it only makes sense to install strace as setuid to root when the
----- users who can execute it are restricted to those users who have
----- this trust. For example, it makes sense to install a special
----- version of strace with mode 'rwsr-xr--', user root and group
----- trace, where members of the trace group are trusted users. If
----- you do use this feature, please remember to install a regular
----- non-setuid version of strace for ordinary users to use.
MULTIPLE PERSONALITIES SUPPORT
-----On some architectures, strace supports decoding of syscalls for
----- processes that use different ABI rather than the one strace uses.
----- Specifically, in addition to decoding native ABI, strace can
----- decode the following ABIs on the following architectures:

----- [1] When strace is built as an x86_64 application
----- [2] When strace is built as an x32 application
----- [3] Big-endian only

----- This support is optional and relies on ability to generate and
----- parse structure definitions during the build time. Please refer
----- to the output of the strace -V command in order to figure out
----- what support is available in your strace build ("non-native"
----- refers to an ABI that differs from the ABI strace has):

----- m32-mpers
----- strace can trace and properly decode non-native 32-bit
----- binaries.
----- no-m32-mpers
----- strace can trace, but cannot properly decode non-native
----- 32-bit binaries.
----- mx32-mpers
----- strace can trace and properly decode non-native
----- 32-on-64-bit binaries.
----- no-mx32-mpers
----- strace can trace, but cannot properly decode non-native
----- 32-on-64-bit binaries.

----- If the output contains neither m32-mpers nor no-m32-mpers, then
----- decoding of non-native 32-bit binaries is not implemented at all
----- or not applicable.

----- Likewise, if the output contains neither mx32-mpers nor no-
----- mx32-mpers, then decoding of non-native 32-on-64-bit binaries is
----- not implemented at all or not applicable.
NOTES
----- It is a pity that so much tracing clutter is produced by systems
----- employing shared libraries.

----- It is instructive to think about system call inputs and outputs
----- as data flow across the user/kernel boundary. Because user-space
----- and kernel-space are separate and address-protected, it is
----- sometimes possible to make deductive inferences about process
----- behavior using inputs and outputs as propositions.

----- In some cases, a system call will differ from the documented
----- behavior or have a different name. For example, the faccessat(2)
----- system call does not have flags argument, and the setrlimit(2)
----- library function uses prlimit64(2) system call on modern
----- (2.6.38+) kernels. These discrepancies are normal but
----- idiosyncratic characteristics of the system call interface and

```

-----are accounted for by C-library wrapper functions.

-----Some system calls have different names in different architectures and personalities. In these cases, system call filtering and printing uses the names that match corresponding NR*_kernel macros of the tracee's architecture and personality. There are two exceptions from this general rule: arm.fadvise64-64(2) ARM syscall and xtensa.fadvise64-64(2) Xtensa syscall are filtered and printed as fadvise64-64(2).

-----On x32, syscalls that are intended to be used by 64-bit processes and not x32 ones (for example, readv(2), that has syscall number 19 on x86_64, with its x32 counterpart has syscall number 515), but called with X32.SYSCALL_BIT flag being set, are designated with #64 suffix.

-----On some platforms a process that is attached to with the -p option may observe a spurious EINTR return from the current system call that is not restartable. (Ideally, all system calls should be restarted on strace attach, making the attach invisible to the traced process, but a few system calls aren't. Arguably, every instance of such behavior is a kernel bug.) This may have an unpredictable effect on the process if the process takes no action to restart the system call.

-----As strace executes the specified command directly and does not employ a shell for that, scripts without shebang that usually run just fine when invoked by shell fail to execute with ENOEXEC error. It is advisable to manually supply a shell as a command with the script as its argument.

BUGS

-----Programs that use the setuid bit do not have effective user ID privileges while being traced.

-----A traced process runs slowly (but check out the --seccomp-bpf option).

-----Unless --kill-on-exit option is used (or --seccomp-bpf option is used in a way that implies --kill-on-exit), traced processes which are descended from command may be left running after an interrupt signal (CTRL-C).

-----By using CLONE_UNTRACED flag of clone system call a tracee can break the guarantee that --seccomp-bpf will not leave any processes with a seccomp program installed for syscall filtering purposes.

HISTORY

-----The original strace was written by Paul Kranenburg for SunOS and was inspired by its trace utility. The SunOS version of strace was ported to Linux and enhanced by Branko Lankester, who also wrote the Linux kernel support. Even though Paul released strace 2.5 in 1992, Branko's work was based on Paul's strace 1.5 release from 1991. In 1993, Rick Sladkey merged strace 2.5 for SunOS and the second release of strace for Linux, added many of the features of truss(1) from SVR4, and produced an strace that worked on both platforms. In 1994 Rick ported strace to SVR4 and Solaris and wrote the automatic configuration support. In 1995 he ported strace to Irix and became tired of writing about himself in the third person.

-----Beginning with 1996, strace was maintained by Wichert Akkerman. During his tenure, strace development migrated to CVS, ports to FreeBSD and many architectures on Linux (including ARM, IA-64, MIPS, PA-RISC, PowerPC, s390, SPARC) were introduced. In 2002, the burden of strace maintainership was transferred to Roland McGrath. Since then, strace gained support for several new Linux architectures (AMD64, s390x, SuperH), bi-architecture support for some of them, and received numerous additions and improvements in syscalls decoders on Linux; strace development migrated to Git during that period. Since 2009, strace is actively maintained by Dmitry Levin. strace gained support for AArch64, ARC, AVR32, Blackfin, Meta, Nios-II, OpenRISC1000, RISC-V, Tile/TileGx, Xtensa architectures since that time. In 2012, unmaintained and apparently broken support for non-Linux operating systems was removed. Also, in 2012 strace gained support for path tracing and file descriptor path decoding. In 2014, support for stack trace printing was added. In 2016, syscall fault injection was implemented.

-----For the additional information, please refer to the NEWS file and strace repository commit log.

REPORTING BUGS

-----Problems with strace should be reported to the strace mailing list mailto:strace-devel@lists.strace.io.

SEE ALSO

-----strace-log-merge(1), ltrace(1), perf-trace(1), trace-cmd(1), time(1), ptrace(2), seccomp(2), syscall(2), proc(5), signal(7)

-----strace Home Page <https://strace.io/>

AUTHORS

-----The complete list of strace contributors can be found in the CREDITS file.

```

COLOPHON
-----This page is part of the strace (system call tracer) project.
-----Information about the project can be found at
-----http://strace.io/. If you have a bug report for this manual
-----page, send it to strace-devel@lists.sourceforge.net. This page
-----was obtained from the project's upstream Git repository
-----https://github.com/strace/strace.git on 2024-06-14. (At that
-----time, the date of the most recent commit that was found in the
-----repository was 2024-06-04.) If you discover any rendering
-----problems in this HTML version of the page, or you believe there
-----is a better or more up-to-date source for the page, or you have
-----corrections or improvements to the information in this COLOPHON
----- (which is not part of the original manual page), send a mail to
-----man-pages@man7.org

strace 6.9.0.16.2a4c4-----2024-06-04-----STRACE(1)

```

3.10 strings: Print Sequences Of Printable Characters

```

NAME
    strings - print the sequences of printable characters in files

SYNOPSIS
    strings [-afovV] [--min-len]
            [-n min-len] [--bytes=min-len]
            [-t radix] [--radix=radix]
            [-e encoding] [--encoding=encoding]
            [-U method] [--unicode=method]
            [-] [--all] [--print-file-name]
            [-T bfdname] [--target=bfdname]
            [-w] [--include-all-whitespace]
            [-s] [--output-separator sep.string]
            [--help] [--version] file...

DESCRIPTION
    For each file given, GNU strings prints the printable character
    sequences that are at least 4 characters long (or the number
    given with the options below) and are followed by an unprintable
    character.

    Depending upon how the strings program was configured it will
    default to either displaying all the printable sequences that it
    can find in each file, or only those sequences that are in
    loadable, initialized data sections. If the file type is
    unrecognizable, or if strings is reading from stdin then it will
    always display all of the printable sequences that it can find.

    For backwards compatibility any file that occurs after a command-
    line option of just - will also be scanned in full, regardless of
    the presence of any -d option.

    strings is mainly useful for determining the contents of non-text
    files.

OPTIONS
    -a
    --all
    - Scan the whole file, regardless of what sections it contains
    or whether those sections are loaded or initialized.
    Normally this is the default behaviour, but strings can be
    configured so that the -d is the default instead.

    The - option is position dependent and forces strings to
    perform full scans of any file that is mentioned after the -
    on the command line, even if the -d option has been
    specified.

    -d
    --data
    Only print strings from initialized, loaded data sections in
    the file. This may reduce the amount of garbage in the
    output, but it also exposes the strings program to any
    security flaws that may be present in the BFD library used to
    scan and load sections. Strings can be configured so that
    this option is the default behaviour. In such cases the -a
    option can be used to avoid using the BFD library and instead
    just print all of the strings found in the file.

    -f
    --print-file-name
    Print the name of the file before each string.

    --help
    Print a summary of the program usage on the standard output
    and exit.

    -min-len
    -n min-len
    --bytes=min-len
    Print sequences of displayable characters that are at least
    min-len characters long. If not specified a default minimum

```

```

length of 4 is used. The distinction between displayable and
non-displayable characters depends upon the setting of the -e
and -U options. Sequences are always terminated at control
characters such as new-line and carriage-return, but not the
tab character.

-o Like -t o. Some other versions of strings have -o act like
-t d instead. Since we can not be compatible with both ways,
we simply chose one.

-t radix
--radix=radix
    Print the offset within the file before each string. The
    single character argument specifies the radix of the
    offset ---o for octal, x for hexadecimal, or d for decimal.

-e encoding
--encoding=encoding
    Select the character encoding of the strings that are to be
    found. Possible values for encoding are: s =
    single-7-bit-byte characters (default), S = single-8-bit-byte
    characters, b = 16-bit bigendian, l = 16-bit littleendian, B
    = 32-bit bigendian, L = 32-bit littleendian. Useful for
    finding wide character strings. (l and b apply to, for
    example, Unicode UTF-16/UCS-2 encodings).

-U [d|i|l|e|x|h]
--unicode=[default|invalid|locale|escape|hex|highlight]
    Controls the display of UTF-8 encoded multibyte characters in
    strings. The default (--unicode=default) is to give them no
    special treatment, and instead rely upon the setting of the
    --encoding option. The other values for this option
    automatically enable --encoding=S.

    The --unicode=invalid option treats them as non-graphic
    characters and hence not part of a valid string. All the
    remaining options treat them as valid string characters.

    The --unicode=locale option displays them in the current
    locale, which may or may not support UTF-8 encoding. The
    --unicode=hex option displays them as hex byte sequences
    enclosed between <> characters. The --unicode=escape option
    displays them as escape sequences (\uxxxx) and the
    --unicode=highlight option displays them as escape sequences
    highlighted in red (if supported by the output device). The
    colouring is intended to draw attention to the presence of
    unicode sequences where they might not be expected.

-T bfdname
--target=bfdname
    Specify an object code format other than your system's
    -----default-format.

-----v
-----V
-----version
-----Print the program version number on the standard output and
-----exit.

-----w
-----include-all-whitespace
    -----By default, tab and space characters are included in the
    -----strings that are displayed, but other whitespace characters,
    -----such as newlines and carriage returns, are not. The -w option
    -----changes this so that all whitespace characters are considered
    -----to be part of a string.

-----s
-----output-separator
    -----By default, output strings are delimited by a new-line. This
    -----option allows you to supply any string to be used as the
    -----output record separator. Useful with
    -----include-all-whitespace where strings may contain new-lines
    -----internally.

-----@file
    -----Read command-line options from file. The options read are
    -----inserted in place of the original @file option. If file does
    -----not exist, or cannot be read, then the option will be treated
    -----literally, and not removed.

    -----Options in file are separated by whitespace. A whitespace
    -----character may be included in an option by surrounding the
    -----entire option in either single or double quotes. Any
    -----character (including a backslash) may be included by
    -----prefixing the character to be included with a backslash. The
    -----file may itself contain additional @file options; any such
    -----options will be processed recursively.
SEE_ALSO
-----ar(1), nm(1), objdump(1), ranlib(1), readelf(1) and the Info
-----entries for binutils.
COPYRIGHT

```

```

-----Copyright (c) 1991-2024 Free Software Foundation, Inc.

-----Permission is granted to copy, distribute and/or modify this
-----document under the terms of the GNU Free Documentation License,
-----Version 1.3 or any later version published by the Free Software
-----Foundation; with no Invariant Sections, with no Front-Cover
-----Texts, and with no Back-Cover Texts. A copy of the license is
-----included in the section entitled "GNU Free Documentation
-----License".
COLOPHON
-----This page is part of the binutils (a collection of tools for
-----working with executable binaries) project. Information about the
-----project can be found at http://www.gnu.org/software/binutils/.
-----If you have a bug report for this manual page, see
-----http://sourceware.org/bugzilla/enter\_bug.cgi?product=binutils.
-----This page was obtained from the tarball binutils-2.42.tar.gz
-----fetched from https://ftp.gnu.org/gnu/binutils/ on 2024-06-14.
-----If you discover any rendering problems in this HTML version of
-----the page, or you believe there is a better or more up-to-date
-----source for the page, or you have corrections or improvements to
-----the information in this COLOPHON (which is not part of the
-----original manual page), send a mail to man-pages@man7.org

binutils-2.42-----2024-06-14-----STRINGS(1)

```